
A Robot Policy Designed as a Weight-Space Artifact Exact Tensor Structure in a Capable VLA

Rational conversion, reduced-Gram analysis, and exact weight access

Anonymous Author(s)

Affiliation

Address

email

Abstract

Neural checkpoints are usually treated as opaque containers for functions. We instead design a vision-language-action checkpoint as an analyzable weight-space artifact. χ -VLA uses only linear, bilinear, and rational learned operations, so its parameters define an explicit rational tensor program. A runtime audit of the seed-0 20.1M ViT checkpoint finds no incompatible learned operation, while weight-derived reconstructions reproduce deployed branches near the float32 precision floor. Exact reduced-Gram analysis derives attention subspaces directly from trained coefficients without materializing the full high-order tensor. Its rank-128 subspaces beat fixed equal-rank random controls in all 36 primary block-head comparisons. The artifact is also behaviorally meaningful. Three policies reach $85.3\% \pm 4.6\%$ closed-loop success on LIBERO-Object, and a fixed three-pass ensemble reaches 93.4%. In a separate same-recipe experiment over 40 familiar tasks, three joint checkpoints reach $70.35\% \pm 1.32\%$ task-macro success versus $69.12\% \pm 2.33\%$ for conventional controls differing by only 0.006% in parameter count. Prospectively frozen audits further show bounded familiar-task language-path necessity, including a drop from 87.7% to 21.0% when only post-lookup token vectors are zeroed. An independent sequential replay rebuilds the learned path through all 12 transformer blocks to the action output on 48 fixed inputs with worst relative error 3.9822×10^{-7} . The anonymous artifact binds raw outcomes, checkpoint metadata, provenance, and independently recomputed gates. It supports exact layerwise structural access in a capable familiar-task policy family, not unseen-task generalization, validated coefficient editing, or a compact input-general symbolic decomposition of the complete policy.

1 Introduction

Neural checkpoints are usually treated as opaque containers for functions. We can run them and collect their activations, but the weights alone rarely provide a readable map from input features to outputs. A post-hoc probe may reveal that an activation correlates with behavior, yet it does not show how the weights construct that feature.

A tensor is a multidimensional table, and tensor decomposition is the analogue of low-rank matrix factorization for such tables. When a network uses only additions, multiplications, and ratios of polynomials, its weights can be expanded into coefficient tensors whose entries record input interactions. Decomposition ranks simpler factors inside those coefficients. The intended benefit is an inspectable model artifact: candidate structure comes from the trained parameters. By fully tensor-decomposable, we mean that every learned block preserves this algebraic form. This is an operator-closure claim, not a claim that we materialize one compact end-to-end coefficient tensor. Its

degree and representation size grow rapidly with depth. For example, one coefficient can connect an image coordinate, an instruction coordinate, and an action coordinate. Factoring the coefficient table ranks high-energy cross-modal interactions directly from the checkpoint without first fitting an activation probe. Those factors are candidate structure, not automatically human-readable concepts or causal explanations. Softmax, ordinary activations, and square-root normalization obstruct this direct conversion in conventional networks. Bilinear language and vision models suggest that changing the primitives can recover it [3, 4]. Whether the resulting artifact can still control a robot is unresolved.

We study two linked questions:

Can the learned weights of a compact VLA specify an explicit rational tensor program, and can exact weight analysis coexist with strong closed-loop control?

Our contributions are:

1. We construct a complete VLA from tensor-compatible primitives and runtime-audit its seed-0 20.1M ViT instantiation. A rational normalization supplies input-specific scaling without leaving the rational function class.
2. We derive exact weight-based structure through each softmax-free attention layer and an $O(d^2)$ reduced-Gram calculation that avoids materializing its high-order coefficient tensor.
3. We test exact attention subspaces against fixed equal-rank random controls. The weight-derived subspaces win all 36 primary block-head comparisons in a seed-0 checkpoint.
4. We evaluate the same three ViT checkpoints on 1,500 canonical LIBERO-Object trials. They reach $85.3\% \pm 4.6\%$, while a fixed prediction-mean ensemble reaches $467/500 = 93.4\%$. Three jointly trained checkpoints average $70.35\% \pm 1.32\%$ over 40 familiar tasks, 1.23 points above matched conventional controls. Their fixed ensemble reaches 74.75%. A frozen seed-0 comparison retains 95.16% of four suite specialists.

Artifact query	Returned object	Current guarantee
Runtime graph census	Operator inventory and module identities	Membership for the audited seed-0 20.1M ViT checkpoint
Attention-coefficient query	Reduced Gram and leading directions	Exact layerwise construction from saved weights
Numerical replay	Per-case residuals and frozen gates	Independent implementation check on recorded inputs
Artifact integrity	Result and source digests	Immutable SHA-256 manifest and one-command verification

Table 1: **Checkpoint interface as a data artifact.** The trained checkpoint supports structured queries beyond ordinary forward execution. Guarantees remain layerwise and scoped to the named artifacts.

This paper does not claim current state of the art. It also does not claim that comparable mechanisms are impossible to obtain from conventional policies. The empirical question is whether the algebraic policy makes exact weight structure directly accessible.

2 A rational tensor robot policy

2.1 Tensor-compatible primitives

Each primitive replaces a familiar transformer component while keeping explicit coefficients. Linear maps contribute first-order terms, while multiplying learned projections creates higher-order interactions that remain determined by the weights.

Homogeneous coordinate. We write $\bar{x} = [1; x]$. The constant coordinate allows a multiplicative layer to represent bias, linear, and quadratic terms in one contraction.

69 **Bilinear feed-forward block.** For $L, R \in \mathbb{R}^{r \times (d+1)}$ and $D \in \mathbb{R}^{d \times r}$,

$$\text{BFFN}(x) = D[(L\bar{x}) \odot (R\bar{x})], \quad y_o = \sum_{i,j} \left(\sum_{\rho} D_{o\rho} L_{\rho i} R_{\rho j} \right) \bar{x}_i \bar{x}_j. \quad (1)$$

70 The inner sum is a table indexed by one output and two input coordinates. The three matrices give its
71 CP factorization, the tensor analogue of a low-rank matrix factorization.

72 **Softmax-free bilinear attention.** Each head forms two query-key systems and one value stream:

$$Y = C [M \odot (Q_1 K_1^\top) \odot (Q_2 K_2^\top) / s] V. \quad (2)$$

73 Here M is a fixed causal mask, C is fixed row scaling, and s is fixed. All learned query, key,
74 value, and output maps are linear. Equation 2 is therefore a polynomial in its input activations, with
75 interaction coefficients determined by the weights.

76 **Rational per-instance normalization.** RMS normalization contains an input-dependent square
77 root. Let $m(x) = d^{-1} \sum_{j=1}^d x_j^2 + \epsilon$, and let $s_0 > 0$ be the frozen running mean-square scale. We
78 deploy a fixed rational approximation $r(v) = P(v)/Q(v)$ to $v^{-1/2}$ over $v \in [0.1, 10]$:

$$v = \frac{m(x)}{s_0}, \quad \text{RNorm}(x) = \frac{x}{\sqrt{s_0}} r(v). \quad (3)$$

79 Here rational means a ratio of polynomials. Projective coordinates carry each activation as a
80 numerator-denominator pair (N, D) representing N/D . Updating this pair preserves input-specific
81 rescaling without leaving the rational tensor representation. Training and deployment use r itself. The
82 exactness claims below therefore concern faithful evaluation of the deployed rational computation,
83 not uniform agreement with RMSNorm. Appendix A.4 certifies pole-free denominators and deployed
84 arithmetic fidelity while retaining that boundary.

85 2.2 Architecture and training

86 Both encoder variants follow the high-level VLA sequence of visual encoding, language and state
87 embedding, joint processing, and action decoding [1]. They use a 64×64 agent-view image, an
88 eight-dimensional robot state, a 26-token task vocabulary, and eight action-query tokens. Both
89 share a width-384, eight-block, twelve-head joint backbone and a linear head that maps each action-
90 query state to seven continuous action dimensions. The 20.1M ViT uses four bilinear-attention vision
91 blocks and has 20,137,352 learned parameters. It supplies the closed-loop capability, runtime operator
92 census, exact-attention, and reduced-Gram results. The 19.2M convolutional variant instead uses three
93 stride-2 bilinear convolution stages that produce an 8×8 token grid. Three separate convolutional
94 files receive the joint-attention and rational-normalizer certificates reported in Appendix A.4.

95 Training uses 66,984 demonstration frames, 40,000 AdamW updates, batch size 256, cosine decay,
96 bfloat16 arithmetic, and an exponential moving average. Closed-loop evaluation uses a 180° camera
97 rotation to match the demonstration convention, canonical LIBERO initial states, and fail-loud checks
98 that verify every requested state was loaded.

99 2.3 Mechanical certificate on the seed-0 20.1M ViT checkpoint

100 Architecture definitions can differ from deployed code. We therefore audit the deployed seed-0 20.1M
101 ViT checkpoint at runtime and reconstruct representative operations from saved weights.

Mechanical operator audit of the seed-0 20.1M ViT checkpoint




Figure 1: **Mechanical audit of the seed-0 20.1M ViT rational checkpoint.** The runtime path invokes 112 linear maps, 74 rational normalizers, two embeddings, and one linear patch projection. No softmax, standard activation, or ordinary normalization is present. Reconstruction errors are measured over 54,784 real activation rows.

102 The normalization reconstruction has maximum absolute error 1.64×10^{-6} . A bilinear FFN re-
 103 construction has median relative error 1.40×10^{-7} . These residuals are bounded by the deployed
 104 module’s float32 statistic. The audit certifies operator membership and local reconstruction. It does
 105 not materialize one compact polynomial for the full eight-layer transformer.

106 3 The artifact is a capable closed-loop policy

107 3.1 Evaluation protocol

108 We evaluate all ten LIBERO-Object tasks with the following fixed protocol:

- 109 • 50 canonical trials per task
- 110 • a 280-step cap
- 111 • three independently trained ViT checkpoints
- 112 • 500 rollouts per checkpoint and 1,500 rollouts in total




Figure 2: **Verified ViT capability.** Left: three independently trained checkpoints and their fixed elementwise prediction-mean ensemble. Right: per-task three-checkpoint mean and ensemble success. The ensemble uses three forward passes.

3.2 Results

Method	Parameters	Inference	Success
χ -VLA, rational ViT	20.1M	1 pass	85.3% $\pm$ 4.6%
χ -VLA, prediction-mean ensemble	$3 \times 20.1M$	3 passes	93.4%

Table 2: **Closed-loop LIBERO-Object success.** Each ViT checkpoint receives 500 canonical trials. The single-model row is the mean $\pm$ sample standard deviation across three checkpoints. The ensemble is one fixed composite evaluated on the same 500 states.

The three checkpoints reach $426/500 = 85.2\%$, $404/500 = 80.8\%$, and $450/500 = 90.0\%$. Elementwise averaging of their predicted action chunks reaches $467/500 = 93.4\%$. This is 8.1 points above the member mean and 3.4 points above the strongest member. The ensemble requires three complete forward passes and is not one 20.1M-parameter policy. Every episode-level record, task index, and protocol field is retained. The artifact manifest binds each immutable result file and recorded checkpoint basename to a checkpoint hash verified separately. Member and ensemble shards span different recorded GPU families and predate explicit matrix-precision fields, so the arithmetic gaps are descriptive rather than a matched-hardware ensemble ablation.

3.3 What this result does and does not establish

The result shows that the rational architecture supports high closed-loop specialist capability. It is an existence claim for a compact tensor-decomposable policy, not a claim of broad robotic generalization or state of the art.

3.4 One policy across four familiar suites

We next train one 20.1M χ -VLA checkpoint jointly on the ten tasks from each of LIBERO-Object, LIBERO-Spatial, LIBERO-Goal, and LIBERO-10. The balanced recipe uses 160,000 updates, equal to the aggregate $4 \times 40,000$ updates used by four separately trained suite specialists. Evaluation uses 50 canonical trials per task, so each checkpoint receives 2,000 closed-loop trials.

Policy	Object	Spatial	Goal	LIBERO-10	40-task macro
Joint χ -VLA, three-seed mean	80.9%	73.1%	81.9%	45.5%	70.35% $\pm$ 1.32%
Matched conventional, three-seed mean	79.2%	77.0%	84.4%	35.9%	69.12% $\pm$ 2.33%
Fixed three-checkpoint mean ensemble	88.6%	74.0%	85.6%	50.8%	74.75%
Joint χ -VLA, seed 0	83.6%	73.0%	82.0%	44.8%	70.85%
Four suite specialists, seed 0	83.0%	80.0%	88.0%	46.8%	74.45%

Table 3: **Closed-loop breadth over 40 familiar tasks.** The first row is the mean across three jointly trained checkpoints, with sample standard deviation across their 40-task macro scores. The matched conventional row uses the same balanced recipe. The ensemble averages normalized χ -VLA action predictions and costs three forward passes. The final two rows form the separately frozen seed-0 retention comparison.

131 The three joint checkpoints score 70.85%, 71.35%, and 68.85%, and their seed-mean success is
132 at least 50% on 32/40 tasks. In the separately frozen same-recipe comparison, three near-equal-
133 parameter conventional checkpoints average 69.12% $\pm$ 2.33%. The χ -VLA mean is 1.23 percentage
134 points higher and passes the fixed requirement of being within five points of the conventional mean.
135 The model sizes differ by only 1,280 parameters, or 0.006%. In the separately frozen mean-prediction
136 ensemble, success rises from the 70.35% member mean to 74.75%, with every suite at or above 50%
137 and 34/40 tasks at or above 50%. Its protocol and gates were fixed after the member checkpoints
138 were selected but before ensemble evaluation. In the separately frozen seed-0 comparison, the single
139 joint checkpoint retains 70.85/74.45 = 95.16% of the four specialists’ pooled task-macro success,
140 above the fixed 85% retention threshold. The suite ratios are 100.7%, 91.3%, 93.2%, and 95.7% for
141 Object, Spatial, Goal, and LIBERO-10. This certificate compares familiar tasks seen during training.
142 It establishes multi-suite capacity and parameter sharing, not transfer to an unseen task, instruction,
143 embodiment, or domain.

144 3.5 Prospectively frozen closed-loop instruction necessity

145 High task success alone does not show that a policy uses its instruction. Before these outcomes
146 were observed, we fixed a paired test on canonical episodes 40–49 for every task and the same three
147 checkpoints. Each checkpoint-state pair was rolled out for the full 280-step horizon under the correct
148 familiar instruction, one deterministic outcome-independent co-present control instruction, and an
149 empty instruction. All conditions shared the exact initial physical input and their order was chosen by
150 a frozen hash. Co-presence was verified from simulator observation fields and does not guarantee
151 unoccluded camera visibility. The empty string encodes tokenizer BOS plus 31 unmasked PAD
152 tokens, together with the model’s learned common BOS.




Figure 3: **Closed-loop instruction necessity on familiar tasks.** Left: exact successes among 100 paired canonical states per checkpoint. Right: correct-instruction success minus each control, pooling the three checkpoint outcomes for each task. All ten task gaps are positive for both controls at the task-aggregate level.

153 Across 100 unique canonical states and three checkpoints, the correct instruction succeeds on
154 247/300 = 82.3% of checkpoint-state pairs. The co-present control succeeds on 72/300 = 24.0%

and the empty instruction on $78/300 = 26.0\%$, giving paired gaps of 58.3 and 56.3 points. Correct-instruction success is 83/100, 77/100, and 87/100 by checkpoint, and both paired gaps are positive for every checkpoint aggregate. All ten pooled task aggregates are strictly positive for both controls. One-sided exact McNemar tests give $p = 2.02 \times 10^{-45}$ and $p = 1.52 \times 10^{-44}$. These pooled tests use 300 checkpoint-state pairs and are not cluster-robust inference. A frozen 20,000-draw task-stratified state-cluster bootstrap resamples 100 unique task-episode states while retaining all three checkpoint outcomes. Its 95% intervals are [51.0, 65.7] and [50.0, 62.3] points. Every prospectively frozen gate passes.

This establishes instruction necessity for familiar LIBERO-Object prompts, objects, scenes, tasks, and checkpoints. Appendix A.1 gives a disjoint paired test that isolates the learned lexical-embedding path. Neither test establishes counterfactual goal following, paraphrase robustness, unseen-object or compositional grounding, cross-suite or physical transfer, or a benefit unique to tensor decomposability.

For Neural Artifacts, this adds a prospectively specified behavioral audit to the algebraic reconstruction checks. The result constrains checkpoint behavior, but it does not make the decomposed factors human-nameable or semantic.

4 Exact weight-derived structure through attention

4.1 Exact layerwise construction

Equation 2 contains two query-key products and one value stream. Expanding the linear maps gives a coefficient table whose entries measure how combinations of query, key, and value coordinates contribute to an output. No examples or fitted probe are needed to construct it. On a tiny four-token, width-four, one-head layer, direct evaluation and an independent explicit polynomial agree to 2.00×10^{-15} . The internal core identity agrees to 1.24×10^{-14} .

The deployed attention also normalizes each query and key with Equation 3. Each normalization multiplies all channels of one token by the same scalar rational factor. Those factors can be pulled outside the bilinear dot products. The polynomial attention core is unchanged, while explicit numerator and denominator factors carry the normalization. The rational extension agrees to 3.61×10^{-16} . On one fixed corrected cached input per checkpoint, we numerically replay all four vision and eight joint attention modules in each of three records that name the same ViT checkpoint artifacts as the capability results. This covers 36 module-input cases and 384 architectural heads. The worst relative ℓ_2 error is 2.016119×10^{-7} and the worst absolute difference is 2.44×10^{-5} . This residual comes from the deployed implementation’s intentional float32 statistic, while the reconstruction itself uses the learned coefficients. The algebraic identity is input-general. The one-input-per-checkpoint records provide a layerwise numerical replay audit, not an exact end-to-end policy decomposition.

We separately certify three convolutional checkpoints from the same rational architecture family. Their vision stacks contain no attention, so we independently rebuild all eight joint modules and all twelve heads per module from the saved weights. Sixteen deterministically selected cache inputs per checkpoint yield 4,608 head-input cases. Every output is finite and the worst relative ℓ_2 error is 6.36×10^{-16} , below the frozen 10^{-6} gate. On the same inputs, independent CP-factor evaluation covers all 384 joint-FFN module-input cases at worst relative error 1.3871×10^{-15} against a 10^{-10} gate. Independent complete-block equations cover all 384 block-input cases at worst relative error 1.5024×10^{-7} against a 10^{-4} gate. The complete-block comparison crosses a deployed PyTorch float32 block call and a NumPy float64 reconstruction. It is input-conditioned and per-block, not a sequential end-to-end reconstruction of the joint transformer or policy.

A separately frozen ViT audit partitions joint attention into vision, instruction, robot-state, and action-query source tokens. Across 128 deterministic, official-task-balanced inputs for each of the three capability checkpoints, all 36,864 head-input cases and 147,456 source-group evaluations are finite, with worst reconstruction error 5.0741×10^{-7} against a 10^{-6} gate. Across the 3,072 projected module-input cases, mean coherent-energy shares are 67.7% vision, 19.2% robot state, 7.9% action query, and 5.1% instruction. These input-conditioned magnitudes are descriptive rather than causal importance. Familiar-prompt permutations are also descriptive and do not establish grounding or closed-loop success.

Finally, a prospectively frozen sequential replay reaches the learned action output on the same three capable ViT checkpoints. For 16 deterministic, official-task-balanced corrected inputs per checkpoint, an independent NumPy float64 evaluation starts from the exact prepared model tensors and rebuilds the learned patch projection, four vision blocks, multimodal embeddings and projections, eight joint blocks, final RationalNorm, and linear action head. Across 48 inputs and 816 recorded stage evaluations, the worst final-action relative ℓ_2 error is 3.9822×10^{-7} against a frozen 10^{-3} gate. Every manual PyTorch traversal is bitwise identical to an unmodified full model call. This is an input-conditioned full learned-forward numerical certificate, not one compact symbolic contraction or an input-general proof. It excludes preprocessing, fixed action unnormalization, chunk execution, simulator dynamics, and success.

Naively materializing the real width-384 coefficient tensor would require approximately 3.2×10^{15} values. We instead form an $O(d^2)$ reduced Gram. Like a covariance matrix, its leading eigenvectors rank important directions without constructing the full table. It matches brute-force materialization to 2.13×10^{-14} at toy scale.

4.2 Trained policy result

The reconstruction certificate above spans three ViT checkpoints. The following subspace-ranking experiment is a descriptive analysis of the seed-0 ViT checkpoint. We apply the exact weight-derived calculation to all twelve heads in early, middle-late, and final backbone blocks 0, 6, and 7. Random projector banks are sampled once, then fixed and reused for all examples. The weights choose the subspace, while cached activations are used only afterward to score how well it preserves the layer’s computation.




Figure 4: Exact weight-derived attention structure. Ratios above one favor the exact weight-derived subspace over a random equal-rank subspace. Each point in the right panel is one trained attention head.

At ranks 8, 32, and 128, respectively, 80.6%, 91.7%, and 100% of the 36 block-head pairs favor the exact subspace. Their median ratios are 1.028, 1.171, and 2.441. At rank 128, the mean ratio is 3.033 and the strongest head reaches 10.135. The rank-128 block means are 3.08, 3.45, and 2.56 for blocks 0, 6, and 7. The 36 block-head pairs are architectural units within one checkpoint, not independent model-level replicates.

This extends exact weight-only construction through individual attention layers. It is not an exact compact symbolic end-to-end decomposition of the complete policy. Symbolic composition across all eight blocks still faces rapid degree and representation growth.

4.3 What the artifact lets a reviewer test

The structural artifact separates four questions that can otherwise be conflated. First, does a factorized operator implement its claimed equation? The attention and FFN records answer this by comparing independent evaluations of the saved factors. Second, does the claimed equation survive the numerical boundary of the deployed implementation? The complete-block records compare a NumPy float64 reconstruction with a deployed PyTorch float32 block call. Third, does the learned path preserve

that agreement under sequential composition to the action output? The full-forward record answers this numerically on fixed prepared inputs. Fourth, is the deployed rational normalizer pole-free and arithmetically faithful? The RationalNorm record recomputes normalized and physical actions from stored arrays. Passing one level does not imply the next, so the verifier reports each decision separately.

This ladder also clarifies what a failed control would mean. Failure of an operator identity would invalidate the corresponding convertibility claim before any behavioral interpretation. Agreement at the operator and block levels would instead localize a later discrepancy to composition, preprocessing, decoding, or simulator interaction. The complete-block certificate remains per-block. The separate full learned-forward record closes the learned numerical path on fixed prepared inputs, but it still excludes preprocessing, action unnormalization, chunk execution, and simulator dynamics. This distinction is especially important for multimodal policies, where an exact internal equation can coexist with an out-of-graph mode choice or a continuous-action convention that changes behavior.

The modality ledger serves a different role. It assigns every joint-attention term to vision, instruction, robot-state, or action-query source tokens and verifies that the four reconstructed groups sum to the deployed output on fixed inputs. The reported energy shares expose where the trained computation is concentrated, but they are not causal importance scores. Prompt permutations likewise provide a reproducible sensitivity distribution, not evidence of grounding. A reviewer can therefore reproduce the algebraic statement without inheriting a stronger semantic interpretation.

Every record is tied to frozen checkpoint, cache, task, provenance, runtime, and source identities. Inputs are selected by deterministic hashes rather than by observed residuals. The standard-library verifier then recomputes row counts, extrema, distributions, and gate decisions from the released rows. These controls make the artifact useful beyond this model. A future conversion method can reuse the same ladder and state exactly which boundary its certificate crosses.

5 Related work and limitations

Related work. OpenVLA and OpenVLA-OFT establish pretrained VLA and continuous-control baselines [1, 2]. Bilinear MLPs, χ -nets, and bilinear IMPALA policies show that multiplicative weights can expose low-rank structure and support subsequent controls [3–5]. Recent work also steers pretrained VLAs through activation directions [6]. Our distinction is the exact weight-derived construction in a policy designed as a rational tensor program. Tensor and polynomial networks provide the broader multilinear foundation [10–12]. VLA robustness benchmarks show that standard LIBERO success can coexist with weak language use [7–9].

Limitations and conclusion. The breadth benchmark covers 40 familiar training tasks across four suites, so it does not establish unseen-task generalization. Exact layer equations are certified directly, and the learned forward path is independently replayed sequentially on fixed inputs. This does not yet produce one compact symbolic contraction or an input-general proof of the full policy. The ViT records connect capability, exact attention, and the reduced-Gram screen through common checkpoint basenames. Separate convolutional files replicate the joint-attention equation certificate across another visual encoder. A paired familiar-task audit establishes instruction necessity for the original goal, but not counterfactual goal following, zero-shot language use, or broad grounding. Within that scope, rational restrictions support strong LIBERO-Object control, substantial multi-suite familiar-task breadth, and exact trained-weight certificates. We do not claim a direct coefficient-edit interface.

Reproducibility. The anonymous artifact releases raw structural-certificate JSONs and all 2,000 episode rows behind the ViT member and ensemble results. Immutable SHA-256 manifests and standard-library verifiers independently recompute capability totals, the matched conventional comparison, four-suite ensemble and retention arithmetic, visual-bottleneck counts and intervals, and all structural-certificate decisions. Historical Modal analyses read dataset-native task metadata. The Athena replacement pins this metadata and maps dataset to official identifiers before task splits. A source audit matched task identifiers, language joins, states, actions, and resized images for all 271,996 cached frames to their pinned LeRobot revisions. Canonical-state checks fail loudly when requested states are unavailable.

References

- [1] M. J. Kim et al. *OpenVLA: An Open-Source Vision-Language-Action Model*. CoRL, 2024.
- [2] M. J. Kim, C. Finn, and P. Liang. *Fine-Tuning Vision-Language-Action Models: Optimizing Speed and Success*. 2025.
- [3] A. Rigg, T. Doods, M. T. Pearce, J. Oramas, and L. Sharkey. *Bilinear MLPs Enable Weight-Based Mechanistic Interpretability*. ICLR, 2025.
- [4] T. Doods, W. Gauderis, G. A. Wiggins, and J. Oramas. *Compositionality Unlocks Deep Interpretable Models*. 2025.
- [5] F. Oozeer, Y. Erisken, A. Rigg, J. Oramas, and L. Sharkey. *Bilinear Convolution Decomposition for Causal RL Interpretability*. 2024.
- [6] B. Häon, K. Stocking, I. Chuang, and C. Tomlin. *Mechanistic Interpretability for Steering Vision-Language-Action Models*. 2025.
- [7] X. Zhou et al. *LIBERO-PRO: Towards Robust and Fair Evaluation of Vision-Language-Action Models Beyond Memorization*. 2025.
- [8] Y. Fang et al. *When Vision Overrides Language: Evaluating and Mitigating Counterfactual Failures in VLAs*. 2026.
- [9] C. Kim et al. *LIBERO-Para: A Diagnostic Benchmark and Metrics for Paraphrase Robustness in VLA Models*. 2026.
- [10] T. Kolda and B. Bader. *Tensor Decompositions and Applications*. SIAM Review, 2009.
- [11] E. Stoudenmire and D. Schwab. *Supervised Learning with Tensor Networks*. NeurIPS, 2016.
- [12] G. Chrysos et al. *Deep Polynomial Neural Networks*. CVPR and TPAMI, 2020 to 2021.

A Supporting results and decisions

A.1 Causal isolation of the learned instruction-embedding path

Token IDs are lookup addresses. A learned embedding table turns those addresses into vectors that the policy can combine with vision and robot state. Before observing outcomes, we fixed a paired test that isolates this conversion. On canonical episodes 30–39, disjoint from the instruction-necessity episodes, both arms received the identical correct 32-token ID tensor. The ablated arm replaced only the output of the token-embedding lookup with exact zeros. Sequence length, token positions, vision, robot state, embodiment, the separate learned common BOS, action queries, transformer weights, and action head were preserved and hash-checked.




Figure 5: **Learned lexical-embedding-path necessity on familiar tasks.** Left: closed-loop success for the intact path and exact post-lookup zeroing, using 100 paired states per checkpoint and 300 pooled pairs. Right: every task aggregate has a positive paired gap. The pooled marker shows the frozen task-stratified state-cluster 95% bootstrap interval.

324 The intact path succeeds on $263/300 = 87.7\%$ of checkpoint-state pairs, compared with $63/300 =$
 325 21.0% after zeroing the learned token vectors. The paired gap is 66.7 points, with 203 intact-only
 326 and 3 zeroed-only outcomes. Full-path success is 84%, 87%, and 92% by checkpoint, and the
 327 corresponding gaps are 78, 57, and 65 points. All ten task aggregates are positive. A pooled one-sided
 328 exact paired test gives $p = 1.42 \times 10^{-56}$, but this test is not cluster-robust because checkpoints
 329 reuse the same 100 states. The frozen 20,000-draw task-stratified state-cluster bootstrap interval is
 330 $[61.7, 71.3]$ points and is our cluster-aware inference. All eight prospectively fixed gates pass.

331 This establishes that the learned lexical-embedding path is necessary for these familiar LIBERO-
 332 Object tasks and three fixed policies. It does not establish semantic grounding, paraphrase robustness,
 333 unseen-object or compositional generalization, cross-suite or physical transfer, or a benefit unique to
 334 tensor decomposability.

335 A.2 Expanded attention screen

336 Within the same seed-0 ViT checkpoint, the corrected fixed-control attention analysis covers all
 337 eight blocks and twelve heads. At rank 128, 83/96 block-head pairs favor the exact weight-derived
 338 subspace. The median ratio is 1.94, the mean is 2.58, and blocks 2 and 3 provide useful depth
 339 controls. At ranks 8 and 32, only 54/96 and 60/96 pairs favor the exact subspace. The effect is
 340 therefore depth-dependent and strongest at the wider tested rank.

341 A.3 Prospectively fixed-rank visual bottleneck

342 The behavioral intervention acts at the width-384 visual-to-joint interface and keeps rank 96. Projector
 343 construction uses official Object tasks 0–3. Rank 96 was fixed before any corrected-basis outcome
 344 was observed, and the frozen evaluation uses tasks 8–9. Across three checkpoints and 300 canonical
 345 episodes, the full policy succeeds on 272 trials. The action-Jacobian projector retains $259/272 =$
 346 95.2% of those successes, with a paired checkpoint-task-episode resampling interval of 91.2–98.6%.
 347 The activation-energy projector retains $267/272 = 98.2\%$, with an interval of 95.8–100.0%. Three
 348 fixed equal-rank random projectors retain $37/272 = 13.6\%$, $9/272 = 3.3\%$, and $22/272 = 8.1\%$.




Figure 6: **Corrected closed-loop visual bottleneck.** Both learned rank-96 projectors preserve nearly all baseline successes across three checkpoints, while equal-rank random projectors do not. Error bars in panel (a) are deterministic paired cluster-resampling intervals. Tasks 8–9 are absent from corrected projector construction, but not from policy training. This is not a clean rank-selection-holdout result.

349 The action-Jacobian basis does not outperform the matched activation-energy basis. The stronger
 350 action-specificity gate therefore does not pass. The supported claim is structured visual sufficiency
 351 relative to random projectors. This intervention is data-derived and remains separate from the exact
 352 weight-only attention decomposition.

353 A.4 Numerical scope of rational conversion

354 The deployed degree- $[2/2]$ least-squares rational fit is exact as a rational computation even though it
 355 approximates RMS normalization. Its numerator has coefficients (5.32997, 7.53552, 0.327274) and
 356 the denominator has coefficients (1, 9.64827, 2.60664). Every denominator coefficient is positive,
 357 which certifies $Q(v) > 0$ for all $v \geq 0$. We instrument all 53 RationalNorm module instances,
 358 each invoked once per forward pass, on 4,096 deterministic inputs for each of three convolutional
 359 checkpoints. Across 544,542,720 activation rows, the minimum observed denominator is 1.0016
 360 and no row is nonfinite. Comparing deployed float32 normalization with an otherwise identical pass
 361 that recomputes only RationalNorm scales in float64, the largest checkpoint-level primary metric,
 362 defined as the maximum of normalized- and physical-action NRMSE, is 2.414×10^{-6} against a
 363 frozen 10^{-3} gate. This primary certificate concerns pole-free denominators and floating-point fidelity
 364 to the deployed rational operator. It does not claim uniform agreement with exact RMS scaling or
 365 measure matched alternative-normalizer runtime overhead.

366 The reconstruction ladder extends from individual attention heads through FFNs and per-block replay
 367 to a separate sequential learned-forward certificate. That final record starts at exact prepared model
 368 tensors and reaches the linear action head on 48 fixed inputs. It is numerical and input-conditioned,
 369 not a compact symbolic contraction or an input-general full-policy proof. The ViT source-partition
 370 audit remains layerwise and descriptive. Figure 7 reports the frozen numerical gates and the depth
 371 profile without assigning causal importance.




Figure 7: **Frozen reconstruction ladder and descriptive ViT source partition.** Panel (a) reports observed-to-gate ratios for independent Conv attention, joint-FFN, and complete joint-block reconstruction, the ViT source-partition identity, the sequential ViT learned-forward action replay, and primary RationalNorm action fidelity. All plotted primary numerical gates pass. Panel (b) shows descriptive mean coherent-energy shares after the attention output projection across 384 checkpoint-input cases per joint block. The depth trend is not a causal importance or grounding result. The learned-forward replay excludes preprocessing, fixed action unnormalization, chunk execution, simulator dynamics, and success.