
A Robot Policy Designed as a Weight-Space Artifact Exact Tensor Structure in a Capable VLA

Rational conversion, reduced-Gram analysis, and exact weight access

Anonymous Author(s)

Affiliation

Address

email

Abstract

Tensor decomposition rewrites a multidimensional table as a small set of low-rank factors. In a neural network, those factors can expose which combinations of inputs and outputs are encoded by the weights. Conventional nonlinear operations make such a table difficult to obtain directly. We introduce χ -VLA, a rational vision-language-action model family with a 19.2M convolutional policy and a 20.1M ViT policy, each containing only bilinear, linear, and rational learned operations. Its checkpoint therefore specifies an explicit rational tensor program rather than only an opaque executable function. A runtime operator audit of the seed-0 20.1M ViT checkpoint and local reconstructions verify its operator-level form. We derive an exact weight-based decomposition of each softmax-free attention layer and, on one fixed corrected cached input per checkpoint, numerically replay all 12 modules in each of three ViT checkpoint artifacts with worst relative error 2.016119×10^{-7} . An $O(d^2)$ reduced Gram avoids materializing the coefficient tensor. In a seed-0 ViT checkpoint, all 36 primary block-head tests favor the rank-128 weight-derived subspace over fixed random controls. Across the same three ViT checkpoints, closed-loop LIBERO-Object success is $85.3\% \pm 4.6\%$ over 500 canonical trials per checkpoint, and a fixed prediction-mean ensemble reaches $467/500 = 93.4\%$. Separate Conv audits reconstruct 4,608 attention-head cases, 384 joint-FFN module-input cases, and 384 complete joint-block input cases, with worst relative errors 6.36×10^{-16} , 1.39×10^{-15} , and 1.5024×10^{-7} . Learned rank-96 visual projectors retain 95.2–98.2% of baseline successes, versus 3.3–13.6% for random projectors. We therefore establish a capable model family with exact layerwise structure derived from trained coefficients. We do not yet provide an exact end-to-end decomposition of the complete policy.

1 Introduction

Neural checkpoints are usually treated as opaque containers for functions. We can run them and collect their activations, but the weights alone rarely provide a readable map from input features to outputs. A post-hoc probe may reveal that an activation correlates with behavior, yet it does not show how the weights construct that feature.

A tensor is a multidimensional table, and tensor decomposition is the analogue of low-rank matrix factorization for such tables. When a network uses only additions, multiplications, and ratios of polynomials, its weights can be expanded into coefficient tensors whose entries record input interactions. Decomposition ranks simpler factors inside those coefficients. The intended benefit is an inspectable model artifact: candidate structure comes from the trained parameters. By fully tensor-decomposable, we mean that every learned block preserves this algebraic form. This is an

operator-closure claim, not a claim that we materialize one compact end-to-end coefficient tensor. Its degree and representation size grow rapidly with depth. For example, one coefficient can connect an image coordinate, an instruction coordinate, and an action coordinate. Factoring the coefficient table ranks high-energy cross-modal interactions directly from the checkpoint without first fitting an activation probe. Those factors are candidate structure, not automatically human-readable concepts or causal explanations. Softmax, ordinary activations, and square-root normalization obstruct this direct conversion in conventional networks. Bilinear language and vision models suggest that changing the primitives can recover it [3, 4]. Whether the resulting artifact can still control a robot is unresolved.

We study two linked questions:

Can the learned weights of a specialist VLA specify an explicit rational tensor program, and can exact weight analysis coexist with strong closed-loop control?

Our contributions are:

1. We construct a complete VLA from tensor-compatible primitives and runtime-audit its seed-0 20.1M ViT instantiation. A rational normalization supplies input-specific scaling without leaving the rational function class.
2. We derive exact weight-based structure through each softmax-free attention layer and an $O(d^2)$ reduced-Gram calculation that avoids materializing its high-order coefficient tensor.
3. We test exact attention subspaces against fixed equal-rank random controls. The weight-derived subspaces win all 36 primary block-head comparisons in a seed-0 checkpoint.
4. We evaluate the same three ViT checkpoints on 1,500 canonical LIBERO-Object trials. They reach $85.3\% \pm 4.6\%$, while a fixed prediction-mean ensemble reaches $467/500 = 93.4\%$.

Artifact query	Returned object	Current guarantee
Runtime graph census	Operator inventory and module identities	Membership for the audited seed-0 20.1M ViT checkpoint
Attention-coefficient query	Reduced Gram and leading directions	Exact layerwise construction from saved weights
Numerical replay	Per-case residuals and frozen gates	Independent implementation check on recorded inputs
Artifact integrity	Result and source digests	Immutable SHA-256 manifest and one-command verification

Table 1: **Checkpoint interface as a data artifact.** The trained checkpoint supports structured queries beyond ordinary forward execution. Guarantees remain layerwise and scoped to the named artifacts.

This paper does not claim current state of the art. It also does not claim that comparable mechanisms are impossible to obtain from conventional policies. The empirical question is whether the algebraic policy makes exact weight structure directly accessible.

2 A rational tensor robot policy

2.1 Tensor-compatible primitives

Each primitive replaces a familiar transformer component while keeping explicit coefficients. Linear maps contribute first-order terms, while multiplying learned projections creates higher-order interactions that remain determined by the weights.

Homogeneous coordinate. We write $\bar{x} = [1; x]$. The constant coordinate allows a multiplicative layer to represent bias, linear, and quadratic terms in one contraction.

Bilinear feed-forward block. For $L, R \in \mathbb{R}^{r \times (d+1)}$ and $D \in \mathbb{R}^{d \times r}$,

$$\text{BFFN}(x) = D[(L\bar{x}) \odot (R\bar{x})], \quad y_o = \sum_{i,j} \left(\sum_{\rho} D_{o\rho} L_{\rho i} R_{\rho j} \right) \bar{x}_i \bar{x}_j. \quad (1)$$

68 The inner sum is a table indexed by one output and two input coordinates. The three matrices give its
 69 CP factorization, the tensor analogue of a low-rank matrix factorization.

70 **Softmax-free bilinear attention.** Each head forms two query-key systems and one value stream:

$$Y = C [M \odot (Q_1 K_1^\top) \odot (Q_2 K_2^\top) / s] V. \quad (2)$$

71 Here M is a fixed causal mask, C is fixed row scaling, and s is fixed. All learned query, key,
 72 value, and output maps are linear. Equation 2 is therefore a polynomial in its input activations, with
 73 interaction coefficients determined by the weights.

74 **Rational per-instance normalization.** RMS normalization contains an input-dependent square
 75 root. Let $m(x) = d^{-1} \sum_{j=1}^d x_j^2 + \epsilon$, and let $s_0 > 0$ be the frozen running mean-square scale. We
 76 deploy a fixed rational approximation $r(v) = P(v)/Q(v)$ to $v^{-1/2}$ over $v \in [0.1, 10]$:

$$v = \frac{m(x)}{s_0}, \quad \text{RNorm}(x) = \frac{x}{\sqrt{s_0}} r(v). \quad (3)$$

77 Here rational means a ratio of polynomials. Projective coordinates carry each activation as a
 78 numerator-denominator pair (N, D) representing N/D . Updating this pair preserves input-specific
 79 rescaling without leaving the rational tensor representation. Training and deployment use r itself. The
 80 exactness claims below therefore concern faithful evaluation of the deployed rational computation,
 81 not uniform agreement with RMSNorm. Appendix A.3 certifies pole-free denominators and deployed
 82 arithmetic fidelity while retaining that boundary.

83 2.2 Architecture and training

84 Both encoder variants follow the high-level VLA sequence of visual encoding, language and state
 85 embedding, joint processing, and action decoding [1]. They use a 64×64 agent-view image, an
 86 eight-dimensional robot state, a 26-token task vocabulary, and eight action-query tokens. Both
 87 share a width-384, eight-block, twelve-head joint backbone and a linear head that maps each action-
 88 query state to seven continuous action dimensions. The 20.1M ViT uses four bilinear-attention vision
 89 blocks and has 20,137,352 learned parameters. It supplies the closed-loop capability, runtime operator
 90 census, exact-attention, and reduced-Gram results. The 19.2M convolutional variant instead uses three
 91 stride-2 bilinear convolution stages that produce an 8×8 token grid. Three separate convolutional
 92 files receive the joint-attention and rational-normalizer certificates reported in Appendix A.3.

93 Training uses 66,984 demonstration frames, 40,000 AdamW updates, batch size 256, cosine decay,
 94 bfloat16 arithmetic, and an exponential moving average. Closed-loop evaluation uses a 180° camera
 95 rotation to match the demonstration convention, canonical LIBERO initial states, and fail-loud checks
 96 that verify every requested state was loaded.

97 2.3 Mechanical certificate on the seed-0 20.1M ViT checkpoint

98 Architecture definitions can differ from deployed code. We therefore audit the deployed seed-0 20.1M
 99 ViT checkpoint at runtime and reconstruct representative operations from saved weights.

Mechanical operator audit of the seed-0 20.1M ViT checkpoint




Figure 1: **Mechanical audit of the seed-0 20.1M ViT rational checkpoint.** The runtime path invokes 112 linear maps, 74 rational normalizers, two embeddings, and one linear patch projection. No softmax, standard activation, or ordinary normalization is present. Reconstruction errors are measured over 54,784 real activation rows.

100 The normalization reconstruction has maximum absolute error 1.64×10^{-6} . A bilinear FFN re-
 101 construction has median relative error 1.40×10^{-7} . These residuals are bounded by the deployed
 102 module’s float32 statistic. The audit certifies operator membership and local reconstruction. It does
 103 not materialize one compact polynomial for the full eight-layer transformer.

104 3 The artifact is a capable closed-loop policy

105 3.1 Evaluation protocol

106 We evaluate all ten LIBERO-Object tasks with the following fixed protocol:

- 107 • 50 canonical trials per task
- 108 • a 280-step cap
- 109 • three independently trained ViT checkpoints
- 110 • 500 rollouts per checkpoint and 1,500 rollouts in total




Figure 2: **Verified ViT capability.** Left: three independently trained checkpoints and their fixed elementwise prediction-mean ensemble. Right: per-task three-checkpoint mean and ensemble success. The ensemble uses three forward passes.

3.2 Results

Method	Parameters	Inference	Success
χ -VLA, rational ViT	20.1M	1 pass	85.3% $\pm$ 4.6%
χ -VLA, prediction-mean ensemble	$3 \times 20.1M$	3 passes	93.4%

Table 2: **Closed-loop LIBERO-Object success.** Each ViT checkpoint receives 500 canonical trials. The single-model row is the mean $\pm$ sample standard deviation across three checkpoints. The ensemble is one fixed composite evaluated on the same 500 states.

The three checkpoints reach $426/500 = 85.2\%$, $404/500 = 80.8\%$, and $450/500 = 90.0\%$. Elementwise averaging of their predicted action chunks reaches $467/500 = 93.4\%$. This is 8.1 points above the member mean and 3.4 points above the strongest member. The ensemble requires three complete forward passes and is not one 20.1M-parameter policy. Every episode-level record, task index, and protocol field is retained. The artifact manifest binds each immutable result file and recorded checkpoint basename to a checkpoint hash verified separately. Member and ensemble shards span different recorded GPU families and predate explicit matrix-precision fields, so the arithmetic gaps are descriptive rather than a matched-hardware ensemble ablation.

3.3 What this result does and does not establish

The result shows that the rational architecture supports high closed-loop specialist capability. It is an existence claim for a compact tensor-decomposable policy, not a claim of broad robotic generalization or state of the art.

4 Exact weight-derived structure through attention

4.1 Exact layerwise construction

Equation 2 contains two query-key products and one value stream. Expanding the linear maps gives a coefficient table whose entries measure how combinations of query, key, and value coordinates

128 contribute to an output. No examples or fitted probe are needed to construct it. On a tiny four-
 129 token, width-four, one-head layer, direct evaluation and an independent explicit polynomial agree to
 130 2.00×10^{-15} . The internal core identity agrees to 1.24×10^{-14} .

131 The deployed attention also normalizes each query and key with Equation 3. Each normalization
 132 multiplies all channels of one token by the same scalar rational factor. Those factors can be pulled
 133 outside the bilinear dot products. The polynomial attention core is unchanged, while explicit
 134 numerator and denominator factors carry the normalization. The rational extension agrees to $3.61 \times$
 135 10^{-16} . On one fixed corrected cached input per checkpoint, we numerically replay all four vision
 136 and eight joint attention modules in each of three records that name the same ViT checkpoint artifacts
 137 as the capability results. This covers 36 module-input cases and 384 architectural heads. The worst
 138 relative ℓ_2 error is 2.016119×10^{-7} and the worst absolute difference is 2.44×10^{-5} . This residual
 139 comes from the deployed implementation’s intentional float32 statistic, while the reconstruction itself
 140 uses the learned coefficients. The algebraic identity is input-general. The one-input-per-checkpoint
 141 records provide a layerwise numerical replay audit, not an exact end-to-end policy decomposition.

142 We separately certify three convolutional checkpoints from the same rational architecture family.
 143 Their vision stacks contain no attention, so we independently rebuild all eight joint modules and all
 144 twelve heads per module from the saved weights. Sixteen deterministically selected cache inputs
 145 per checkpoint yield 4,608 head-input cases. Every output is finite and the worst relative ℓ_2 error is
 146 6.36×10^{-16} , below the frozen 10^{-6} gate. On the same inputs, independent CP-factor evaluation
 147 covers all 384 joint-FFN module-input cases at worst relative error 1.3871×10^{-15} against a 10^{-10}
 148 gate. Independent complete-block equations cover all 384 block-input cases at worst relative error
 149 1.5024×10^{-7} against a 10^{-4} gate. The complete-block comparison crosses a deployed PyTorch
 150 float32 block call and a NumPy float64 reconstruction. It is input-conditioned and per-block, not a
 151 sequential end-to-end reconstruction of the joint transformer or policy.

152 A separately frozen ViT audit partitions joint attention into vision, instruction, robot-state, and
 153 action-query source tokens. Across 128 deterministic, official-task-balanced inputs for each of the
 154 three capability checkpoints, all 36,864 head-input cases and 147,456 source-group evaluations are
 155 finite, with worst reconstruction error 5.0741×10^{-7} against a 10^{-6} gate. Across the 3,072 projected
 156 module-input cases, mean coherent-energy shares are 67.7% vision, 19.2% robot state, 7.9% action
 157 query, and 5.1% instruction. These input-conditioned magnitudes are descriptive rather than causal
 158 importance. Familiar-prompt permutations are also descriptive and do not establish grounding or
 159 closed-loop success.

160 Naively materializing the real width-384 coefficient tensor would require approximately 3.2×10^{15}
 161 values. We instead form an $O(d^2)$ reduced Gram. Like a covariance matrix, its leading eigenvectors
 162 rank important directions without constructing the full table. It matches brute-force materialization to
 163 2.13×10^{-14} at toy scale.

164 4.2 Trained policy result

165 The reconstruction certificate above spans three ViT checkpoints. The following subspace-ranking
 166 experiment is a descriptive analysis of the seed-0 ViT checkpoint. We apply the exact weight-derived
 167 calculation to all twelve heads in early, middle-late, and final backbone blocks 0, 6, and 7. Random
 168 projector banks are sampled once, then fixed and reused for all examples. The weights choose the
 169 subspace, while cached activations are used only afterward to score how well it preserves the layer’s
 170 computation.




Figure 3: Exact weight-derived attention structure. Ratios above one favor the exact weight-derived subspace over a random equal-rank subspace. Each point in the right panel is one trained attention head.

At ranks 8, 32, and 128, respectively, 80.6%, 91.7%, and 100% of the 36 block-head pairs favor the exact subspace. Their median ratios are 1.028, 1.171, and 2.441. At rank 128, the mean ratio is 3.033 and the strongest head reaches 10.135. The rank-128 block means are 3.08, 3.45, and 2.56 for blocks 0, 6, and 7. The 36 block-head pairs are architectural units within one checkpoint, not independent model-level replicates.

This extends exact weight-only construction through individual attention layers. It is not an exact end-to-end decomposition of the complete policy. Composition across all eight blocks still faces rapid degree and representation growth.

4.3 What the artifact lets a reviewer test

The structural artifact separates three questions that can otherwise be conflated. First, does a factorized operator implement its claimed equation? The attention and FFN records answer this by comparing independent evaluations of the saved factors. Second, does the claimed equation survive the numerical boundary of the deployed implementation? The complete-block records compare a NumPy float64 reconstruction with a deployed PyTorch float32 block call. Third, does this agreement reach the action interface? The RationalNorm record recomputes normalized and physical actions from stored arrays. Passing one level does not imply the next, so the verifier reports each decision separately.

This ladder also clarifies what a failed control would mean. Failure of an operator identity would invalidate the corresponding convertibility claim before any behavioral interpretation. Agreement at the operator and block levels would instead localize a later discrepancy to composition, preprocessing, decoding, or simulator interaction. The current complete-block certificate is not a sequential policy replay, so it cannot by itself rule out those later sources. This distinction is especially important for multimodal policies, where an exact internal equation can coexist with an out-of-graph mode choice or a continuous-action convention that changes behavior.

The modality ledger serves a different role. It assigns every joint-attention term to vision, instruction, robot-state, or action-query source tokens and verifies that the four reconstructed groups sum to the deployed output on fixed inputs. The reported energy shares expose where the trained computation is concentrated, but they are not causal importance scores. Prompt permutations likewise provide a reproducible sensitivity distribution, not evidence of grounding. A reviewer can therefore reproduce the algebraic statement without inheriting a stronger semantic interpretation.

Every record is tied to frozen checkpoint, cache, task, provenance, runtime, and source identities. Inputs are selected by deterministic hashes rather than by observed residuals. The standard-library verifier then recomputes row counts, extrema, distributions, and gate decisions from the released rows. These controls make the artifact useful beyond this model. A future conversion method can reuse the same ladder and state exactly which boundary its certificate crosses.

5 Related work and limitations

Related work. OpenVLA and OpenVLA-OFT establish pretrained VLA and continuous-control baselines [1, 2]. Bilinear MLPs, χ -nets, and bilinear IMPALA policies show that multiplicative weights can expose low-rank structure and support subsequent controls [3–5]. Recent work also steers pretrained VLAs through activation directions [6]. Our distinction is the exact weight-derived construction in a policy designed as a rational tensor program. Tensor and polynomial networks provide the broader multilinear foundation [12–14]. VLA robustness benchmarks show that standard LIBERO success can coexist with weak language use [9–11].

Limitations and conclusion. The benchmark is one specialist suite, so it does not establish broad robotic generalization. Exact reconstruction is certified layer by layer and does not yet produce one compact symbolic contraction of the full policy. The ViT records connect capability, exact attention, and the reduced-Gram screen through common checkpoint basenames. Separate convolutional files replicate the joint-attention equation certificate across another visual encoder. Evaluation uses fixed task instructions and does not establish robust instruction grounding. Within that scope, rational restrictions support strong LIBERO-Object control and exact trained-weight certificates. We do not claim a direct coefficient-edit interface.

Reproducibility. The anonymous artifact releases raw structural-certificate JSONs and all 2,000 episode rows behind the ViT member and ensemble results. Immutable SHA-256 manifests and standard-library verifiers independently recompute the capability totals, visual-bottleneck counts and intervals, and the attention, FFN, complete-block, modality-partition, and RationalNorm certificate decisions. Historical Modal cache-backed analyses read dataset-native task metadata. The replacement Athena pipeline pins that metadata and maps dataset identifiers to official identifiers before defining task splits. A source-level audit compared all 271,996 cached frames across Object, Spatial, Goal, and LIBERO-10 with their pinned LeRobot revisions. Task identifiers, language joins, states, actions, and resized images matched exactly for every frame. Canonical-state checks fail loudly when requested states are unavailable.

References

- [1] M. J. Kim et al. *OpenVLA: An Open-Source Vision-Language-Action Model*. CoRL, 2024.
- [2] M. J. Kim, C. Finn, and P. Liang. *Fine-Tuning Vision-Language-Action Models: Optimizing Speed and Success*. 2025.
- [3] A. Rigg, T. Dooms, M. T. Pearce, J. Oramas, and L. Sharkey. *Bilinear MLPs Enable Weight-Based Mechanistic Interpretability*. ICLR, 2025.
- [4] T. Dooms, W. Gauderis, G. A. Wiggins, and J. Oramas. *Compositionality Unlocks Deep Interpretable Models*. 2025.
- [5] F. Oozeer, Y. Erisken, A. Rigg, J. Oramas, and L. Sharkey. *Bilinear Convolution Decomposition for Causal RL Interpretability*. 2024.
- [6] B. Häon, K. Stocking, I. Chuang, and C. Tomlin. *Mechanistic Interpretability for Steering Vision-Language-Action Models*. 2025.
- [7] S. Marks et al. *Sparse Feature Circuits: Discovering and Editing Interpretable Causal Graphs in Language Models*. ICLR, 2025.
- [8] M. Kekić et al. *Learning Nonlinear Causal Reductions to Explain Reinforcement Learning Policies*. ICLR, 2026.
- [9] X. Zhou et al. *LIBERO-PRO: Towards Robust and Fair Evaluation of Vision-Language-Action Models Beyond Memorization*. 2025.
- [10] Y. Fang et al. *When Vision Overrides Language: Evaluating and Mitigating Counterfactual Failures in VLAs*. 2026.
- [11] C. Kim et al. *LIBERO-Para: A Diagnostic Benchmark and Metrics for Paraphrase Robustness in VLA Models*. 2026.
- [12] T. Kolda and B. Bader. *Tensor Decompositions and Applications*. SIAM Review, 2009.
- [13] E. Stoudenmire and D. Schwab. *Supervised Learning with Tensor Networks*. NeurIPS, 2016.
- [14] G. Chrysos et al. *Deep Polynomial Neural Networks*. CVPR and TPAMI, 2020 to 2021.

A Supporting results and decisions

A.1 Expanded attention screen

Within the same seed-0 ViT checkpoint, the corrected fixed-control attention analysis covers all eight blocks and twelve heads. At rank 128, 83/96 block-head pairs favor the exact weight-derived subspace. The median ratio is 1.94, the mean is 2.58, and blocks 2 and 3 provide useful depth controls. At ranks 8 and 32, only 54/96 and 60/96 pairs favor the exact subspace. The effect is therefore depth-dependent and strongest at the wider tested rank.

A.2 Prospectively fixed-rank visual bottleneck

The behavioral intervention acts at the width-384 visual-to-joint interface and keeps rank 96. Projector construction uses official Object tasks 0–3. Rank 96 was fixed before any corrected-basis outcome was observed, and the frozen evaluation uses tasks 8–9. Across three checkpoints and 300 canonical episodes, the full policy succeeds on 272 trials. The action-Jacobian projector retains $259/272 = 95.2\%$ of those successes, with a paired checkpoint-task-episode resampling interval of 91.2–98.6%. The activation-energy projector retains $267/272 = 98.2\%$, with an interval of 95.8–100.0%. Three fixed equal-rank random projectors retain $37/272 = 13.6\%$, $9/272 = 3.3\%$, and $22/272 = 8.1\%$.




Figure 4: **Corrected closed-loop visual bottleneck.** Both learned rank-96 projectors preserve nearly all baseline successes across three checkpoints, while equal-rank random projectors do not. Error bars in panel (a) are deterministic paired cluster-resampling intervals. Tasks 8–9 are absent from corrected projector construction, but not from policy training. This is not a clean rank-selection-holdout result.

271 The action-Jacobian basis does not outperform the matched activation-energy basis. The stronger
 272 action-specificity gate therefore does not pass. The supported claim is structured visual sufficiency
 273 relative to random projectors. This intervention is data-derived and remains separate from the exact
 274 weight-only attention decomposition.

275 A.3 Numerical scope of rational conversion

276 The deployed degree-[2/2] least-squares rational fit is exact as a rational computation even though it
 277 approximates RMS normalization. Its numerator has coefficients (5.32997, 7.53552, 0.327274) and
 278 the denominator has coefficients (1, 9.64827, 2.60664). Every denominator coefficient is positive,
 279 which certifies $Q(v) > 0$ for all $v \geq 0$. We instrument all 53 RationalNorm module instances,
 280 each invoked once per forward pass, on 4,096 deterministic inputs for each of three convolutional
 281 checkpoints. Across 544,542,720 activation rows, the minimum observed denominator is 1.0016
 282 and no row is nonfinite. Comparing deployed float32 normalization with an otherwise identical pass
 283 that recomputes only RationalNorm scales in float64, the largest checkpoint-level primary metric,
 284 defined as the maximum of normalized- and physical-action NRMSE, is 2.414×10^{-6} against a
 285 frozen 10^{-3} gate. This primary certificate concerns pole-free denominators and floating-point fidelity
 286 to the deployed rational operator. It does not claim uniform agreement with exact RMS scaling or
 287 measure matched alternative-normalizer runtime overhead.

288 The joint reconstruction ladder therefore extends from individual attention heads through FFNs to
 289 complete joint blocks on the same fixed inputs. The full-block certificate remains per-block and does
 290 not sequentially reconstruct the full transformer. The ViT source-partition audit is likewise layerwise
 291 and input-conditioned. Figure 5 reports the frozen numerical gates and the descriptive depth profile
 292 without assigning causal importance.




Figure 5: **Frozen reconstruction ladder and descriptive ViT source partition.** Panel (a) reports observed-to-gate ratios for independent Conv attention, joint-FFN, and complete joint-block reconstruction, plus the ViT source-partition identity and primary RationalNorm action fidelity. All plotted primary numerical gates pass. Panel (b) shows descriptive mean coherent-energy shares after the attention output projection across 384 checkpoint-input cases per joint block. The depth trend is not a causal importance or grounding result. Complete-block replay is per-block rather than sequential end-to-end policy reconstruction.