
A Robot Policy Designed as a Weight-Space Artifact Exact Tensor Structure in a Capable VLA

Rational conversion, reduced-Gram analysis, and causal intervention

Anonymous Author(s)

Affiliation

Address

email

Abstract

Tensor decomposition rewrites a multidimensional table as a small set of low-rank factors. In a neural network, those factors can expose which combinations of inputs and outputs are encoded by the weights. Conventional nonlinear operations make such a table difficult to obtain directly. We introduce χ -VLA, a 20.1M-parameter vision-language-action policy containing only bilinear, linear, and rational learned operations. Its checkpoint therefore specifies an explicit rational tensor program rather than only an opaque executable function. A mechanical audit and local reconstructions verify that claim. We derive an exact weight-based decomposition of each softmax-free attention layer and reconstruct all 12 modules in three checkpoints with worst relative error 2.56×10^{-7} . An $O(d^2)$ reduced Gram avoids materializing the coefficient tensor. All 36 primary block-head tests favor the rank-128 weight-derived subspace over fixed random controls. The full appendix screen finds the effect in 83/96 pairs and identifies where the signal concentrates across depth. Across three checkpoints, a causally ranked 96-dimensional visual subspace succeeds in 50/60 matched trials, compared with 52/60 for the full representation. Three equal-rank random controls reach only 1/60, 3/60, and 3/60. The selected subspace retains 49 of the 52 full-policy successes. On disjoint cache samples, it has $33.5\text{--}110.6\times$ lower reconstruction error than the median random projector across action groups and checkpoints. The artifact is behaviorally nontrivial, reaching $93.7\% \pm 0.8\%$ on LIBERO-Object, and a three-checkpoint ensemble reaches 96.2%. We therefore establish a capable model artifact with exact layerwise structure and a compact causal policy bottleneck.

1 Introduction

Neural checkpoints are usually treated as opaque containers for functions. We can run them and collect their activations, but the weights alone rarely provide a readable map from input features to outputs. A post-hoc probe may reveal that an activation correlates with behavior, yet it does not show how the weights construct that feature or whether editing it will have a selective causal effect.

A tensor is a multidimensional table, and tensor decomposition is the analogue of low-rank matrix factorization for such tables. When a network uses only additions, multiplications, and ratios of polynomials, its weights can be expanded into coefficient tensors whose entries record input interactions. Decomposition ranks simpler factors inside those coefficients. The intended benefit is an inspectable model artifact: candidate mechanisms come from the trained parameters, then activation data and interventions test whether they matter. Softmax, ordinary activations, and square-root normalization obstruct this direct conversion in conventional networks. Bilinear language and vision

models suggest that changing the primitives can recover it [3, 4]. Whether the resulting artifact can still control a robot is unresolved.

We study two linked questions:

Can the learned weights of a specialist VLA specify an explicit rational tensor program, and can that artifact expose exact structure with a measurable causal consequence?

Our contributions are:

1. We construct and mechanically certify a complete VLA from tensor-compatible primitives. A rational normalization supplies input-specific scaling without leaving the rational function class.
2. We derive exact weight-based structure through each softmax-free attention layer and an $O(d^2)$ reduced-Gram calculation that avoids materializing its high-order coefficient tensor.
3. We test the artifact at behavioral scale. Exact subspaces beat fixed random controls, and a causally ranked rank-96 visual subspace retains 49 of 52 full-policy successes across three checkpoints. Three random controls total only 7/180.
4. We establish that the analyzed artifact is a capable policy, reaching $93.7\% \pm 0.8\%$ over three LIBERO-Object seeds. A three-checkpoint ensemble reaches 96.2% at three times the inference cost.

This paper does not claim current state of the art. It also does not claim that comparable mechanisms are impossible to obtain from conventional policies. The empirical question is whether the algebraic policy makes weight-derived structure direct, stable, and behaviorally testable.

2 A rational tensor robot policy

2.1 Tensor-compatible primitives

Each primitive replaces a familiar transformer component while keeping explicit coefficients. Linear maps contribute first-order terms, while multiplying learned projections creates higher-order interactions that remain determined by the weights.

Homogeneous coordinate. We write $\bar{x} = [1; x]$. The constant coordinate allows a multiplicative layer to represent bias, linear, and quadratic terms in one contraction.

Bilinear feed-forward block. For $L, R \in \mathbb{R}^{r \times (d+1)}$ and $D \in \mathbb{R}^{d \times r}$,

$$\text{BFFN}(x) = D[(L\bar{x}) \odot (R\bar{x})], \quad y_o = \sum_{i,j} \left(\sum_{\rho} D_{o\rho} L_{\rho i} R_{\rho j} \right) \bar{x}_i \bar{x}_j. \quad (1)$$

The inner sum is a table indexed by one output and two input coordinates. The three matrices give its CP factorization, the tensor analogue of a low-rank matrix factorization.

Softmax-free bilinear attention. Each head forms two query-key systems and one value stream:

$$Y = C [M \odot (Q_1 K_1^\top) \odot (Q_2 K_2^\top) / s] V. \quad (2)$$

Here M is a fixed causal mask, C is fixed row scaling, and s is fixed. All learned query, key, value, and output maps are linear. Equation 2 is therefore a polynomial in its input activations, with interaction coefficients determined by the weights.

Rational per-instance normalization. RMS normalization contains an input-dependent square root. We instead deploy a fixed rational approximation $r(z) = P(z)/Q(z)$ to $1/\sqrt{z + \epsilon}$:

$$\text{RNorm}(x) = x r \left(\frac{1}{d} \sum_{j=1}^d x_j^2 \right). \quad (3)$$

Here rational means a ratio of polynomials. Projective coordinates carry each activation as a numerator-denominator pair (N, D) representing N/D . Updating this pair preserves input-specific rescaling without leaving the rational tensor representation.

2.2 Architecture and training

The policy follows the high-level VLA sequence of visual encoding, language and state embedding, joint processing, and action decoding [1]. It uses a 64×64 agent-view image, an eight-dimensional robot state, a 26-token task vocabulary, and eight action-query tokens. The joint backbone has width 384, eight blocks, and twelve heads. The action head maps each action-query state directly to seven continuous action dimensions. The complete model has 20,137,352 learned parameters.

Training uses 66,984 demonstration frames, 40,000 AdamW updates, batch size 256, cosine decay, bfloat16 arithmetic, and an exponential moving average. Closed-loop evaluation uses a 180° camera rotation to match the demonstration convention, canonical LIBERO initial states, and fail-loud checks that verify every requested state was loaded.

2.3 Mechanical certificate on the trained checkpoint

Architecture definitions can differ from deployed code. We therefore audit the real 189/200 checkpoint at runtime and reconstruct representative operations from saved weights.




Figure 1: **Mechanical audit of the trained rational checkpoint.** The runtime path invokes 112 linear maps, 74 rational normalizers, two embeddings, and one linear patch projection. No softmax, standard activation, or ordinary normalization is present. Reconstruction errors are measured over 54,784 real activation rows.

The normalization reconstruction has maximum absolute error 1.64×10^{-6} . A bilinear FFN reconstruction has median relative error 1.40×10^{-7} . These residuals are bounded by the deployed module’s float32 statistic. The audit certifies operator membership and local reconstruction. It does not materialize one compact polynomial for the full eight-layer transformer.

3 The artifact is a capable closed-loop policy

3.1 Evaluation protocol

We evaluate all ten LIBERO-Object tasks with:

- 50 canonical trials per task
- a 280-step cap
- three independently trained seeds per architecture
- 500 rollouts per seed and 1,500 rollouts per architecture

Published OpenVLA and Diffusion Policy results are protocol-aligned references. They were not rerun in our codebase, so differences may still include implementation and training choices.

3.2 Results

Method	Parameters	Training seeds	Success
χ -VLA, rational conv	20.1M	3	93.7% $\pm$ 0.8%
χ -VLA, conv ensemble	$3 \times 20.1\text{M}$	$3 \rightarrow 1$	96.2%
OpenVLA-7B [1]	7B	3	88.4% $\pm$ 0.8%
Diffusion Policy [1]	not matched	3	92.5% $\pm$ 0.7%

Table 1: **Closed-loop success under aligned evaluation settings.** The two external rows are historical references rather than current state-of-the-art claims.

The convolutional seeds score 94.0%, 94.4%, and 92.8%. Their narrow spread shows that high specialist capability is reproduced across three independently trained checkpoints.

Elementwise averaging of the three convolutional action predictions reaches 96.2%. This is 2.3 points above the mean of its members and 1.2 points above the strongest member. The gain is largest on tomato sauce, butter, and ketchup. The ensemble requires three complete forward passes and is not one 20.1M-parameter policy.

3.3 What this result does and does not establish

The result shows that the rational architecture supports high closed-loop specialist capability. It is an existence and reproducibility claim for a compact tensor-decomposable policy, not a claim of broad robotic generalization or state of the art.

4 Exact weight-derived structure through attention

4.1 Exact layerwise construction

Equation 2 contains two query-key products and one value stream. Expanding the linear maps gives a coefficient table whose entries measure how combinations of query, key, and value coordinates contribute to an output. No examples or fitted probe are needed to construct it. On a tiny four-token, width-four, one-head layer, direct evaluation and an independent explicit polynomial agree to 2.00×10^{-15} . The internal core identity agrees to 1.24×10^{-14} .

The deployed attention also normalizes each query and key with Equation 3. Each normalization multiplies all channels of one token by the same scalar rational factor. Those factors can be pulled outside the bilinear dot products. The polynomial attention core is unchanged, while explicit numerator and denominator factors carry the normalization. The rational extension agrees to 3.61×10^{-16} . We then reconstruct all four vision and eight joint attention modules in each of three trained checkpoints, covering 384 heads in total. The worst relative ℓ_2 error is 2.56×10^{-7} and the worst absolute difference is 2.23×10^{-5} . This residual comes from the deployed implementation’s intentional float32 statistic, while the reconstruction itself uses the learned coefficients.

Naively materializing the real width-384 coefficient tensor would require approximately 3.2×10^{15} values. We instead form an $O(d^2)$ reduced Gram. Like a covariance matrix, its leading eigenvectors rank important directions without constructing the full table. It matches brute-force materialization to 2.13×10^{-14} at toy scale.

4.2 Trained policy result

For the subspace-ranking experiment, we apply the exact weight-derived calculation to all twelve heads in early, middle-late, and final backbone blocks 0, 6, and 7. Random projector banks are sampled once, then fixed and reused for all examples. The weights choose the subspace, while cached activations are used only afterward to score how well it preserves the layer’s computation.




Figure 2: Exact weight-derived attention structure. Ratios above one favor the exact weight-derived subspace over a random equal-rank subspace. Each point in the right panel is one trained attention head.

At ranks 8, 32, and 128, respectively, 80.6%, 91.7%, and 100% of the 36 block-head pairs favor the exact subspace. Their median ratios are 1.028, 1.171, and 2.441. At rank 128, the mean ratio is 3.033 and the strongest head reaches 10.135. The rank-128 block means are 3.08, 3.45, and 2.56 for blocks 0, 6, and 7.

This extends exact weight-only construction through individual attention layers. It is not an exact end-to-end decomposition of the complete policy. Composition across all eight blocks still faces rapid degree and representation growth.

5 Causal policy structure

We next test whether the identified structure affects closed-loop behavior.

5.1 Causal visual subspace

At the visual bond before the joint backbone, we rank directions separately for translation, rotation, and gripper actions. A subspace is a smaller coordinate system inside the model’s 384-dimensional visual representation. Retaining it means zeroing all other directions, then running the actual policy. We compare selected and random equal-width subspaces. Because the ranking uses cached training demonstrations and downstream action gradients, offline reconstruction remains in-cache. Its evaluation samples are disjoint from the samples that estimate the ranking. Closed-loop intervention uses independent canonical simulator states, so it tests a causal bottleneck but is not a weight-only discovery result.

The action-balanced rank screen identifies a 96-dimensional subspace, one quarter of the visual bond. Full representations succeed in 17/20, 15/20, and 20/20 trials across checkpoints. Selected subspaces reach 17/20, 14/20, and 19/20, and retain 49 of the 52 exact full-policy successes. Their conditional retention is 94.1%, 93.3%, and 95.0%. Three random controls per checkpoint total only 1/60, 3/60, and 3/60. Every checkpoint and random-control comparison has the same selected-over-random direction, with exact paired $p \leq 2.4 \times 10^{-4}$. Using 1,024 discovery samples and 1,024 disjoint evaluation samples per checkpoint, the selected basis has 53.0–110.6 $\times$ lower translation error, 37.5–62.1 $\times$ lower rotation error, and 33.5–103.0 $\times$ lower gripper error than the median equal-rank random projector. This repairs sample reuse but does not turn the training-cache diagnostic into held-out task generalization. This data-driven full-policy ranking complements the exact layerwise analysis in Section 4.

6 Related work and limitations

Related work. OpenVLA and OpenVLA-OFT establish pretrained VLA and continuous-control baselines [1, 2]. Bilinear MLPs, χ -nets, and bilinear IMPALA policies show that multiplicative weights can expose low-rank structure and causal controls [3–5]. Recent work also steers pretrained

169 VLAs through activation directions [6]. Our distinction is the exact weight-derived construction in a
170 policy designed as a rational tensor program. Tensor and polynomial networks provide the broader
171 multilinear foundation [12–14]. VLA robustness benchmarks show that standard LIBERO success
172 can coexist with weak language use [9–11].

173 **Limitations and conclusion.** The benchmark is one specialist suite, so it does not establish broad
174 robotic generalization. Exact reconstruction is certified layer by layer, while the full-policy visual
175 ranking also uses cached activations and action gradients. Within that scope, rational restrictions
176 support competitive LIBERO-Object control, exact trained-weight certificates, and a rank-96 causal
177 bottleneck that retains 49 of 52 full-policy successes across three checkpoints.

178 **Reproducibility.** The final artifact will provide code, configurations, checkpoint hashes, exact
179 certificates, and per-episode logs. Canonical-state checks fail loudly when requested states are
180 unavailable.

References

- [1] M. J. Kim et al. *OpenVLA: An Open-Source Vision-Language-Action Model*. CoRL, 2024.
- [2] M. J. Kim, C. Finn, and P. Liang. *Fine-Tuning Vision-Language-Action Models: Optimizing Speed and Success*. 2025.
- [3] A. Rigg, T. Doods, M. T. Pearce, J. Oramas, and L. Sharkey. *Bilinear MLPs Enable Weight-Based Mechanistic Interpretability*. ICLR, 2025.
- [4] T. Doods, W. Gauderis, G. A. Wiggins, and J. Oramas. *Compositionality Unlocks Deep Interpretable Models*. 2025.
- [5] F. Oozeer, Y. Eriskin, A. Rigg, J. Oramas, and L. Sharkey. *Bilinear Convolution Decomposition for Causal RL Interpretability*. 2024.
- [6] B. Häon, K. Stocking, I. Chuang, and C. Tomlin. *Mechanistic Interpretability for Steering Vision-Language-Action Models*. 2025.
- [7] S. Marks et al. *Sparse Feature Circuits: Discovering and Editing Interpretable Causal Graphs in Language Models*. ICLR, 2025.
- [8] M. Kekić et al. *Learning Nonlinear Causal Reductions to Explain Reinforcement Learning Policies*. ICLR, 2026.
- [9] X. Zhou et al. *LIBERO-PRO: Towards Robust and Fair Evaluation of Vision-Language-Action Models Beyond Memorization*. 2025.
- [10] Y. Fang et al. *When Vision Overrides Language: Evaluating and Mitigating Counterfactual Failures in VLAs*. 2026.
- [11] C. Kim et al. *LIBERO-Para: A Diagnostic Benchmark and Metrics for Paraphrase Robustness in VLA Models*. 2026.
- [12] T. Kolda and B. Bader. *Tensor Decompositions and Applications*. SIAM Review, 2009.
- [13] E. Stoudenmire and D. Schwab. *Supervised Learning with Tensor Networks*. NeurIPS, 2016.
- [14] G. Chrysos et al. *Deep Polynomial Neural Networks*. CVPR and TPAMI, 2020 to 2021.

A Supporting results and decisions

A.1 Original 94.5% capability result

Before the final protocol-aligned benchmark, one rational ViT checkpoint succeeded in 189/200 trials under a 600-step horizon and 20 trials per task. This result established that the rational implementation could support all ten tasks. It should not be compared directly with the protocol-aligned external references because it uses one training seed, fewer trials, and a longer horizon.




Figure 3: **One quarter of the visual bond preserves policy behavior.** Each checkpoint uses four tasks and five paired canonical trials per task. Rank 96 retains 49 of 52 full-policy successes. Three equal-rank random controls reach 1/60, 3/60, and 3/60. The pooled interval is descriptive because trials within a checkpoint share learned weights.

212 A.2 Per-task ensemble behavior




Figure 4: **Per-task effect of prediction averaging.** Gray points are the three convolutional checkpoints, evaluated with 50 canonical trials per task. Curves compare their taskwise mean with elementwise action-prediction averaging. The ensemble raises overall success from a 93.9% member mean to 96.2%, with its largest gains on tomato sauce, butter, and ketchup. It requires three forward passes and is not a single 20.1M-parameter policy.

213 A.3 Expanded attention screen

214 The corrected fixed-control attention analysis covers all eight blocks and twelve heads. At rank 128,
 215 83/96 block-head pairs favor the exact weight-derived subspace. The median ratio is 1.94, the mean
 216 is 2.58, and blocks 2 and 3 provide useful depth controls. At ranks 8 and 32, only 54/96 and 60/96
 217 pairs favor the exact subspace. The effect is therefore depth-dependent and strongest at the wider
 218 tested rank.

219 **A.4 Numerical scope of rational conversion**

220 The deployed Padé computation is exactly rational even though it approximates RMS normalization.
221 Final certification must report its polynomial degrees, coefficients, denominator margin, observed
222 activation range, tail error, action drift, and runtime overhead. Operator membership alone does not
223 establish practical numerical safety under distribution shift.