A Robot Policy Designed as a Weight-Space Artifact
Exact Tensor Structure in a Capable VLA

Rational conversion, reduced-Gram analysis, and exact weight access

Anonymous Author(s)
Affiliation
Address
email

Abstract

Tensor decomposition rewrites a multidimensional table as a small set of low-rank factors. In a neural network, those factors can expose which combinations of inputs and outputs are encoded by the weights. Conventional nonlinear operations make such a table difficult to obtain directly. We introduce χ-VLA, a rational vision-language-action model family with a 19.2M convolutional policy and a 20.1M ViT policy, each containing only bilinear, linear, and rational learned operations. Its checkpoint therefore specifies an explicit rational tensor program rather than only an opaque executable function. A runtime operator audit of the seed-0 20.1M ViT checkpoint and local reconstructions verify its operator-level form. We derive an exact weight-based decomposition of each softmax-free attention layer and, on one fixed cached input per checkpoint, numerically replay all 12 modules in each of three ViT checkpoint artifacts with worst relative error 2.56×10^{-7}. An $O(d^2)$ reduced Gram avoids materializing the coefficient tensor. In a seed-0 ViT checkpoint, all 36 primary block-head tests favor the rank-128 weight-derived subspace over fixed random controls. Across the same three ViT checkpoints, closed-loop LIBERO-Object success is $85.3\% \pm 4.6\%$ over 500 canonical trials per checkpoint, and a fixed prediction-mean ensemble reaches $467/500 = 93.4\%$. A separate Conv audit reconstructs $4{,}608$ attention-head cases, while learned rank-96 visual projectors retain 95.2–98.2% of baseline successes, versus 3.3–13.6% for random projectors. We therefore establish a capable model family with exact layerwise structure derived from trained coefficients. We do not yet provide an exact end-to-end decomposition of the complete policy.

1 Introduction

Neural checkpoints are usually treated as opaque containers for functions. We can run them and collect their activations, but the weights alone rarely provide a readable map from input features to outputs. A post-hoc probe may reveal that an activation correlates with behavior, yet it does not show how the weights construct that feature.

A tensor is a multidimensional table, and tensor decomposition is the analogue of low-rank matrix factorization for such tables. When a network uses only additions, multiplications, and ratios of polynomials, its weights can be expanded into coefficient tensors whose entries record input interactions. Decomposition ranks simpler factors inside those coefficients. The intended benefit is an inspectable model artifact: candidate structure comes from the trained parameters. By fully tensor-decomposable, we mean that every learned block preserves this algebraic form. This is an operator-closure claim, not a claim that we materialize one compact end-to-end coefficient tensor. Its degree and representation size grow rapidly with depth. For example, one coefficient can connect an image coordinate, an instruction coordinate, and an action coordinate. Factoring the coefficient table ranks high-energy cross-modal interactions directly from the checkpoint without first fitting an activation probe. Those factors are candidate structure, not automatically human-readable concepts or causal explanations. Softmax, ordinary activations, and square-root normalization obstruct this direct conversion in conventional networks. Bilinear language and vision models suggest that changing the primitives can recover it [3, 4]. Whether the resulting artifact can still control a robot is unresolved.

We study two linked questions:

> *Can the learned weights of a specialist VLA specify an explicit rational tensor program, and can exact weight analysis coexist with strong closed-loop control?*

Our contributions are:

1. We construct a complete VLA from tensor-compatible primitives and runtime-audit its seed-0 20.1M ViT instantiation. A rational normalization supplies input-specific scaling without leaving the rational function class.
2. We derive exact weight-based structure through each softmax-free attention layer and an $O(d^2)$ reduced-Gram calculation that avoids materializing its high-order coefficient tensor.
3. We test exact attention subspaces against fixed equal-rank random controls. The weight-derived subspaces win all 36 primary block-head comparisons in a seed-0 checkpoint.
4. We evaluate the same three ViT checkpoints on $1{,}500$ canonical LIBERO-Object trials. They reach $85.3\% \pm 4.6\%$, while a fixed prediction-mean ensemble reaches $467/500 = 93.4\%$.

Artifact query	Returned object	Current guarantee
Runtime graph census	Operator inventory and module identities	Membership for the audited seed-0 20.1M ViT checkpoint
Attention-coefficient query	Reduced Gram and leading directions	Exact layerwise construction from saved weights
Numerical replay	Per-case residuals and frozen gates	Independent implementation check on recorded inputs
Artifact integrity	Result and source digests	Immutable SHA-256 manifest and one-command verification

Table 1: **Checkpoint interface as a data artifact.** The trained checkpoint supports structured queries beyond ordinary forward execution. Guarantees remain layerwise and scoped to the named artifacts.

This paper does not claim current state of the art. It also does not claim that comparable mechanisms are impossible to obtain from conventional policies. The empirical question is whether the algebraic policy makes exact weight structure directly accessible.

2 A rational tensor robot policy

2.1 Tensor-compatible primitives

Each primitive replaces a familiar transformer component while keeping explicit coefficients. Linear maps contribute first-order terms, while multiplying learned projections creates higher-order interactions that remain determined by the weights.

Homogeneous coordinate. We write $\bar{x} = [1; x]$. The constant coordinate allows a multiplicative layer to represent bias, linear, and quadratic terms in one contraction.

Bilinear feed-forward block. For $L, R \in \mathbb{R}^{r \times (d+1)}$ and $D \in \mathbb{R}^{d \times r}$,

$$\mathrm{BFFN}(x) = D\left[(L\bar{x}) \odot (R\bar{x})\right], \qquad y_o = \sum_{i,j} \left(\sum_{\rho} D_{o\rho} L_{\rho i} R_{\rho j} \right) \bar{x}_i \bar{x}_j. \tag{1}$$

The inner sum is a table indexed by one output and two input coordinates. The three matrices give its CP factorization, the tensor analogue of a low-rank matrix factorization.

Softmax-free bilinear attention. Each head forms two query-key systems and one value stream:

$$Y = C\left[M \odot (Q_1 K_1^\top) \odot (Q_2 K_2^\top)/s\right] V. \tag{2}$$

Here M is a fixed causal mask, C is fixed row scaling, and s is fixed. All learned query, key, value, and output maps are linear. Equation 2 is therefore a polynomial in its input activations, with interaction coefficients determined by the weights.

Rational per-instance normalization. RMS normalization contains an input-dependent square root. Let $m(x) = d^{-1}\sum_{j=1}^{d} x_j^2 + \epsilon$, and let $s_0 > 0$ be the frozen running mean-square scale. We deploy a fixed rational approximation $r(v) = P(v)/Q(v)$ to $v^{-1/2}$ over $v \in [0.1, 10]$:

$$v = \frac{m(x)}{s_0}, \qquad \mathrm{RNorm}(x) = \frac{x}{\sqrt{s_0}}\, r(v). \tag{3}$$

Here rational means a ratio of polynomials. Projective coordinates carry each activation as a numerator-denominator pair (N, D) representing N/D. Updating this pair preserves input-specific rescaling without leaving the rational tensor representation. Training and deployment use r itself. The exactness claims below therefore concern faithful evaluation of the deployed rational computation, not uniform agreement with RMSNorm. Appendix A.3 quantifies that distinction.

2.2 Architecture and training

Both encoder variants follow the high-level VLA sequence of visual encoding, language and state embedding, joint processing, and action decoding [1]. They use a 64×64 agent-view image, an eight-dimensional robot state, a 26-token task vocabulary, and eight action-query tokens. Both share a width-384, eight-block, twelve-head joint backbone and a linear head that maps each action-query state to seven continuous action dimensions. The 20.1M ViT uses four bilinear-attention vision blocks and has 20,137,352 learned parameters. It supplies the closed-loop capability, runtime operator census, exact-attention, and reduced-Gram results. The 19.2M convolutional variant instead uses three stride-2 bilinear convolution stages that produce an 8×8 token grid. Three separate convolutional files receive the joint-attention and rational-normalizer certificates reported in Appendix A.3.

Training uses 66,984 demonstration frames, 40,000 AdamW updates, batch size 256, cosine decay, bfloat16 arithmetic, and an exponential moving average. Closed-loop evaluation uses a 180° camera rotation to match the demonstration convention, canonical LIBERO initial states, and fail-loud checks that verify every requested state was loaded.

2.3 Mechanical certificate on the seed-0 20.1M ViT checkpoint

Architecture definitions can differ from deployed code. We therefore audit the deployed seed-0 20.1M ViT checkpoint at runtime and reconstruct representative operations from saved weights.

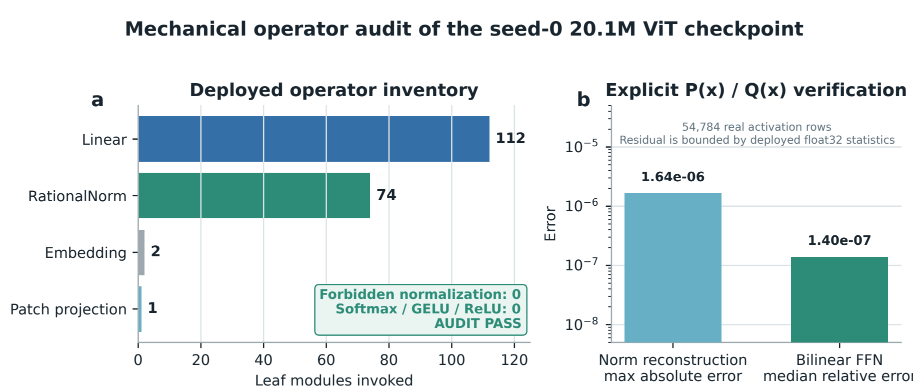

Figure 1: **Mechanical audit of the seed-0 20.1M ViT rational checkpoint.** The runtime path invokes 112 linear maps, 74 rational normalizers, two embeddings, and one linear patch projection. No softmax, standard activation, or ordinary normalization is present. Reconstruction errors are measured over 54,784 real activation rows.

The normalization reconstruction has maximum absolute error 1.64×10^{-6}. A bilinear FFN reconstruction has median relative error 1.40×10^{-7}. These residuals are bounded by the deployed module's float32 statistic. The audit certifies operator membership and local reconstruction. It does not materialize one compact polynomial for the full eight-layer transformer.

3 The artifact is a capable closed-loop policy

3.1 Evaluation protocol

We evaluate all ten LIBERO-Object tasks with the following fixed protocol:

- 50 canonical trials per task
- a 280-step cap
- three independently trained ViT checkpoints
- 500 rollouts per checkpoint and 1,500 rollouts in total

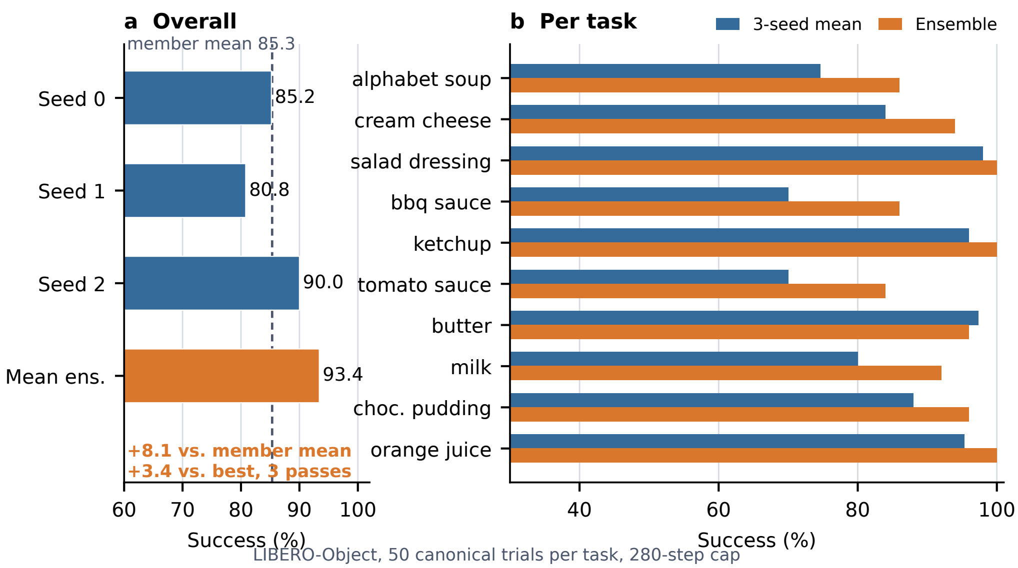

Figure 2: **Verified ViT capability.** Left: three independently trained checkpoints and their fixed elementwise prediction-mean ensemble. Right: per-task three-checkpoint mean and ensemble success. The ensemble uses three forward passes.

3.2 Results

Method	Parameters	Inference	Success
χ-VLA, rational ViT	20.1M	1 pass	$\mathbf{85.3\% \pm 4.6\%}$
χ-VLA, prediction-mean ensemble	3×20.1M	3 passes	$\mathbf{93.4\%}$

Table 2: **Closed-loop LIBERO-Object success.** Each ViT checkpoint receives 500 canonical trials. The single-model row is the mean $\pm$ sample standard deviation across three checkpoints. The ensemble is one fixed composite evaluated on the same 500 states.

The three checkpoints reach $426/500 = 85.2\%$, $404/500 = 80.8\%$, and $450/500 = 90.0\%$. Elementwise averaging of their predicted action chunks reaches $467/500 = 93.4\%$. This is 8.1 points above the member mean and 3.4 points above the strongest member. The ensemble requires three complete forward passes and is not one 20.1M-parameter policy. Every episode-level record, task index, and protocol field is retained. The artifact manifest binds each immutable result file and recorded checkpoint basename to a checkpoint hash verified separately. Member and ensemble shards span different recorded GPU families and predate explicit matrix-precision fields, so the arithmetic gaps are descriptive rather than a matched-hardware ensemble ablation.

3.3 What this result does and does not establish

The result shows that the rational architecture supports high closed-loop specialist capability. It is an existence claim for a compact tensor-decomposable policy, not a claim of broad robotic generalization or state of the art.

4 Exact weight-derived structure through attention

4.1 Exact layerwise construction

Equation 2 contains two query-key products and one value stream. Expanding the linear maps gives a coefficient table whose entries measure how combinations of query, key, and value coordinates contribute to an output. No examples or fitted probe are needed to construct it. On a tiny four-token, width-four, one-head layer, direct evaluation and an independent explicit polynomial agree to 2.00×10^{-15}. The internal core identity agrees to 1.24×10^{-14}.

The deployed attention also normalizes each query and key with Equation 3. Each normalization multiplies all channels of one token by the same scalar rational factor. Those factors can be pulled outside the bilinear dot products. The polynomial attention core is unchanged, while explicit numerator and denominator factors carry the normalization. The rational extension agrees to 3.61×10^{-16}. On one fixed cached input per checkpoint, we numerically replay all four vision and eight joint attention modules in each of three records that name the same ViT checkpoint artifacts as the capability results. This covers 36 module-input cases and 384 architectural heads. The worst relative ℓ_2 error is 2.56×10^{-7} and the worst absolute difference is 2.23×10^{-5}. This residual comes from the deployed implementation's intentional float32 statistic, while the reconstruction itself uses the learned coefficients. The algebraic identity is input-general. The fixed inputs provide a numerical replay audit.

We separately certify three convolutional checkpoints from the same rational architecture family. Their vision stacks contain no attention, so we independently rebuild all eight joint modules and all twelve heads per module from the saved weights. Sixteen deterministically selected cache inputs per checkpoint yield 384 module cases and 4,608 head-input cases. Every output is finite and the worst relative ℓ_2 error is 6.36×10^{-16}, far below the frozen 10^{-6} gate. This same-checkpoint certificate covers the joint-attention equation at each normalized layer input. It does not rebuild the convolutional stem, feed-forward blocks, residual composition, final normalization, or action head as one end-to-end contraction.

Naively materializing the real width-384 coefficient tensor would require approximately 3.2×10^{15} values. We instead form an $O(d^2)$ reduced Gram. Like a covariance matrix, its leading eigenvectors rank important directions without constructing the full table. It matches brute-force materialization to 2.13×10^{-14} at toy scale.

4.2 Trained policy result

The reconstruction certificate above spans three ViT checkpoints. The following subspace-ranking experiment is a descriptive analysis of the seed-0 ViT checkpoint. We apply the exact weight-derived calculation to all twelve heads in early, middle-late, and final backbone blocks 0, 6, and 7. Random projector banks are sampled once, then fixed and reused for all examples. The weights choose the subspace, while cached activations are used only afterward to score how well it preserves the layer's computation.

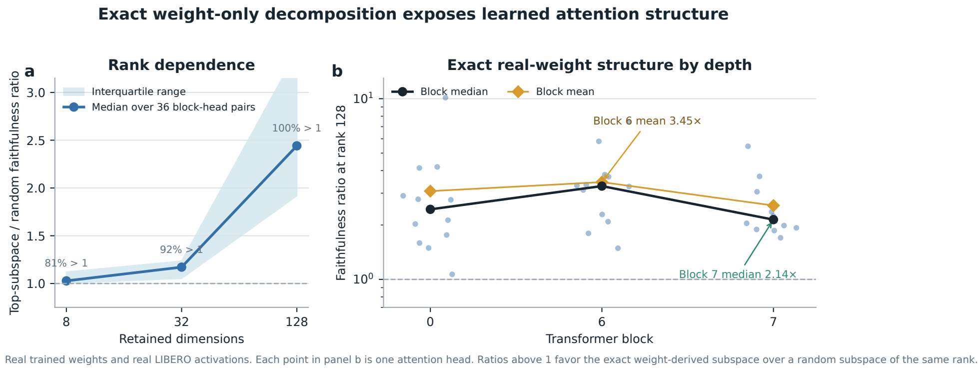

Figure 3: **Exact weight-derived attention structure.** Ratios above one favor the exact weight-derived subspace over a random equal-rank subspace. Each point in the right panel is one trained attention head.

At ranks 8, 32, and 128, respectively, 80.6%, 91.7%, and 100% of the 36 block-head pairs favor the exact subspace. Their median ratios are 1.028, 1.171, and 2.441. At rank 128, the mean ratio is 3.033 and the strongest head reaches 10.135. The rank-128 block means are 3.08, 3.45, and 2.56 for blocks 0, 6, and 7. The 36 block-head pairs are architectural units within one checkpoint, not independent model-level replicates.

This extends exact weight-only construction through individual attention layers. It is not an exact end-to-end decomposition of the complete policy. Composition across all eight blocks still faces rapid degree and representation growth.

5 Related work and limitations

Related work. OpenVLA and OpenVLA-OFT establish pretrained VLA and continuous-control baselines [1, 2]. Bilinear MLPs, χ-nets, and bilinear IMPALA policies show that multiplicative weights can expose low-rank structure and support subsequent controls [3–5]. Recent work also steers pretrained VLAs through activation directions [6]. Our distinction is the exact weight-derived construction in a policy designed as a rational tensor program. Tensor and polynomial networks provide the broader multilinear foundation [12–14]. VLA robustness benchmarks show that standard LIBERO success can coexist with weak language use [9–11].

Limitations and conclusion. The benchmark is one specialist suite, so it does not establish broad robotic generalization. Exact reconstruction is certified layer by layer and does not yet produce one compact symbolic contraction of the full policy. The ViT records connect capability, exact attention, and the reduced-Gram screen through common checkpoint basenames. Separate convolutional files replicate the joint-attention equation certificate across another visual encoder. Evaluation uses fixed task instructions and does not establish robust instruction grounding. Within that scope, rational restrictions support strong LIBERO-Object control and exact trained-weight certificates. We do not claim a direct coefficient-edit interface.

Reproducibility. The anonymous artifact releases raw structural-certificate JSONs and all 2,000 episode rows behind the ViT member and ensemble results. Immutable SHA-256 manifests and standard-library verifiers independently recompute the capability totals, visual-bottleneck counts and intervals, and the convolutional-attention and RationalNorm certificate decisions. Historical Modal cache-backed analyses read dataset-native task metadata. The replacement Athena pipeline pins that metadata and maps dataset identifiers to official identifiers before defining task splits. A source-level audit compared all 271,996 cached frames across Object, Spatial, Goal, and LIBERO-10 with their pinned LeRobot revisions. Task identifiers, language joins, states, actions, and resized images matched exactly for every frame. Canonical-state checks fail loudly when requested states are unavailable.

References

[1] M. J. Kim et al. *OpenVLA: An Open-Source Vision-Language-Action Model.* CoRL, 2024.

[2] M. J. Kim, C. Finn, and P. Liang. *Fine-Tuning Vision-Language-Action Models: Optimizing Speed and Success.* 2025.

[3] A. Rigg, T. Dooms, M. T. Pearce, J. Oramas, and L. Sharkey. *Bilinear MLPs Enable Weight-Based Mechanistic Interpretability.* ICLR, 2025.

[4] T. Dooms, W. Gauderis, G. A. Wiggins, and J. Oramas. *Compositionality Unlocks Deep Interpretable Models.* 2025.

[5] F. Oozeer, Y. Erisken, A. Rigg, J. Oramas, and L. Sharkey. *Bilinear Convolution Decomposition for Causal RL Interpretability.* 2024.

[6] B. Häon, K. Stocking, I. Chuang, and C. Tomlin. *Mechanistic Interpretability for Steering Vision-Language-Action Models.* 2025.

[7] S. Marks et al. *Sparse Feature Circuits: Discovering and Editing Interpretable Causal Graphs in Language Models.* ICLR, 2025.

[8] M. Kekić et al. *Learning Nonlinear Causal Reductions to Explain Reinforcement Learning Policies.* ICLR, 2026.

[9] X. Zhou et al. *LIBERO-PRO: Towards Robust and Fair Evaluation of Vision-Language-Action Models Beyond Memorization.* 2025.

[10] Y. Fang et al. *When Vision Overrides Language: Evaluating and Mitigating Counterfactual Failures in VLAs.* 2026.

[11] C. Kim et al. *LIBERO-Para: A Diagnostic Benchmark and Metrics for Paraphrase Robustness in VLA Models.* 2026.

[12] T. Kolda and B. Bader. *Tensor Decompositions and Applications.* SIAM Review, 2009.

[13] E. Stoudenmire and D. Schwab. *Supervised Learning with Tensor Networks.* NeurIPS, 2016.

[14] G. Chrysos et al. *Deep Polynomial Neural Networks.* CVPR and TPAMI, 2020 to 2021.

A Supporting results and decisions

A.1 Expanded attention screen

Within the same seed-0 ViT checkpoint, the corrected fixed-control attention analysis covers all eight blocks and twelve heads. At rank 128, $83/96$ block-head pairs favor the exact weight-derived subspace. The median ratio is 1.94, the mean is 2.58, and blocks 2 and 3 provide useful depth controls. At ranks 8 and 32, only $54/96$ and $60/96$ pairs favor the exact subspace. The effect is therefore depth-dependent and strongest at the wider tested rank.

A.2 Prospectively fixed-rank visual bottleneck

The behavioral intervention acts at the width-384 visual-to-joint interface and keeps rank 96. Projector construction uses official Object tasks 0–3. Rank 96 was fixed before any corrected-basis outcome was observed, and the frozen evaluation uses tasks 8–9. Across three checkpoints and 300 canonical episodes, the full policy succeeds on 272 trials. The action-Jacobian projector retains $259/272 = 95.2\%$ of those successes, with a paired checkpoint-task-episode resampling interval of 91.2–98.6%. The activation-energy projector retains $267/272 = 98.2\%$, with an interval of 95.8–100.0%. Three fixed equal-rank random projectors retain $37/272 = 13.6\%$, $9/272 = 3.3\%$, and $22/272 = 8.1\%$.

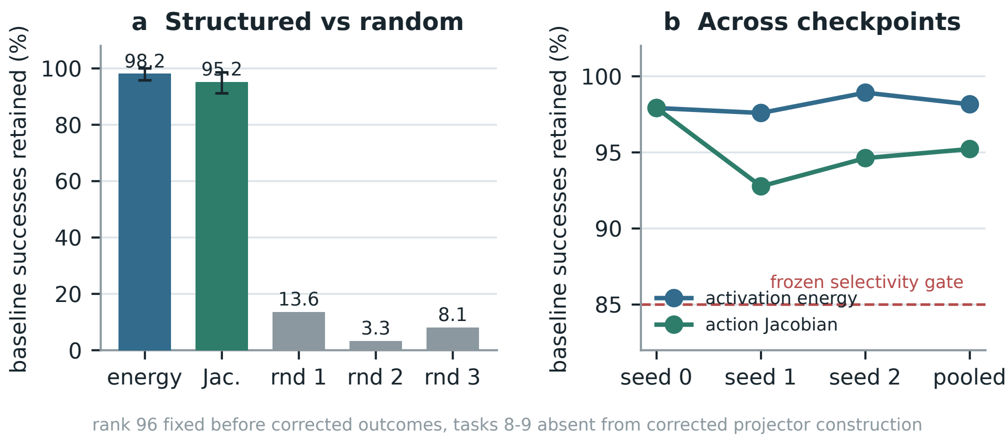

Figure 4: **Corrected closed-loop visual bottleneck.** Both learned rank-96 projectors preserve nearly all baseline successes across three checkpoints, while equal-rank random projectors do not. Error bars in panel (a) are deterministic paired cluster-resampling intervals. Tasks 8–9 are absent from corrected projector construction, but not from policy training. This is not a clean rank-selection-holdout result.

The action-Jacobian basis does not outperform the matched activation-energy basis. The stronger action-specificity gate therefore does not pass. The supported claim is structured visual sufficiency relative to random projectors. This intervention is data-derived and remains separate from the exact weight-only attention decomposition.

A.3 Numerical scope of rational conversion

The deployed degree-$[2/2]$ least-squares rational fit is exact as a rational computation even though it approximates RMS normalization. Its numerator has coefficients $(5.32997, 7.53552, 0.327274)$ and the denominator has coefficients $(1, 9.64827, 2.60664)$. Every denominator coefficient is positive, which certifies $Q(v) > 0$ for all $v \geq 0$. We instrument all 53 RationalNorm module instances, each invoked once per forward pass, on $4{,}096$ deterministic inputs for each of three convolutional checkpoints. Across $544{,}542{,}720$ activation rows, the minimum observed denominator is 1.0016 and no row is nonfinite. Comparing deployed float32 normalization with an otherwise identical pass that recomputes only RationalNorm scales in float64, the largest checkpoint-level primary metric, defined as the maximum of normalized- and physical-action NRMSE, is 2.414×10^{-6} against a frozen 10^{-3} gate.

The certificate also bounds what the operator approximates. Observed normalized mean-square ratios v span 1.66×10^{-4} to 17.17, outside the nominal fit interval $[0.1, 10]$. Across site-checkpoint pairs, the minimum coverages are 60.018% in that interval and 60.022% with rational-versus-exact-RMS scale error at most 3.4%. The frozen 99% secondary coverage gates therefore do not pass. The positive claim is therefore global absence of nonnegative denominator poles and cache-conditioned floating-point fidelity to the deployed rational operator, not uniform agreement with exact RMS scaling. Matched alternative-normalizer runtime overhead remains unmeasured.

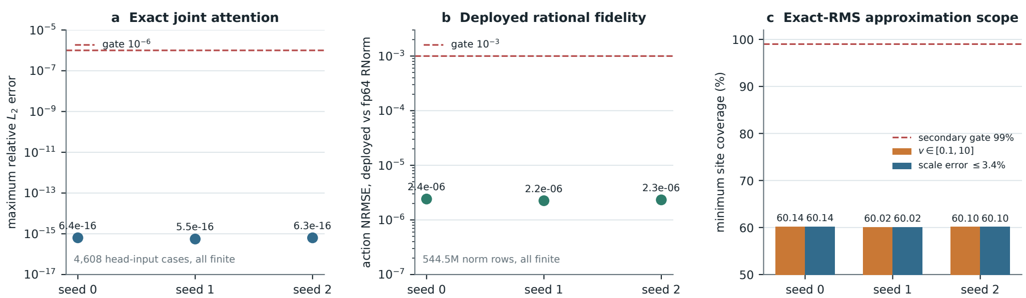

Figure 5: **Frozen structural certificates on three independent convolutional checkpoints.** Panels (a) and (b) pass their primary gates with large margins. Panel (a) reconstructs all $4{,}608$ audited joint-attention head-input cases. Panel (b) compares deployed float32 normalization with an otherwise identical pass that recomputes only RationalNorm scales in float64. Panel (c) reports the separately frozen exact-RMS approximation coverage, which does not pass its 99% secondary gate.