
Toward Auditable Vision-Language Control

A Fully Tensor-Decomposable VLA

Specialist capability and exact weight-level diagnostics

Anonymous Author(s)

Affiliation

Address

email

Abstract

Tensor decomposition factors a large table of interactions into a small set of low-rank components that can be ranked and inspected. Applying it directly to a conventional robot policy is difficult because softmax, ordinary activations, and normalization interrupt the required algebraic form. We introduce χ -VLA, a compact vision-language-action policy whose learned map uses only bilinear, linear, and rational operations. Its checkpoint therefore defines an explicit rational tensor program, making interactions among image, instruction, state, and action coordinates accessible from the weights. A runtime operator audit of the seed-0 20.1M ViT checkpoint finds no incompatible learned operation, and local reconstructions reproduce its deployed rational and bilinear branches near the float32 precision floor. Across three independently trained 20.1M ViT policies, closed-loop LIBERO-Object success is $85.3\% \pm 4.6\%$ over 500 canonical trials per checkpoint. A preregistered elementwise prediction-mean ensemble reaches $467/500 = 93.4\%$ at the cost of three forward passes. On a separate three-checkpoint convolutional instantiation, independent reconstruction covers all 4,608 joint-attention head-input cases with worst relative error 6.36×10^{-16} . Exact-attention records naming the same three ViT checkpoint artifacts derive the softmax-free attention decomposition and numerically replay every attention module on one fixed cached input per checkpoint, with worst relative error 2.56×10^{-7} . In a seed-0 ViT checkpoint, all 36 primary block-head tests favor its rank-128 subspace over equal-rank random controls, with median fidelity ratio 2.44. Separately, prospectively fixed rank-96 visual projectors retain 95.2–98.2% of baseline successes on tasks absent from corrected projector construction, while three equal-rank random projectors retain only 3.3–13.6%. The evidence establishes strong specialist capability together with exact, checkpoint-native access to learned attention structure. Such access is relevant when deployment failures may involve visual shortcuts or weak language use. The present results do not establish broad robotic generalization, robust instruction grounding, selective behavioral editing, or an exact end-to-end decomposition of the full policy.

1 Introduction

Deployable vision-language robot policies need more than aggregate success. When a policy fails, developers need to distinguish possible visual shortcuts, weak instruction use, state-estimation errors, and unstable action generation. Conventional checkpoints expose little of this structure directly. Post-hoc probes can identify correlated activations, but their conclusions depend on a chosen dataset and need not reveal how the deployed weights construct an action.

Tensor decomposition offers a complementary, checkpoint-native analysis surface. A tensor is a multidimensional table, and a decomposition rewrites that table as a sum of simpler low-rank factors, much as a matrix factorization finds dominant directions in a matrix. If a network is built from additions, multiplications, and ratios of polynomials, its weights define an explicit table of input interactions. By *fully tensor-decomposable*, we mean that every learned block preserves this algebraic form. This is an operator-closure claim, not a claim that we materialize one compact end-to-end coefficient tensor. Its degree and representation size grow rapidly with depth. For example, one coefficient can connect an image coordinate, an instruction coordinate, and an action coordinate. Factoring the coefficient table ranks high-energy cross-modal interactions directly from the checkpoint without first fitting an activation probe. Those factors are candidate structure, not automatically human-readable concepts or causal explanations. This structure does not certify semantic grounding by itself. It provides a direct weight-level surface for subsequent grounding and failure analyses.

Robot control makes this goal useful and difficult. Visual shortcuts can produce high benchmark success while ignoring the instruction, and small action errors compound in closed loop. Standard softmax, nonlinear activations, and square-root normalization break the explicit polynomial form. Continuous actions also cannot rely on the external categorical sampling used by a language model. Exact structural access is useful only if the restricted policy still controls reliably.

We therefore study a prerequisite for deployment-oriented diagnostics:

Can a specialist VLA retain strong closed-loop capability when every deployed learned operation is polynomial or rational, while exposing exact layerwise structure from its weights?

We test capability and structural access in two fully rational encoder instantiations. The same three ViT checkpoints establish closed-loop capability and support the exact attention audit. A separate convolutional instantiation tests whether structural certificates transfer across visual encoders. These experiments test whether a useful diagnostic surface can coexist with control, not whether that surface already certifies grounding or supports a validated edit.

Our contributions are:

1. We construct a complete VLA from tensor-compatible primitives. A rational normalization supplies input-specific scaling without leaving the rational function class.
2. We evaluate three fully rational 20.1M ViT policies on 1,500 canonical LIBERO-Object trials. They reach $85.3\% \pm 4.6\%$, while a fixed prediction-mean ensemble reaches $467/500 = 93.4\%$.
3. We mechanically audit the seed-0 20.1M ViT checkpoint and explicitly reconstruct deployed rational and bilinear branches from their weights.
4. We extend exact weight-derived decomposition through individual attention layers with a reduced-Gram calculation that avoids materializing the full high-order tensor. In a separate closed-loop audit, learned rank-96 visual projectors retain 95.2–98.2% of baseline successes, versus 3.3–13.6% for random projectors.

This paper does not claim current state of the art. It also does not claim that comparable mechanisms are impossible to obtain from conventional policies. Our narrower result is that an algebraic VLA can make exact layerwise weight structure directly accessible without giving up strong specialist control.

Certified by the current evidence	Not certified by the current evidence
Operator membership in the audited ViT checkpoint	Semantic instruction grounding
Exact coefficients for individual bilinear attention layers	Causal visual-shortcut attribution
Weight-derived attention subspaces and numerical replay	Safe coefficient editing and physical robustness

Table 1: **Deployment diagnostic contract.** Exact checkpoint access narrows what can be inspected, but it is not itself a grounding or deployment certificate.

77 2 A rational tensor robot policy

78 2.1 Tensor-compatible primitives

79 Each primitive below replaces a familiar transformer component while keeping an explicit coefficient description. Linear maps contribute first-order terms. Multiplying learned linear projections
80 creates higher-order interactions whose coefficients remain functions of the weights, rather than an
81 unstructured nonlinear black box.
82

83 **Homogeneous coordinate.** We write $\bar{x} = [1; x]$. The constant coordinate allows a multiplicative
84 layer to represent bias, linear, and quadratic terms in one contraction.

85 **Bilinear feed-forward block.** For $L, R \in \mathbb{R}^{r \times (d+1)}$ and $D \in \mathbb{R}^{d \times r}$,

$$\text{BFFN}(x) = D[(L\bar{x}) \odot (R\bar{x})], \quad y_o = \sum_{i,j} \left(\sum_{\rho} D_{o\rho} L_{\rho i} R_{\rho j} \right) \bar{x}_i \bar{x}_j. \quad (1)$$

86 The inner sum is an order-three table indexed by output coordinate and two input coordinates. The
87 three matrices give its CP factorization, the tensor analogue of a low-rank matrix factorization,
88 without materializing every table entry.

89 **Softmax-free bilinear attention.** Each head forms two query-key systems and one value stream:

$$Y = C [M \odot (Q_1 K_1^T) \odot (Q_2 K_2^T) / s] V. \quad (2)$$

90 Here M is a fixed causal mask, C is fixed row scaling, and s is fixed. All learned query, key, value,
91 and output maps are linear. Equation 2 is therefore a polynomial in its input activations. Unlike
92 softmax attention, every multiplicative interaction can be expanded into coefficients determined by
93 the trained weights.

94 **Rational per-instance normalization.** RMS normalization contains an input-dependent square
95 root. Let $m(x) = d^{-1} \sum_{j=1}^d x_j^2 + \epsilon$, and let $s_0 > 0$ be the frozen running mean-square scale. We
96 deploy a fixed rational approximation $r(v) = P(v)/Q(v)$ to $v^{-1/2}$ over $v \in [0.1, 10]$:

$$v = \frac{m(x)}{s_0}, \quad \text{RNorm}(x) = \frac{x}{\sqrt{s_0}} r(v). \quad (3)$$

97 Here *rational* means a ratio of two polynomials. Projective coordinates carry each activation as a
98 numerator-denominator pair (N, D) representing N/D . Linear, bilinear, residual, and normalization
99 operations update this pair. The model therefore keeps input-specific rescaling while remaining
100 exactly representable as a rational tensor network. Training and deployment use r itself. The
101 exactness claims below therefore concern faithful evaluation of the deployed rational computation,
102 not uniform agreement with RMSNorm. Appendix A.3 quantifies that distinction.




Figure 1: **Deployed χ -VLA map.** The strongest checkpoints use one vision stream and a linear continuous-action head. Every learned component inside the dashed boundary is linear, bilinear, or rational.

2.2 Architecture and training

Both encoder variants follow the high-level VLA sequence of visual encoding, language and state embedding, joint processing, and action decoding [1]. They use a 64×64 agent-view image, an eight-dimensional robot state, a 26-token task vocabulary, and eight action-query tokens. Both share a width-384, eight-block, twelve-head joint backbone and a linear head that maps each action-query state to seven continuous action dimensions. The 20.1M ViT uses four bilinear-attention vision blocks and has 20,137,352 learned parameters. It supplies the closed-loop capability, runtime operator census, exact-attention, and reduced-Gram results. The 19.2M convolutional variant instead uses three stride-2 bilinear convolution stages that produce an 8×8 token grid. Three separate convolutional files receive the joint-attention and rational-normalizer certificates reported in Appendix A.3.

Training uses 66,984 demonstration frames, 40,000 AdamW updates, batch size 256, cosine decay, bfloat16 arithmetic, and an exponential moving average. Closed-loop evaluation uses a 180° camera rotation to match the demonstration convention, canonical LIBERO initial states, and fail-loud checks that verify every requested state was loaded.

2.3 Mechanical certificate on the seed-0 20.1M ViT checkpoint

Architecture definitions can differ from deployed code. We therefore audit the deployed seed-0 20.1M ViT checkpoint at runtime and reconstruct representative operations from saved weights.




Figure 2: **Mechanical audit of the seed-0 20.1M ViT rational checkpoint.** The runtime path invokes 112 linear maps, 74 rational normalizers, two embeddings, and one linear patch projection. No softmax, standard activation, or ordinary normalization is present. Reconstruction errors are measured over 54,784 real activation rows.

The normalization reconstruction has maximum absolute error 1.64×10^{-6} . A bilinear FFN reconstruction has median relative error 1.40×10^{-7} . These residuals are bounded by the deployed module’s float32 statistic. The audit certifies operator membership and local reconstruction. It does not materialize one compact polynomial for the full eight-layer transformer.

3 Protocol-aligned closed-loop capability

3.1 Evaluation protocol

We evaluate all ten LIBERO-Object tasks with the following fixed protocol:

- 50 canonical trials per task
- a 280-step cap
- three independently trained ViT checkpoints
- 500 rollouts per checkpoint and 1,500 rollouts in total




Figure 3: **Verified ViT capability.** Left: three independently trained checkpoints and their fixed elementwise prediction-mean ensemble. Right: per-task three-checkpoint mean and ensemble success. The ensemble uses three forward passes.

3.2 Results

Method	Parameters	Inference	Success
χ -VLA, rational ViT	20.1M	1 pass	85.3% $\pm$ 4.6%
χ -VLA, prediction-mean ensemble	$3 \times 20.1M$	3 passes	93.4%

Table 2: **Closed-loop LIBERO-Object success.** Each ViT checkpoint receives 500 canonical trials. The single-model row is the mean $\pm$ sample standard deviation across three checkpoints. The ensemble is one fixed composite evaluated on the same 500 states.

The three checkpoints reach $426/500 = 85.2\%$, $404/500 = 80.8\%$, and $450/500 = 90.0\%$. Elementwise averaging of their predicted action chunks reaches $467/500 = 93.4\%$. This is 8.1 points above the member mean and 3.4 points above the strongest member. The ensemble requires three complete forward passes and is not one 20.1M-parameter policy. Every episode-level record, task index, and protocol field is retained. The artifact manifest binds each immutable result file and recorded checkpoint basename to a checkpoint hash verified separately. Member and ensemble shards span different recorded GPU families and predate explicit matrix-precision fields, so the arithmetic gaps are descriptive rather than a matched-hardware ensemble ablation.

3.3 What this result does and does not establish

The result shows that the rational architecture supports high closed-loop specialist capability. It is an existence claim for a compact tensor-decomposable policy, not a claim of broad robotic generalization or state of the art.

4 Exact weight-derived structure through attention

4.1 Exact layerwise construction

Equation 2 contains two query-key products and one value stream. Expanding the linear maps gives a high-order coefficient table. Each entry says how strongly one combination of query, key, and

value coordinates contributes to an output. It is *weight-derived* because no examples or fitted probe are needed to construct it. On a tiny four-token, width-four, one-head layer, direct evaluation and an independent explicit polynomial agree to 2.00×10^{-15} . The internal core identity agrees to 1.24×10^{-14} .

The deployed attention also normalizes each query and key with Equation 3. Each normalization multiplies all channels of one token by the same scalar rational factor. Those factors can be pulled outside the bilinear dot products. The polynomial attention core is unchanged, while explicit numerator and denominator factors carry the normalization. The rational extension agrees to 3.61×10^{-16} . On one fixed cached input per checkpoint, we numerically replay all four vision and eight joint attention modules in each of three records that name the same ViT checkpoint artifacts as the capability results. This covers 36 module-input cases and 384 architectural heads. The worst relative ℓ_2 error is 2.56×10^{-7} and the worst absolute difference is 2.23×10^{-5} . This residual comes from the deployed implementation’s intentional float32 statistic, while the reconstruction itself uses the learned coefficients. The algebraic identity is input-general. The fixed inputs provide a numerical replay audit.

We separately certify three convolutional checkpoints from the same rational architecture family. Their vision stacks contain no attention, so we independently rebuild all eight joint modules and all twelve heads per module from the saved weights. Sixteen deterministically selected cache inputs per checkpoint yield 384 module cases and 4,608 head-input cases. Every output is finite and the worst relative ℓ_2 error is 6.36×10^{-16} , far below the frozen 10^{-6} gate. This same-checkpoint certificate covers the joint-attention equation at each normalized layer input. It does not rebuild the convolutional stem, feed-forward blocks, residual composition, final normalization, or action head as one end-to-end contraction.

Naively materializing the real width-384 coefficient tensor would require approximately 3.2×10^{15} values. We instead contract the CP factors into an $O(d^2)$ reduced Gram. This smaller matrix plays the role of a covariance matrix: its leading eigenvectors rank the directions that carry the most coefficient energy without constructing the full table. At toy multi-head scale, it matches brute-force materialization to 2.13×10^{-14} .

4.2 Trained policy result

The reconstruction certificate above spans three ViT checkpoints. The following subspace-ranking experiment is a descriptive analysis of the seed-0 ViT checkpoint. We apply the exact weight-derived calculation to all twelve heads in early, middle-late, and final backbone blocks 0, 6, and 7. Random projector banks are sampled once, then fixed and reused for all evaluation examples. The weights choose the subspace. Real cached activations are used only afterward to measure how faithfully that subspace preserves the layer’s computation.




Figure 4: **Exact weight-derived attention structure.** Ratios above one favor the exact weight-derived subspace over a random equal-rank subspace. Each point in the right panel is one trained attention head.

At ranks 8, 32, and 128, respectively, 80.6%, 91.7%, and 100% of the 36 block-head pairs favor the exact subspace. Their median ratios are 1.028, 1.171, and 2.441. At rank 128, the mean ratio is 3.033 and the strongest head reaches 10.135. The rank-128 block means are 3.08, 3.45, and 2.56 for blocks 0, 6, and 7. The 36 block-head pairs are architectural units within one checkpoint, not independent model-level replicates.

This extends exact weight-only construction through individual attention layers. It is not an exact end-to-end decomposition of the complete policy. Composition across all eight blocks still faces rapid degree and representation growth.

A separate intervention tests behavioral sufficiency at the visual-to-joint interface rather than deriving a projector from the exact attention factors. An action-Jacobian projector estimated on official Object tasks 0–3, with rank 96 fixed before corrected outcomes, retains 259/272 full-policy successes on tasks 8–9 across three ViT checkpoints. These tasks are absent from corrected projector construction. An activation-energy projector retains 267/272, while three equal-rank random projectors retain 37/272, 9/272, and 22/272. All ten tasks appeared during policy training. This diagnoses a structured low-dimensional visual dependence relative to random projectors, not robust instruction grounding, unseen-task generalization, or unique action-Jacobian specificity. Appendix A.2 reports the paired controls and stability intervals.

5 Related work

VLA capability. OpenVLA established an open pretrained VLA and standardized LIBERO comparisons [1]. OpenVLA-OFT showed that continuous action chunks and simple regression can substantially improve adaptation and control efficiency [2]. These systems use large pretrained backbones. χ -VLA instead studies a compact from-scratch specialist under an algebraic restriction.

Tensor and polynomial networks. Polynomial networks and tensor networks use explicit multilinear structure for representation and learning [12–14]. Bilinear MLP work showed that weight eigendecomposition can expose low-rank features and circuits [3]. χ -nets introduced ODT for compositional bilinear networks [4]. We extend the setting to a continuous closed-loop policy and give an exact reduced construction for individual softmax-free attention layers.

Causal mechanisms and policy analysis. Bilinear IMPALA policies have combined competitive ProcGen control with weight-derived eigenfilters, followed by activation probing and control [5]. This is the closest architectural precedent for weight-based policy analysis. Recent work steers pretrained VLAs through causally identified activation directions [6]. Sparse feature circuits identify and edit causal feature graphs in conventional language models [7], while nonlinear causal reductions recover behavioral patterns in reinforcement-learning policies from rollout perturbations [8]. These results make a universal separation from conventional policies untenable. Our intended distinction is empirical: exact weight auditing in a strong continuous vision-language policy.

VLA robustness. Recent evaluations show that standard LIBERO success can coexist with weak robustness and language use [9–11]. This motivates model-internal diagnostics, but our exact weight analysis does not itself certify grounding. We therefore treat LIBERO-Object as a bounded specialist capability setting rather than evidence of broad deployment readiness.

6 Limitations and conclusion

χ -VLA demonstrates that strong specialist closed-loop control can coexist with a mechanically auditable rational graph and exact layerwise attention analysis. This combination is relevant to deployment diagnostics because it exposes trained coefficient structure without fitting a separate probe.

The scope is deliberately specific. LIBERO-Object is one specialist suite with fixed task instructions. The experiments do not show that the policy follows language when visual cues conflict, generalizes across task distributions, or transfers to physical robots. Exact reconstruction is certified layer by layer and does not yet yield one compact symbolic contraction of the full policy. The three-checkpoint ensemble trades three forward passes for its strongest capability result.

We therefore present χ -VLA as diagnostic infrastructure, not as a deployment-ready or robustly grounded robot policy. The ViT records connect capability, exact attention, and the reduced-Gram screen through common checkpoint basenames. Separate convolutional files replicate the joint-attention equation certificate across another visual encoder. Broader suites, counterfactual instruction tests, and physical-robot evaluation remain necessary for deployment claims. We do not claim a direct coefficient-edit interface.

Reproducibility. Training and evaluation use fixed task lists, explicit camera and action conventions, canonical-state manifests, checkpointed moving-average weights, and per-run result files. Historical Modal cache-backed analyses read dataset-native task metadata. The replacement Athena pipeline pins that metadata and maps dataset identifiers to official identifiers before defining task splits. A source-level audit compared all 271,996 cached frames across Object, Spatial, Goal, and LIBERO-10 with their pinned LeRobot revisions. Task identifiers, language joins, states, actions, and resized images matched exactly for every frame. The anonymous artifact releases raw structural-certificate JSONs and all 2,000 episode rows behind the ViT member and ensemble results. Immutable SHA-256 manifests and standard-library verifiers independently recompute the capability totals, visual-bottleneck counts and intervals, and the convolutional-attention and RationalNorm certificate decisions. Preflight checks fail if canonical initial states cannot be loaded.

References

- [1] M. J. Kim et al. *OpenVLA: An Open-Source Vision-Language-Action Model*. CoRL, 2024.
- [2] M. J. Kim, C. Finn, and P. Liang. *Fine-Tuning Vision-Language-Action Models: Optimizing Speed and Success*. 2025.
- [3] A. Rigg, T. Dooms, M. T. Pearce, J. Oramas, and L. Sharkey. *Bilinear MLPs Enable Weight-Based Mechanistic Interpretability*. ICLR, 2025.
- [4] T. Dooms, W. Gauderis, G. A. Wiggins, and J. Oramas. *Compositionality Unlocks Deep Interpretable Models*. 2025.
- [5] F. Oozeer, Y. Erisken, A. Rigg, J. Oramas, and L. Sharkey. *Bilinear Convolution Decomposition for Causal RL Interpretability*. 2024.
- [6] B. Häon, K. Stocking, I. Chuang, and C. Tomlin. *Mechanistic Interpretability for Steering Vision-Language-Action Models*. 2025.
- [7] S. Marks et al. *Sparse Feature Circuits: Discovering and Editing Interpretable Causal Graphs in Language Models*. ICLR, 2025.
- [8] M. Kekić et al. *Learning Nonlinear Causal Reductions to Explain Reinforcement Learning Policies*. ICLR, 2026.
- [9] X. Zhou et al. *LIBERO-PRO: Towards Robust and Fair Evaluation of Vision-Language-Action Models Beyond Memorization*. 2025.
- [10] Y. Fang et al. *When Vision Overrides Language: Evaluating and Mitigating Counterfactual Failures in VLAs*. 2026.
- [11] C. Kim et al. *LIBERO-Para: A Diagnostic Benchmark and Metrics for Paraphrase Robustness in VLA Models*. 2026.
- [12] T. Kolda and B. Bader. *Tensor Decompositions and Applications*. SIAM Review, 2009.
- [13] E. Stoudenmire and D. Schwab. *Supervised Learning with Tensor Networks*. NeurIPS, 2016.
- [14] G. Chrysos et al. *Deep Polynomial Neural Networks*. CVPR and TPAMI, 2020 to 2021.

A Supporting results and decisions

A.1 Expanded attention screen

Within the same seed-0 ViT checkpoint, the corrected fixed-control attention analysis covers all eight blocks and twelve heads. At rank 128, 83/96 block-head pairs favor the exact weight-derived subspace. The median ratio is 1.94, the mean is 2.58, and blocks 2 and 3 provide useful depth controls. At ranks 8 and 32, only 54/96 and 60/96 pairs favor the exact subspace. The effect is therefore depth-dependent and strongest at the wider tested rank.

A.2 Prospectively fixed-rank visual bottleneck

The behavioral intervention acts at the width-384 visual-to-joint interface and keeps rank 96. Projector construction uses official Object tasks 0–3. Rank 96 was fixed before any corrected-basis outcome was observed, and the frozen evaluation uses tasks 8–9. Across three checkpoints and 300 canonical episodes, the full policy succeeds on 272 trials. The action-Jacobian projector retains $259/272 = 95.2\%$ of those successes, with a paired checkpoint-task-episode resampling interval of 91.2–98.6%. The activation-energy projector retains $267/272 = 98.2\%$, with an interval of 95.8–100.0%. Three fixed equal-rank random projectors retain $37/272 = 13.6\%$, $9/272 = 3.3\%$, and $22/272 = 8.1\%$.




Figure 5: **Corrected closed-loop visual bottleneck.** Both learned rank-96 projectors preserve nearly all baseline successes across three checkpoints, while equal-rank random projectors do not. Error bars in panel (a) are deterministic paired cluster-resampling intervals. Tasks 8–9 are absent from corrected projector construction, but not from policy training. This is not a clean rank-selection-holdout result.

The action-Jacobian basis does not outperform the matched activation-energy basis. The stronger action-specificity gate therefore does not pass. The supported claim is structured visual sufficiency relative to random projectors. It is not evidence of language grounding, and the data-derived intervention remains separate from the exact weight-only attention decomposition.

293 A.3 Numerical scope of rational conversion

294 The deployed degree- $[2/2]$ least-squares rational fit is exact as a rational computation even though it
 295 approximates RMS normalization. Its numerator has coefficients (5.32997, 7.53552, 0.327274) and
 296 the denominator has coefficients (1, 9.64827, 2.60664). Every denominator coefficient is positive,
 297 which certifies $Q(v) > 0$ for all $v \geq 0$. We instrument all 53 RationalNorm module instances,
 298 each invoked once per forward pass, on 4,096 deterministic inputs for each of three convolutional
 299 checkpoints. Across 544,542,720 activation rows, the minimum observed denominator is 1.0016 and
 300 no row is nonfinite. Comparing deployed float32 normalization with an otherwise identical pass that
 301 recomputes only RationalNorm scales in float64, the largest checkpoint-level primary metric, defined
 302 as the maximum of normalized- and physical-action NRMSE, is 2.414×10^{-6} against a frozen 10^{-3}
 303 gate.

304 The certificate also bounds what the operator approximates. Observed normalized mean-square ratios
 305 v span 1.66×10^{-4} to 17.17, outside the nominal fit interval $[0.1, 10]$. Across site-checkpoint pairs,
 306 the minimum coverages are 60.018% in that interval and 60.022% with rational-versus-exact-RMS
 307 scale error at most 3.4%. The frozen 99% secondary coverage gates therefore do not pass. The
 308 positive claim is therefore global absence of nonnegative denominator poles and cache-conditioned
 309 floating-point fidelity to the deployed rational operator, not uniform agreement with exact RMS
 310 scaling. Matched alternative-normalizer runtime overhead remains unmeasured.




Figure 6: **Frozen structural certificates on three independent convolutional checkpoints.** Panels (a) and (b) pass their primary gates with large margins. Panel (a) reconstructs all 4,608 audited joint-attention head-input cases. Panel (b) compares deployed float32 normalization with an otherwise identical pass that recomputes only RationalNorm scales in float64. Panel (c) reports the separately frozen exact-RMS approximation coverage, which does not pass its 99% secondary gate.