
Toward Auditable Vision-Language Control

A Fully Tensor-Decomposable VLA

Single-suite strength, four-suite breadth, and exact weight-level diagnostics

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Tensor decomposition factors a large table of interactions into a small set of
2 low-rank components that can be ranked and inspected. Applying it directly
3 to a conventional robot policy is difficult because softmax, ordinary activations,
4 and normalization interrupt the required algebraic form. We introduce χ -VLA,
5 a compact vision-language-action policy whose learned map uses only bilinear,
6 linear, and rational operations. Its checkpoint therefore defines an explicit rational
7 tensor program, making interactions among image, instruction, state, and
8 action coordinates accessible from the weights. A runtime operator audit of the
9 seed-0 20.1M ViT checkpoint finds no incompatible learned operation, and local
10 reconstructions reproduce its deployed rational and bilinear branches near the
11 float32 precision floor. Across three independently trained 20.1M ViT policies,
12 closed-loop LIBERO-Object success is $85.3\% \pm 4.6\%$ over 500 canonical trials
13 per checkpoint. A preregistered elementwise prediction-mean ensemble reaches
14 $467/500 = 93.4\%$ at the cost of three forward passes. Separately, three jointly
15 trained checkpoints spanning 40 familiar LIBERO tasks reach $70.35\% \pm 1.32\%$
16 task-macro success, compared with $69.12\% \pm 2.33\%$ for parameter-matched conventional
17 controls under the same recipe. Their frozen three-pass mean-prediction
18 ensemble reaches 74.75% with every suite at or above 50%. At equal aggregate
19 updates, a frozen seed-0 joint checkpoint retains 95.16% of four suite specialists’
20 pooled success. On the three LIBERO-Object checkpoints, a prospectively frozen
21 paired test on 100 unique familiar-task states finds $247/300 = 82.3\%$ correct-
22 instruction success, versus $72/300 = 24.0\%$ with a co-present control instruction
23 and $78/300 = 26.0\%$ with an empty instruction. On disjoint episodes, zeroing
24 only the learned post-lookup token vectors drops success from $263/300 = 87.7\%$
25 to $63/300 = 21.0\%$, with positive gaps for all ten task aggregates. On a separate
26 three-checkpoint convolutional instantiation, independent reconstruction covers all
27 4,608 joint-attention head-input cases, all 384 joint-FFN module-input cases, and
28 all 384 complete joint-block input cases. The respective worst relative errors are
29 6.36×10^{-16} , 1.39×10^{-15} , and 1.5024×10^{-7} . Exact-attention records naming
30 the same three ViT checkpoint artifacts derive the softmax-free attention decomposition
31 and numerically replay every attention module on one fixed corrected
32 cached input per checkpoint, with worst relative error 2.016119×10^{-7} . A separate
33 sequential replay independently rebuilds the learned path through four vision
34 and eight joint blocks to the linear action output on 48 fixed inputs, with worst
35 relative error 3.9822×10^{-7} against a frozen 10^{-3} gate. In a seed-0 ViT checkpoint,
36 all 36 primary block-head tests favor its rank-128 subspace over equal-rank
37 random controls, with median fidelity ratio 2.44. Separately, prospectively fixed
38 rank-96 visual projectors retain 95.2–98.2% of baseline successes on tasks absent

from corrected projector construction, while three equal-rank random projectors retain only 3.3–13.6%. The evidence establishes strong single-suite and four-suite familiar-task capability together with exact, checkpoint-native access to learned attention structure and bounded familiar-task instruction necessity. Such access is relevant when deployment failures may involve visual shortcuts or weak language use. The present results do not establish unseen-task generalization, counterfactual goal following, selective behavioral editing, or an input-general compact symbolic decomposition of the full policy.

1 Introduction

Deployable vision-language robot policies need more than aggregate success. When a policy fails, developers need to distinguish possible visual shortcuts, weak instruction use, state-estimation errors, and unstable action generation. Conventional checkpoints expose little of this structure directly. Post-hoc probes can identify correlated activations, but their conclusions depend on a chosen dataset and need not reveal how the deployed weights construct an action.

Tensor decomposition offers a complementary, checkpoint-native analysis surface. A tensor is a multidimensional table, and a decomposition rewrites that table as a sum of simpler low-rank factors, much as a matrix factorization finds dominant directions in a matrix. If a network is built from additions, multiplications, and ratios of polynomials, its weights define an explicit table of input interactions. By *fully tensor-decomposable*, we mean that every learned block preserves this algebraic form. This is an operator-closure claim, not a claim that we materialize one compact end-to-end coefficient tensor. Its degree and representation size grow rapidly with depth. For example, one coefficient can connect an image coordinate, an instruction coordinate, and an action coordinate. Factoring the coefficient table ranks high-energy cross-modal interactions directly from the checkpoint without first fitting an activation probe. Those factors are candidate structure, not automatically human-readable concepts or causal explanations. This structure does not certify semantic grounding by itself. It provides a direct weight-level surface for subsequent grounding and failure analyses.

Robot control makes this goal useful and difficult. Visual shortcuts can produce high benchmark success while ignoring the instruction, and small action errors compound in closed loop. Standard softmax, nonlinear activations, and square-root normalization break the explicit polynomial form. Continuous actions also cannot rely on the external categorical sampling used by a language model. Exact structural access is useful only if the restricted policy still controls reliably.

We therefore study a prerequisite for deployment-oriented diagnostics:

Can a compact VLA retain strong closed-loop capability when every deployed learned operation is polynomial or rational, while exposing exact layerwise structure from its weights?

We test capability and structural access in two fully rational encoder instantiations. The same three ViT checkpoints establish closed-loop capability and support the exact attention audit. A separate convolutional instantiation tests whether structural certificates transfer across visual encoders. These experiments test whether a useful diagnostic surface can coexist with control, not whether that surface already certifies grounding or supports a validated edit.

Our contributions are:

1. We construct a complete VLA from tensor-compatible primitives. A rational normalization supplies input-specific scaling without leaving the rational function class.
2. We evaluate three fully rational 20.1M ViT policies on 1,500 canonical LIBERO-Object trials. They reach $85.3\% \pm 4.6\%$, while a fixed prediction-mean ensemble reaches $467/500 = 93.4\%$. Three jointly trained checkpoints average $70.35\% \pm 1.32\%$ over 40 familiar tasks, 1.23 points above matched conventional controls. Their fixed ensemble reaches 74.75%. A frozen seed-0 comparison retains 95.16% of four suite specialists.
3. We mechanically audit the seed-0 20.1M ViT checkpoint and explicitly reconstruct deployed rational and bilinear branches from their weights.

89 4. We extend exact weight-derived decomposition through individual attention layers with a
 90 reduced-Gram calculation that avoids materializing the full high-order tensor. In a separate
 91 closed-loop audit, learned rank-96 visual projectors retain 95.2–98.2% of baseline successes,
 92 versus 3.3–13.6% for random projectors.

93 This paper does not claim current state of the art. It also does not claim that comparable mechanisms
 94 are impossible to obtain from conventional policies. Our narrower result is that an algebraic VLA can
 95 make exact layerwise weight structure directly accessible without giving up strong specialist control.

Certified by the current evidence	Not certified by the current evidence
Operator membership in the audited ViT check-point	Semantic instruction grounding
Exact coefficients for individual bilinear attention layers	Causal visual-shortcut attribution
Weight-derived attention subspaces and numerical replay	Safe coefficient editing and physical robustness

Table 1: **Deployment diagnostic contract.** Exact checkpoint access narrows what can be inspected, but it is not itself a grounding or deployment certificate.

96 2 A rational tensor robot policy

97 2.1 Tensor-compatible primitives

98 Each primitive below replaces a familiar transformer component while keeping an explicit coefficient
 99 description. Linear maps contribute first-order terms. Multiplying learned linear projections
 100 creates higher-order interactions whose coefficients remain functions of the weights, rather than an
 101 unstructured nonlinear black box.

102 **Homogeneous coordinate.** We write $\bar{x} = [1; x]$. The constant coordinate allows a multiplicative
 103 layer to represent bias, linear, and quadratic terms in one contraction.

104 **Bilinear feed-forward block.** For $L, R \in \mathbb{R}^{r \times (d+1)}$ and $D \in \mathbb{R}^{d \times r}$,

$$\text{BFFN}(x) = D[(L\bar{x}) \odot (R\bar{x})], \quad y_o = \sum_{i,j} \left(\sum_{\rho} D_{o\rho} L_{\rho i} R_{\rho j} \right) \bar{x}_i \bar{x}_j. \quad (1)$$

105 The inner sum is an order-three table indexed by output coordinate and two input coordinates. The
 106 three matrices give its CP factorization, the tensor analogue of a low-rank matrix factorization,
 107 without materializing every table entry.

108 **Softmax-free bilinear attention.** Each head forms two query-key systems and one value stream:

$$Y = C[M \odot (Q_1 K_1^\top) \odot (Q_2 K_2^\top) / s] V. \quad (2)$$

109 Here M is a fixed causal mask, C is fixed row scaling, and s is fixed. All learned query, key, value,
 110 and output maps are linear. Equation 2 is therefore a polynomial in its input activations. Unlike
 111 softmax attention, every multiplicative interaction can be expanded into coefficients determined by
 112 the trained weights.

113 **Rational per-instance normalization.** RMS normalization contains an input-dependent square
 114 root. Let $m(x) = d^{-1} \sum_{j=1}^d x_j^2 + \epsilon$, and let $s_0 > 0$ be the frozen running mean-square scale. We
 115 deploy a fixed rational approximation $r(v) = P(v)/Q(v)$ to $v^{-1/2}$ over $v \in [0.1, 10]$:

$$v = \frac{m(x)}{s_0}, \quad \text{RNorm}(x) = \frac{x}{\sqrt{s_0}} r(v). \quad (3)$$

116 Here *rational* means a ratio of two polynomials. Projective coordinates carry each activation as a
 117 numerator-denominator pair (N, D) representing N/D . Linear, bilinear, residual, and normalization

operations update this pair. The model therefore keeps input-specific rescaling while remaining exactly representable as a rational tensor network. Training and deployment use r itself. The exactness claims below therefore concern faithful evaluation of the deployed rational computation, not uniform agreement with RMSNorm. Appendix A.5 certifies pole-free denominators and deployed arithmetic fidelity while retaining that boundary.




Figure 1: **Deployed χ -VLA map.** The strongest checkpoints use one vision stream and a linear continuous-action head. Every learned component inside the dashed boundary is linear, bilinear, or rational.

2.2 Architecture and training

Both encoder variants follow the high-level VLA sequence of visual encoding, language and state embedding, joint processing, and action decoding [1]. They use a 64×64 agent-view image, an eight-dimensional robot state, a 26-token task vocabulary, and eight action-query tokens. Both share a width-384, eight-block, twelve-head joint backbone and a linear head that maps each action-query state to seven continuous action dimensions. The 20.1M ViT uses four bilinear-attention vision blocks and has 20,137,352 learned parameters. It supplies the closed-loop capability, runtime operator census, exact-attention, and reduced-Gram results. The 19.2M convolutional variant instead uses three stride-2 bilinear convolution stages that produce an 8×8 token grid. Three separate convolutional files receive the joint-attention and rational-normalizer certificates reported in Appendix A.5.

Training uses 66,984 demonstration frames, 40,000 AdamW updates, batch size 256, cosine decay, bfloat16 arithmetic, and an exponential moving average. Closed-loop evaluation uses a 180° camera rotation to match the demonstration convention, canonical LIBERO initial states, and fail-loud checks that verify every requested state was loaded.

2.3 Mechanical certificate on the seed-0 20.1M ViT checkpoint

Architecture definitions can differ from deployed code. We therefore audit the deployed seed-0 20.1M ViT checkpoint at runtime and reconstruct representative operations from saved weights.

Mechanical operator audit of the seed-0 20.1M ViT checkpoint




Figure 2: **Mechanical audit of the seed-0 20.1M ViT rational checkpoint.** The runtime path invokes 112 linear maps, 74 rational normalizers, two embeddings, and one linear patch projection. No softmax, standard activation, or ordinary normalization is present. Reconstruction errors are measured over 54,784 real activation rows.

140 The normalization reconstruction has maximum absolute error 1.64×10^{-6} . A bilinear FFN re-
 141 construction has median relative error 1.40×10^{-7} . These residuals are bounded by the deployed
 142 module’s float32 statistic. The audit certifies operator membership and local reconstruction. It does
 143 not materialize one compact polynomial for the full eight-layer transformer.

144 3 Protocol-aligned closed-loop capability

145 3.1 Evaluation protocol

146 We evaluate all ten LIBERO-Object tasks with the following fixed protocol:

- 147 • 50 canonical trials per task
- 148 • a 280-step cap
- 149 • three independently trained ViT checkpoints
- 150 • 500 rollouts per checkpoint and 1,500 rollouts in total




Figure 3: **Verified ViT capability.** Left: three independently trained checkpoints and their fixed elementwise prediction-mean ensemble. Right: per-task three-checkpoint mean and ensemble success. The ensemble uses three forward passes.

3.2 Results

Method	Parameters	Inference	Success
χ -VLA, rational ViT	20.1M	1 pass	85.3% $\pm$ 4.6%
χ -VLA, prediction-mean ensemble	$3 \times 20.1M$	3 passes	93.4%

Table 2: **Closed-loop LIBERO-Object success.** Each ViT checkpoint receives 500 canonical trials. The single-model row is the mean $\pm$ sample standard deviation across three checkpoints. The ensemble is one fixed composite evaluated on the same 500 states.

The three checkpoints reach $426/500 = 85.2\%$, $404/500 = 80.8\%$, and $450/500 = 90.0\%$. Elementwise averaging of their predicted action chunks reaches $467/500 = 93.4\%$. This is 8.1 points above the member mean and 3.4 points above the strongest member. The ensemble requires three complete forward passes and is not one 20.1M-parameter policy. Every episode-level record, task index, and protocol field is retained. The artifact manifest binds each immutable result file and recorded checkpoint basename to a checkpoint hash verified separately. Member and ensemble shards span different recorded GPU families and predate explicit matrix-precision fields, so the arithmetic gaps are descriptive rather than a matched-hardware ensemble ablation.

3.3 What this result does and does not establish

The result shows that the rational architecture supports high closed-loop specialist capability. It is an existence claim for a compact tensor-decomposable policy, not a claim of broad robotic generalization or state of the art.

3.4 One policy across four familiar suites

Three jointly trained 20.1M checkpoints reach $70.35\% \pm 1.32\%$ task-macro success over 40 familiar LIBERO tasks and 50 trials per task, 1.23 points above parameter-matched conventional controls at $69.12\% \pm 2.33\%$ under the same recipe. Their frozen three-pass mean-prediction ensemble reaches 74.75% , with every suite and 34/40 tasks at or above 50%. At equal aggregate updates, a frozen

seed-0 joint checkpoint retains 95.16% of four suite specialists, above the fixed 85% gate. Because all 40 tasks appeared in training, this establishes multi-suite capacity, not unseen-task transfer.

3.5 Prospectively frozen closed-loop instruction necessity

High task success alone does not show that a policy uses its instruction. Before these outcomes were observed, we fixed a paired test on canonical episodes 40–49 for every task and the same three checkpoints. Each checkpoint-state pair was rolled out for the full 280-step horizon under the correct familiar instruction, one deterministic outcome-independent co-present control instruction, and an empty instruction. All conditions shared the exact initial physical input and their order was chosen by a frozen hash. Co-presence was verified from simulator observation fields and does not guarantee unoccluded camera visibility. The empty string encodes tokenizer BOS plus 31 unmasked PAD tokens, together with the model’s learned common BOS.




Figure 4: **Closed-loop instruction necessity on familiar tasks.** Left: exact successes among 100 paired canonical states per checkpoint. Right: correct-instruction success minus each control, pooling the three checkpoint outcomes for each task. All ten task gaps are positive for both controls at the task-aggregate level.

Across 100 unique canonical states and three checkpoints, the correct instruction succeeds on $247/300 = 82.3\%$ of checkpoint-state pairs. The co-present control succeeds on $72/300 = 24.0\%$ and the empty instruction on $78/300 = 26.0\%$, giving paired gaps of 58.3 and 56.3 points. Correct-instruction success is 83/100, 77/100, and 87/100 by checkpoint, and both paired gaps are positive for every checkpoint aggregate. All ten pooled task aggregates are strictly positive for both controls. One-sided exact McNemar tests give $p = 2.02 \times 10^{-45}$ and $p = 1.52 \times 10^{-44}$. These pooled tests use 300 checkpoint-state pairs and are not cluster-robust inference. A frozen 20,000-draw task-stratified state-cluster bootstrap resamples 100 unique task-episode states while retaining all three checkpoint outcomes. Its 95% intervals are [51.0, 65.7] and [50.0, 62.3] points. Every prospectively frozen gate passes.

This establishes instruction necessity for familiar LIBERO-Object prompts, objects, scenes, tasks, and checkpoints. Appendix A.1 gives a disjoint paired test that isolates the learned lexical-embedding path. Neither test establishes counterfactual goal following, paraphrase robustness, unseen-object or compositional grounding, cross-suite or physical transfer, or a benefit unique to tensor decomposability.

For the VLM4RWD setting, this separates familiar-task language use from deployment-grade grounding. The control changes the instruction while holding the initial physical input fixed, but it does not test policy-unseen language or scenes.

4 Exact weight-derived structure through attention

4.1 Exact layerwise construction

Equation 2 contains two query-key products and one value stream. Expanding the linear maps gives a high-order coefficient table. Each entry says how strongly one combination of query, key, and

value coordinates contributes to an output. It is *weight-derived* because no examples or fitted probe are needed to construct it. On a tiny four-token, width-four, one-head layer, direct evaluation and an independent explicit polynomial agree to 2.00×10^{-15} . The internal core identity agrees to 1.24×10^{-14} .

The deployed attention also normalizes each query and key with Equation 3. Each normalization multiplies all channels of one token by the same scalar rational factor. Those factors can be pulled outside the bilinear dot products. The polynomial attention core is unchanged, while explicit numerator and denominator factors carry the normalization. The rational extension agrees to 3.61×10^{-16} . On one fixed corrected cached input per checkpoint, we numerically replay all four vision and eight joint attention modules in each of three records that name the same ViT checkpoint artifacts as the capability results. This covers 36 module-input cases and 384 architectural heads. The worst relative ℓ_2 error is 2.016119×10^{-7} and the worst absolute difference is 2.44×10^{-5} . This residual comes from the deployed implementation’s intentional float32 statistic, while the reconstruction itself uses the learned coefficients. The algebraic identity is input-general. The one-input-per-checkpoint records provide a layerwise numerical replay audit, not an exact end-to-end policy decomposition.

We separately certify three convolutional checkpoints from the same rational architecture family. Independent reconstruction covers 4,608 joint-attention head-input cases, 384 joint-FFN module-input cases, and 384 complete joint-block input cases. Their respective worst relative errors are 6.36×10^{-16} , 1.3871×10^{-15} , and 1.5024×10^{-7} , all below their frozen gates. These are input-conditioned layer and block certificates.

A separately frozen ViT audit partitions 36,864 joint-attention head-input cases into vision, instruction, robot-state, and action-query sources. Worst reconstruction error is 5.0741×10^{-7} , while mean coherent-energy shares are 67.7%, 5.1%, 19.2%, and 7.9%. The shares are descriptive rather than causal importance or grounding.

Finally, a prospectively frozen sequential replay rebuilds the learned path from prepared tensors through all 12 transformer blocks to the linear head. Across 48 inputs and 816 stages, worst final-action relative error is 3.9822×10^{-7} against a frozen 10^{-3} gate. This fixed-input numerical certificate is not a compact symbolic or input-general proof.

Naively materializing the real width-384 coefficient tensor would require approximately 3.2×10^{15} values. We instead contract the CP factors into an $O(d^2)$ reduced Gram. This smaller matrix plays the role of a covariance matrix: its leading eigenvectors rank the directions that carry the most coefficient energy without constructing the full table. At toy multi-head scale, it matches brute-force materialization to 2.13×10^{-14} .

4.2 Trained policy result

The reconstruction certificate spans three checkpoints. A descriptive seed-0 analysis applies the exact weight-derived calculation to all twelve heads in blocks 0, 6, and 7. Random projector banks are sampled once and reused. The weights choose the subspace, while cached activations only score how faithfully it preserves the computation.




Figure 5: **Exact weight-derived attention structure.** Ratios above one favor the exact weight-derived subspace over a random equal-rank subspace. Each point in the right panel is one trained attention head.

At ranks 8, 32, and 128, 80.6%, 91.7%, and 100% of the 36 block-head pairs favor the exact subspace, with median ratios 1.028, 1.171, and 2.441. At rank 128, the mean is 3.033, the maximum is 10.135, and block means are 3.08, 3.45, and 2.56. These pairs are architectural units within one checkpoint, not independent model-level replicates.

This extends exact weight-only construction through individual attention layers, not to a compact symbolic end-to-end policy decomposition. Composition still faces rapid degree and representation growth.

A separate intervention tests behavioral sufficiency at the visual-to-joint interface. A rank-96 action-Jacobian projector fixed from Object tasks 0–3 retains 259/272 successes on tasks 8–9 across three checkpoints. Activation energy retains 267/272, while equal-rank random projectors retain 37/272, 9/272, and 22/272. All tasks appeared during policy training. This diagnoses structured visual dependence relative to random projectors, not grounding, unseen-task generalization, or unique action-Jacobian specificity. Appendix A.4 gives the paired controls and intervals.

5 Related work

Capability, mechanisms, and robustness. OpenVLA and OpenVLA-OFT give pretrained VLA baselines [1, 2], while polynomial and tensor networks provide our multilinear foundation [12–14]. Bilinear models expose low-rank circuits from trained weights [3, 4]. Policy analysis and causal activation studies make our distinction empirical [5–8]. Because LIBERO success can coexist with weak robustness and language use [9–11], the claim remains bounded to familiar tasks from the training suites.

6 Limitations and conclusion

χ -VLA shows that strong single-suite control and substantial four-suite familiar-task breadth can coexist with exact access to trained structure. The paired controls establish familiar-task instruction necessity, not counterfactual goal following, unseen-task generalization, or physical transfer. The learned path is replayed on fixed inputs, not given as one compact symbolic contraction or input-general proof, and the strongest ensemble costs three passes. We present χ -VLA as diagnostic infrastructure, not a deployment-ready or robustly grounded policy. Common ViT checkpoints connect capability, exact attention, and reduced-Gram analysis, while separate convolutional files replicate the attention equation certificate. We do not claim a coefficient-edit interface. The anonymous artifact and Appendix A.2 provide full reproducibility details.

References

- [1] M. J. Kim et al. *OpenVLA: An Open-Source Vision-Language-Action Model*. CoRL, 2024.
- [2] M. J. Kim, C. Finn, and P. Liang. *Fine-Tuning Vision-Language-Action Models: Optimizing Speed and Success*. 2025.
- [3] A. Rigg, T. Doods, M. T. Pearce, J. Oramas, and L. Sharkey. *Bilinear MLPs Enable Weight-Based Mechanistic Interpretability*. ICLR, 2025.
- [4] T. Doods, W. Gauderis, G. A. Wiggins, and J. Oramas. *Compositionality Unlocks Deep Interpretable Models*. 2025.
- [5] F. Oozeer, Y. Erisken, A. Rigg, J. Oramas, and L. Sharkey. *Bilinear Convolution Decomposition for Causal RL Interpretability*. 2024.
- [6] B. Häon, K. Stocking, I. Chuang, and C. Tomlin. *Mechanistic Interpretability for Steering Vision-Language-Action Models*. 2025.
- [7] S. Marks et al. *Sparse Feature Circuits: Discovering and Editing Interpretable Causal Graphs in Language Models*. ICLR, 2025.
- [8] M. Kekić et al. *Learning Nonlinear Causal Reductions to Explain Reinforcement Learning Policies*. ICLR, 2026.

- [9] X. Zhou et al. *LIBERO-PRO: Towards Robust and Fair Evaluation of Vision-Language-Action Models Beyond Memorization*. 2025.
- [10] Y. Fang et al. *When Vision Overrides Language: Evaluating and Mitigating Counterfactual Failures in VLAs*. 2026.
- [11] C. Kim et al. *LIBERO-Para: A Diagnostic Benchmark and Metrics for Paraphrase Robustness in VLA Models*. 2026.
- [12] T. Kolda and B. Bader. *Tensor Decompositions and Applications*. SIAM Review, 2009.
- [13] E. Stoudenmire and D. Schwab. *Supervised Learning with Tensor Networks*. NeurIPS, 2016.
- [14] G. Chrysos et al. *Deep Polynomial Neural Networks*. CVPR and TPAMI, 2020 to 2021.

A Supporting results and decisions

A.1 Causal isolation of the learned instruction-embedding path

Token IDs are lookup addresses. A learned embedding table turns those addresses into vectors that the policy can combine with vision and robot state. Before observing outcomes, we fixed a paired test that isolates this conversion. On canonical episodes 30–39, disjoint from the instruction-necessity episodes, both arms received the identical correct 32-token ID tensor. The ablated arm replaced only the output of the token-embedding lookup with exact zeros. Sequence length, token positions, vision, robot state, embodiment, the separate learned common BOS, action queries, transformer weights, and action head were preserved and hash-checked.




Figure 6: **Learned lexical-embedding-path necessity on familiar tasks.** Left: closed-loop success for the intact path and exact post-lookup zeroing, using 100 paired states per checkpoint and 300 pooled pairs. Right: every task aggregate has a positive paired gap. The pooled marker shows the frozen task-stratified state-cluster 95% bootstrap interval.

The intact path succeeds on $263/300 = 87.7\%$ of checkpoint-state pairs, compared with $63/300 = 21.0\%$ after zeroing the learned token vectors. The paired gap is 66.7 points, with 203 intact-only and 3 zeroed-only outcomes. Full-path success is 84%, 87%, and 92% by checkpoint, and the corresponding gaps are 78, 57, and 65 points. All ten task aggregates are positive. A pooled one-sided exact paired test gives $p = 1.42 \times 10^{-56}$, but this test is not cluster-robust because checkpoints reuse the same 100 states. The frozen 20,000-draw task-stratified state-cluster bootstrap interval is $[61.7, 71.3]$ points and is our cluster-aware inference. All eight prospectively fixed gates pass.

This establishes that the learned lexical-embedding path is necessary for these familiar LIBERO-Object tasks and three fixed policies. It does not establish semantic grounding, paraphrase robustness, unseen-object or compositional generalization, cross-suite or physical transfer, or a benefit unique to tensor decomposability.

A.2 Reproducibility details

Training and evaluation use fixed task lists, explicit camera and action conventions, canonical-state manifests, checkpointed moving-average weights, and per-run result files. Historical Modal cache-backed analyses read dataset-native task metadata. The replacement Athena pipeline pins that metadata and maps dataset identifiers to official identifiers before defining task splits. A source-level audit compared all 271,996 cached frames across Object, Spatial, Goal, and LIBERO-10 with their pinned LeRobot revisions. Task identifiers, language joins, states, actions, and resized images matched exactly for every frame. The anonymous artifact releases raw structural-certificate JSONs and all 2,000 episode rows behind the ViT member and ensemble results. Immutable SHA-256 manifests and standard-library verifiers independently recompute the capability totals, matched conventional comparison, four-suite ensemble and retention arithmetic, visual-bottleneck counts and intervals, and the attention, FFN, complete-block, modality-partition, full learned-forward, and RationalNorm certificate decisions. Preflight checks fail if canonical initial states cannot be loaded.

A.3 Expanded attention screen

Within the same seed-0 ViT checkpoint, the corrected fixed-control attention analysis covers all eight blocks and twelve heads. At rank 128, 83/96 block-head pairs favor the exact weight-derived subspace. The median ratio is 1.94, the mean is 2.58, and blocks 2 and 3 provide useful depth controls. At ranks 8 and 32, only 54/96 and 60/96 pairs favor the exact subspace. The effect is therefore depth-dependent and strongest at the wider tested rank.

A.4 Prospectively fixed-rank visual bottleneck

The behavioral intervention acts at the width-384 visual-to-joint interface and keeps rank 96. Projector construction uses official Object tasks 0–3. Rank 96 was fixed before any corrected-basis outcome was observed, and the frozen evaluation uses tasks 8–9. Across three checkpoints and 300 canonical episodes, the full policy succeeds on 272 trials. The action-Jacobian projector retains 259/272 = 95.2% of those successes, with a paired checkpoint-task-episode resampling interval of 91.2–98.6%. The activation-energy projector retains 267/272 = 98.2%, with an interval of 95.8–100.0%. Three fixed equal-rank random projectors retain 37/272 = 13.6%, 9/272 = 3.3%, and 22/272 = 8.1%.




Figure 7: **Corrected closed-loop visual bottleneck.** Both learned rank-96 projectors preserve nearly all baseline successes across three checkpoints, while equal-rank random projectors do not. Error bars in panel (a) are deterministic paired cluster-resampling intervals. Tasks 8–9 are absent from corrected projector construction, but not from policy training. This is not a clean rank-selection-holdout result.

The action-Jacobian basis does not outperform the matched activation-energy basis. The stronger action-specificity gate therefore does not pass. The supported claim is structured visual sufficiency relative to random projectors. It is not evidence of language grounding, and the data-derived intervention remains separate from the exact weight-only attention decomposition.

346 A.5 Numerical scope of rational conversion

347 The deployed degree- $[2/2]$ least-squares rational fit is exact as a rational computation even though it
 348 approximates RMS normalization. Its numerator has coefficients (5.32997, 7.53552, 0.327274) and
 349 the denominator has coefficients (1, 9.64827, 2.60664). Every denominator coefficient is positive,
 350 which certifies $Q(v) > 0$ for all $v \geq 0$. We instrument all 53 RationalNorm module instances,
 351 each invoked once per forward pass, on 4,096 deterministic inputs for each of three convolutional
 352 checkpoints. Across 544,542,720 activation rows, the minimum observed denominator is 1.0016
 353 and no row is nonfinite. Comparing deployed float32 normalization with an otherwise identical pass
 354 that recomputes only RationalNorm scales in float64, the largest checkpoint-level primary metric,
 355 defined as the maximum of normalized- and physical-action NRMSE, is 2.414×10^{-6} against a
 356 frozen 10^{-3} gate. This primary certificate concerns pole-free denominators and floating-point fidelity
 357 to the deployed rational operator. It does not claim uniform agreement with exact RMS scaling or
 358 measure matched alternative-normalizer runtime overhead.

359 The reconstruction ladder extends from individual attention heads through FFNs and per-block replay
 360 to a separate sequential learned-forward certificate. That final record starts at exact prepared model
 361 tensors and reaches the linear action head on 48 fixed inputs. It is numerical and input-conditioned,
 362 not a compact symbolic contraction or an input-general full-policy proof. The ViT source-partition
 363 audit remains layerwise and descriptive. Figure 8 reports the frozen numerical gates and the depth
 364 profile without assigning causal importance.




Figure 8: **Frozen reconstruction ladder and descriptive ViT source partition.** Panel (a) reports observed-to-gate ratios for independent Conv attention, joint-FFN, and complete joint-block reconstruction, the ViT source-partition identity, the sequential ViT learned-forward action replay, and primary RationalNorm action fidelity. All plotted primary numerical gates pass. Panel (b) shows descriptive mean coherent-energy shares after the attention output projection across 384 checkpoint-input cases per joint block. The depth trend is not a causal importance or grounding result. The learned-forward replay excludes preprocessing, fixed action unnormalization, chunk execution, simulator dynamics, and success.