
Specialist Physical AI Need Not Be Opaque

A Fully Tensor-Decomposable VLA Policy

Protocol-aligned control and exact weight analysis

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Tensor decomposition turns a large table of learned interactions into simpler low-
2 rank factors that can be ranked and inspected. Conventional robot policies contain
3 operations that block this direct weight-level view. We introduce χ -VLA, a 20.1M-
4 parameter vision-language-action policy built only from bilinear, linear, and rational
5 learned operations. Its weights can therefore be reorganized into explicit interaction
6 tensors rather than treated only as an opaque checkpoint. A mechanical audit and
7 numerical reconstructions verify that the deployed policy preserves this form.
8 Trained from scratch on LIBERO-Object, convolutional χ -VLA reaches $93.7\% \pm$
9 0.8% over three seeds under the published protocol. In a separate ViT instantiation,
10 we reconstruct every attention module in three checkpoints with worst relative
11 error 2.56×10^{-7} , and exact weight-derived subspaces beat equal-rank random
12 controls in all 36 primary block-head tests. The result is a capable and structurally
13 inspectable specialist family with exact access to trained attention coefficients. It
14 does not yet provide an exact end-to-end decomposition of the full policy.

15 1 Introduction

16 Modern robot policies can achieve high success while giving little indication of why an action was
17 chosen. Their checkpoints contain millions of weights, but those weights do not directly show which
18 image regions, words, and robot-state coordinates interact. Post-hoc probes can find correlated
19 activations, yet a correlation is not automatically a mechanism or a safe editing target.

20 A tensor is simply a multidimensional table. Tensor decomposition rewrites that table as a sum of
21 low-rank factors, analogous to finding dominant directions in a matrix. If every learned block uses
22 additions, multiplications, and ratios of polynomials, the full policy retains an explicit interaction
23 table defined by its weights. Its decomposed factors then provide candidate control mechanisms that
24 can be compared across layers and tested by intervention [3, 4]. This is what we mean by a fully
25 tensor-decomposable policy.

26 The restriction is nontrivial in robotics. Softmax, ordinary activations, and square-root normalization
27 break the algebraic form, while continuous actions require precise closed-loop outputs. The potential
28 payoff is also concrete: a weight-derived mechanism could help diagnose whether a policy follows
29 language, keys on a visual shortcut, or uses a fragile control feature.

30 We study the narrower but testable question:

31 *Can a specialist VLA retain strong closed-loop capability when every deployed*
 32 *learned operation is polynomial or rational, while supporting exact weight analy-*
 33 *sis?*

34 We test capability and structural analysis in two fully rational encoder instantiations. The convo-
 35 lutional checkpoints establish closed-loop capability. Separate ViT checkpoints support the exact
 36 attention audit.

37 Our contributions are:

- 38 1. We construct a complete VLA from tensor-compatible primitives. A rational normalization
 39 supplies input-specific scaling without leaving the rational function class.
- 40 2. We evaluate two fully rational 20.1M-parameter policies under a protocol aligned with
 41 published LIBERO-Object results. The strongest single architecture reaches $93.7\% \pm 0.8\%$
 42 over three seeds, and a three-checkpoint ensemble reaches 96.2%.
- 43 3. We mechanically audit the real trained policy and explicitly reconstruct deployed rational
 44 and bilinear branches from their weights.
- 45 4. We extend exact weight-derived decomposition through individual attention layers with a
 46 reduced-Gram calculation that avoids materializing the full high-order tensor.

47 This paper does not claim current state of the art. It also does not claim that comparable mechanisms
 48 are impossible to obtain from conventional policies. The empirical question is whether the algebraic
 49 policy makes exact weight structure directly accessible.

50 2 A rational tensor robot policy

51 2.1 Tensor-compatible primitives

52 Each primitive replaces a familiar transformer component while keeping explicit coefficients. Lin-
 53 ear maps contribute first-order terms, while multiplying learned projections creates higher-order
 54 interactions that remain determined by the weights.

55 **Homogeneous coordinate.** We write $\bar{x} = [1; x]$. The constant coordinate allows a multiplicative
 56 layer to represent bias, linear, and quadratic terms in one contraction.

57 **Bilinear feed-forward block.** For $L, R \in \mathbb{R}^{r \times (d+1)}$ and $D \in \mathbb{R}^{d \times r}$,

$$\text{BFFN}(x) = D[(L\bar{x}) \odot (R\bar{x})], \quad y_o = \sum_{i,j} \left(\sum_{\rho} D_{o\rho} L_{\rho i} R_{\rho j} \right) \bar{x}_i \bar{x}_j. \quad (1)$$

58 The inner sum is a table indexed by one output and two input coordinates. The three matrices give its
 59 CP factorization, the tensor analogue of a low-rank matrix factorization.

60 **Softmax-free bilinear attention.** Each head forms two query-key systems and one value stream:

$$Y = C[M \odot (Q_1 K_1^\top) \odot (Q_2 K_2^\top) / s] V. \quad (2)$$

61 Here M is a fixed causal mask, C is fixed row scaling, and s is fixed. All learned query, key,
 62 value, and output maps are linear. Equation 2 is therefore a polynomial in its input activations, with
 63 interaction coefficients determined by the weights.

64 **Rational per-instance normalization.** RMS normalization contains an input-dependent square
 65 root. We instead deploy a fixed rational approximation $r(z) = P(z)/Q(z)$ to $1/\sqrt{z + \epsilon}$:

$$\text{RNorm}(x) = x r \left(\frac{1}{d} \sum_{j=1}^d x_j^2 \right). \quad (3)$$

66 Here rational means a ratio of polynomials. Projective coordinates carry each activation as a
 67 numerator-denominator pair (N, D) representing N/D . Updating this pair preserves input-specific
 68 rescaling without leaving the rational tensor representation.

69 2.2 Architecture and training

70 The policy follows the high-level VLA sequence of visual encoding, language and state embedding,
 71 joint processing, and action decoding [1]. It uses a 64×64 agent-view image, an eight-dimensional
 72 robot state, a 26-token task vocabulary, and eight action-query tokens. The joint backbone has width
 73 384, eight blocks, and twelve heads. The action head maps each action-query state directly to seven
 74 continuous action dimensions. The complete model has 20,137,352 learned parameters.

75 Training uses 66,984 demonstration frames, 40,000 AdamW updates, batch size 256, cosine decay,
 76 bfloat16 arithmetic, and an exponential moving average. Closed-loop evaluation uses a 180° camera
 77 rotation to match the demonstration convention, canonical LIBERO initial states, and fail-loud checks
 78 that verify every requested state was loaded.

79 2.3 Mechanical certificate on the trained checkpoint

80 Architecture definitions can differ from deployed code. We therefore audit the real 189/200 check-
 81 point at runtime and reconstruct representative operations from saved weights.

82 The normalization reconstruction has maximum absolute error 1.64×10^{-6} . A bilinear FFN re-
 83 construction has median relative error 1.40×10^{-7} . These residuals are bounded by the deployed
 84 module’s float32 statistic. The audit certifies operator membership and local reconstruction, as
 85 summarized in Appendix Figure 3. It does not materialize one compact polynomial for the full
 86 eight-layer transformer.

87 3 Protocol-aligned closed-loop capability

88 3.1 Evaluation protocol

89 We evaluate all ten LIBERO-Object tasks with:

- 90 • 50 canonical trials per task
- 91 • a 280-step cap
- 92 • three independently trained seeds per architecture
- 93 • 500 rollouts per seed and 1,500 rollouts per architecture

94 Published OpenVLA and Diffusion Policy results are protocol-aligned references. They were not
 95 rerun in our codebase, so differences may still include implementation and training choices.




Figure 1: **Protocol-aligned LIBERO-Object capability.** Small dots are χ -VLA training seeds. The ensemble star averages three convolutional checkpoints and uses three forward passes per action chunk. External references are reported published values.

96 3.2 Results

Method	Parameters	Training seeds	Success
χ -VLA, rational conv	20.1M	3	93.7% $\pm$ 0.8%
χ -VLA, conv ensemble	$3 \times 20.1\text{M}$	$3 \rightarrow 1$	96.2%
OpenVLA-7B [1]	7B	3	88.4% $\pm$ 0.8%
Diffusion Policy [1]	not matched	3	92.5% $\pm$ 0.7%

Table 1: **Closed-loop success under aligned evaluation settings.** The two external rows are historical references rather than current state-of-the-art claims.

97 The convolutional seeds score 94.0%, 94.4%, and 92.8%. Their narrow spread shows that high
 98 specialist capability is reproduced across three independently trained checkpoints. These are historical
 99 Modal aggregate records. A locked replay is producing immutable episode-level outcomes.

100 Separately, the paired ensemble evaluation re-evaluates the three members at 93.6%, 95.0%, and
 101 93.2%, for a 93.9% mean. Elementwise prediction averaging reaches 96.2%. This is 2.3 points
 102 above that paired mean and 1.2 points above its strongest member. The gain is largest on tomato
 103 sauce, butter, and ketchup. The ensemble requires three complete forward passes and is not one
 104 20.1M-parameter policy.

105 3.3 What this result does and does not establish

106 The result shows that the rational architecture supports high closed-loop specialist capability. It is an
 107 existence and reproducibility claim for a compact tensor-decomposable policy, not a claim of broad
 108 robotic generalization or state of the art.

109 4 Exact weight-derived structure through attention

110 4.1 Exact layerwise construction

111 Equation 2 contains two query-key products and one value stream. Expanding the linear maps gives
 112 a coefficient table whose entries measure how combinations of query, key, and value coordinates
 113 contribute to an output. No examples or fitted probe are needed to construct it. On a tiny four-
 114 token, width-four, one-head layer, direct evaluation and an independent explicit polynomial agree to
 115 2.00×10^{-15} . The internal core identity agrees to 1.24×10^{-14} .

116 The deployed attention also normalizes each query and key with Equation 3. Each normalization
 117 multiplies all channels of one token by the same scalar rational factor. Those factors can be pulled
 118 outside the bilinear dot products. The polynomial attention core is unchanged, while explicit
 119 numerator and denominator factors carry the normalization. The rational extension agrees to $3.61 \times$
 120 10^{-16} . We then reconstruct all four vision and eight joint attention modules in each of three
 121 trained checkpoints, covering 384 heads in total. The worst relative ℓ_2 error is 2.56×10^{-7} and the
 122 worst absolute difference is 2.23×10^{-5} . This residual comes from the deployed implementation’s
 123 intentional float32 statistic, while the reconstruction itself uses the learned coefficients.

124 Naively materializing the real width-384 coefficient tensor would require approximately 3.2×10^{15}
 125 values. We instead form an $O(d^2)$ reduced Gram. Like a covariance matrix, its leading eigenvectors
 126 rank important directions without constructing the full table. It matches brute-force materialization to
 127 2.13×10^{-14} at toy scale.

128 4.2 Trained policy result

129 For the subspace-ranking experiment, we apply the exact weight-derived calculation to all twelve
 130 heads in early, middle-late, and final backbone blocks 0, 6, and 7. Random projector banks are
 131 sampled once, then fixed and reused for all examples. The weights choose the subspace, while cached
 132 activations are used only afterward to score how well it preserves the layer’s computation.

133 At ranks 8, 32, and 128, respectively, 80.6%, 91.7%, and 100% of the 36 block-head pairs favor the
 134 exact subspace. Their median ratios are 1.028, 1.171, and 2.441. At rank 128, the mean ratio is 3.033

135 and the strongest head reaches 10.135. The rank-128 block means are 3.08, 3.45, and 2.56 for blocks
136 0, 6, and 7.

137 This extends exact weight-only construction through individual attention layers. It is not an exact
138 end-to-end decomposition of the complete policy. Composition across all eight blocks still faces
139 rapid degree and representation growth.

140 5 Related work and limitations

141 **Related work.** OpenVLA and OpenVLA-OFT establish pretrained VLA and continuous-control
142 baselines [1, 2]. Bilinear MLPs, χ -nets, and bilinear IMPALA policies show that multiplicative
143 weights can expose low-rank structure and support subsequent controls [3–5]. Recent work also
144 steers pretrained VLAs through activation directions [6]. Our distinction is the exact weight-derived
145 construction in a policy designed as a rational tensor program. Tensor and polynomial networks
146 provide the broader multilinear foundation [12–14]. VLA robustness benchmarks show that standard
147 LIBERO success can coexist with weak language use [9–11].

148 **Limitations and conclusion.** The benchmark is one specialist suite, so it does not establish broad
149 robotic generalization. Exact reconstruction is certified layer by layer and does not yet produce one
150 compact symbolic contraction of the full policy. The capability result uses convolutional checkpoints,
151 while the exact analysis uses separate ViT checkpoints. Evaluation uses fixed task instructions and
152 does not establish robust instruction grounding. Within that scope, rational restrictions support strong
153 LIBERO-Object control and exact trained-weight certificates. We do not claim a direct coefficient-edit
154 interface.

155 **Reproducibility.** The final artifact will provide code, configurations, checkpoint hashes, exact
156 certificates, and per-episode logs. Historical Modal cache-backed analyses read dataset-native task
157 metadata. The replacement Athena pipeline pins that metadata and maps dataset identifiers to official
158 identifiers before defining task splits. Canonical-state checks fail loudly when requested states are
159 unavailable.

References

- [1] M. J. Kim et al. *OpenVLA: An Open-Source Vision-Language-Action Model*. CoRL, 2024.
- [2] M. J. Kim, C. Finn, and P. Liang. *Fine-Tuning Vision-Language-Action Models: Optimizing Speed and Success*. 2025.
- [3] A. Rigg, T. Dooms, M. T. Pearce, J. Oramas, and L. Sharkey. *Bilinear MLPs Enable Weight-Based Mechanistic Interpretability*. ICLR, 2025.
- [4] T. Dooms, W. Gauderis, G. A. Wiggins, and J. Oramas. *Compositionality Unlocks Deep Interpretable Models*. 2025.
- [5] F. Oozeer, Y. Erisken, A. Rigg, J. Oramas, and L. Sharkey. *Bilinear Convolution Decomposition for Causal RL Interpretability*. 2024.
- [6] B. Häon, K. Stocking, I. Chuang, and C. Tomlin. *Mechanistic Interpretability for Steering Vision-Language-Action Models*. 2025.
- [7] S. Marks et al. *Sparse Feature Circuits: Discovering and Editing Interpretable Causal Graphs in Language Models*. ICLR, 2025.
- [8] M. Kekić et al. *Learning Nonlinear Causal Reductions to Explain Reinforcement Learning Policies*. ICLR, 2026.
- [9] X. Zhou et al. *LIBERO-PRO: Towards Robust and Fair Evaluation of Vision-Language-Action Models Beyond Memorization*. 2025.
- [10] Y. Fang et al. *When Vision Overrides Language: Evaluating and Mitigating Counterfactual Failures in VLAs*. 2026.
- [11] C. Kim et al. *LIBERO-Para: A Diagnostic Benchmark and Metrics for Paraphrase Robustness in VLA Models*. 2026.
- [12] T. Kolda and B. Bader. *Tensor Decompositions and Applications*. SIAM Review, 2009.
- [13] E. Stoudenmire and D. Schwab. *Supervised Learning with Tensor Networks*. NeurIPS, 2016.
- [14] G. Chrysos et al. *Deep Polynomial Neural Networks*. CVPR and TPAMI, 2020 to 2021.

185 A Supporting results and decisions

186 A.1 Per-task ensemble behavior




Figure 2: **Per-task effect of prediction averaging in the paired ensemble evaluation.** Gray points are the three convolutional checkpoints, evaluated with 50 canonical trials per task. Curves compare their taskwise mean with elementwise action-prediction averaging. The ensemble raises overall success from a 93.9% member mean to 96.2%, with its largest gains on tomato sauce, butter, and ketchup. It requires three forward passes and is not a single 20.1M-parameter policy.

187 A.2 Expanded attention screen

188 The corrected fixed-control attention analysis covers all eight blocks and twelve heads. At rank 128,
 189 83/96 block-head pairs favor the exact weight-derived subspace. The median ratio is 1.94, the mean
 190 is 2.58, and blocks 2 and 3 provide useful depth controls. At ranks 8 and 32, only 54/96 and 60/96
 191 pairs favor the exact subspace. The effect is therefore depth-dependent and strongest at the wider
 192 tested rank.

193 A.3 Numerical scope of rational conversion

194 The deployed Padé computation is exactly rational even though it approximates RMS normalization.
 195 Final certification must report its polynomial degrees, coefficients, denominator margin, observed
 196 activation range, tail error, action drift, and runtime overhead. Operator membership alone does not
 197 establish practical numerical safety under distribution shift.

Mechanical decomposability audit of the trained 189/200 checkpoint




Figure 3: **Mechanical audit of the trained rational checkpoint.** The runtime path invokes 112 linear maps, 74 rational normalizers, two embeddings, and one linear patch projection. No softmax, standard activation, or ordinary normalization is present. Reconstruction errors are measured over 54,784 real activation rows.