

---

# A Capable, Fully Tensor-Decomposable Vision-Language-Action Policy

Exact rational conversion, protocol-aligned control, and weight-derived  
analysis

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

Tensor decomposition factors a large table of interactions into a small set of low-rank components that can be ranked and inspected. Applying it directly to a conventional robot policy is difficult because softmax, ordinary activations, and normalization interrupt the required algebraic form. We introduce  $\chi$ -VLA, a compact vision-language-action policy whose learned map uses only bilinear, linear, and rational operations. Its checkpoint therefore defines an explicit rational tensor program, making interactions among image, instruction, state, and action coordinates accessible from the weights. A runtime operator audit of the seed-0 20.1M ViT checkpoint finds no incompatible learned operation, and local reconstructions reproduce its deployed rational and bilinear branches near the float32 precision floor. Across three independently trained 20.1M ViT policies, closed-loop LIBERO-Object success is  $85.3\% \pm 4.6\%$  over 500 canonical trials per checkpoint. A preregistered elementwise prediction-mean ensemble reaches  $467/500 = 93.4\%$  at the cost of three forward passes. Across the three checkpoints, a prospectively frozen paired test on 100 unique familiar-task states finds  $247/300 = 82.3\%$  correct-instruction success, versus  $72/300 = 24.0\%$  with a co-present control instruction and  $78/300 = 26.0\%$  with an empty instruction. On disjoint episodes, zeroing only the learned post-lookup token vectors drops success from  $263/300 = 87.7\%$  to  $63/300 = 21.0\%$ , with positive gaps for all ten task aggregates. On a separate three-checkpoint convolutional instantiation, independent reconstruction covers all 4,608 joint-attention head-input cases, all 384 joint-FFN module-input cases, and all 384 complete joint-block input cases. The respective worst relative errors are  $6.36 \times 10^{-16}$ ,  $1.39 \times 10^{-15}$ , and  $1.5024 \times 10^{-7}$ . Exact-attention records naming the same three ViT checkpoint artifacts derive the softmax-free attention decomposition and numerically replay every attention module on one fixed corrected cached input per checkpoint, with worst relative error  $2.016119 \times 10^{-7}$ . A separate sequential replay independently rebuilds the learned path through four vision and eight joint blocks to the linear action output on 48 fixed inputs, with worst relative error  $3.9822 \times 10^{-7}$  against a frozen  $10^{-3}$  gate. In a seed-0 ViT checkpoint, all 36 primary block-head tests favor its rank-128 subspace over equal-rank random controls, with median fidelity ratio 2.44. Separately, prospectively fixed rank-96 visual projectors retain 95.2–98.2% of baseline successes on tasks absent from corrected projector construction, while three equal-rank random projectors retain only 3.3–13.6%. The evidence establishes strong specialist capability and exact layerwise access to trained attention coefficients, together with bounded familiar-task instruction necessity. It does not establish

37 broad robotic generalization, counterfactual goal following, or an input-general  
38 compact symbolic decomposition of the full policy.

## 39 1 Introduction

40 Modern robot policies can work well while remaining difficult to inspect. A checkpoint contains  
41 millions of learned numbers, but it does not directly reveal which combinations of image regions,  
42 instruction tokens, and robot state produce an action. Analysts therefore usually collect activations  
43 on a dataset and fit a separate probe. That can identify correlations, but it leaves open whether the  
44 same structure is present in the weights and whether it causes behavior.

45 Tensor decomposition offers a different analysis surface. A tensor is a multidimensional table, and a  
46 decomposition rewrites that table as a sum of simpler low-rank factors, much as a matrix factorization  
47 finds dominant directions in a matrix. If a network is built from additions, multiplications, and ratios  
48 of polynomials, its weights define an explicit table of input interactions. Decomposing that table can  
49 rank the combinations the model is constructed to use. By *fully tensor-decomposable*, we mean that  
50 every learned block preserves this algebraic form. This is an operator-closure claim, not a claim that  
51 we materialize one compact end-to-end coefficient tensor. Its degree and representation size grow  
52 rapidly with depth. For example, one coefficient can connect an image coordinate, an instruction  
53 coordinate, and an action coordinate. Factoring the coefficient table ranks high-energy cross-modal  
54 interactions directly from the checkpoint without first fitting an activation probe. Those factors are  
55 candidate structure, not automatically human-readable concepts or causal explanations. The aim is  
56 not merely compression. It is a direct path from parameters to candidate mechanisms that can be  
57 tested by intervention [3, 4].

58 Robot control makes this goal useful and difficult. Visual shortcuts can produce high benchmark  
59 success while ignoring the instruction, and small action errors compound in closed loop. Standard  
60 softmax, nonlinear activations, and square-root normalization break the explicit polynomial form.  
61 Continuous actions also cannot rely on the external categorical sampling used by a language model.  
62 A practical design must recover these functions without losing precise control.

63 We study the narrower but testable question:

64 *Can a specialist VLA retain strong closed-loop capability when every deployed*  
65 *learned operation is polynomial or rational, while supporting exact weight analy-*  
66 *sis?*

67 We test capability and structural analysis in two fully rational encoder instantiations. The same three  
68 ViT checkpoints establish closed-loop capability and support the exact attention audit. A separate  
69 convolutional instantiation tests whether the structural certificates transfer across visual encoders.

70 Our contributions are:

- 71 1. We construct a complete VLA from tensor-compatible primitives. A rational normalization  
72 supplies input-specific scaling without leaving the rational function class.
- 73 2. We evaluate three fully rational 20.1M ViT policies on 1,500 canonical LIBERO-Object  
74 trials. They reach  $85.3\% \pm 4.6\%$ , while a fixed prediction-mean ensemble reaches  
75  $467/500 = 93.4\%$ .
- 76 3. We mechanically audit the seed-0 20.1M ViT checkpoint and explicitly reconstruct deployed  
77 rational and bilinear branches from their weights.
- 78 4. We extend exact weight-derived decomposition through individual attention layers with a  
79 reduced-Gram calculation that avoids materializing the full high-order tensor. In a separate  
80 closed-loop audit, learned rank-96 visual projectors retain 95.2–98.2% of baseline successes,  
81 versus 3.3–13.6% for random projectors.

82 This paper does not claim current state of the art. It also does not claim that comparable mechanisms  
83 are impossible to obtain from conventional policies. The empirical question is whether the algebraic  
84 policy makes exact weight structure directly accessible.

## 85 2 A rational tensor robot policy

### 86 2.1 Tensor-compatible primitives

87 Each primitive below replaces a familiar transformer component while keeping an explicit coefficient description. Linear maps contribute first-order terms. Multiplying learned linear projections  
 88 creates higher-order interactions whose coefficients remain functions of the weights, rather than an  
 89 unstructured nonlinear black box.  
 90

91 **Homogeneous coordinate.** We write  $\bar{x} = [1; x]$ . The constant coordinate allows a multiplicative  
 92 layer to represent bias, linear, and quadratic terms in one contraction.

93 **Bilinear feed-forward block.** For  $L, R \in \mathbb{R}^{r \times (d+1)}$  and  $D \in \mathbb{R}^{d \times r}$ ,

$$\text{BFFN}(x) = D [(L\bar{x}) \odot (R\bar{x})], \quad y_o = \sum_{i,j} \left( \sum_{\rho} D_{o\rho} L_{\rho i} R_{\rho j} \right) \bar{x}_i \bar{x}_j. \quad (1)$$

94 The inner sum is an order-three table indexed by output coordinate and two input coordinates. The  
 95 three matrices give its CP factorization, the tensor analogue of a low-rank matrix factorization,  
 96 without materializing every table entry.

97 **Softmax-free bilinear attention.** Each head forms two query-key systems and one value stream:

$$Y = C [M \odot (Q_1 K_1^T) \odot (Q_2 K_2^T) / s] V. \quad (2)$$

98 Here  $M$  is a fixed causal mask,  $C$  is fixed row scaling, and  $s$  is fixed. All learned query, key, value,  
 99 and output maps are linear. Equation 2 is therefore a polynomial in its input activations. Unlike  
 100 softmax attention, every multiplicative interaction can be expanded into coefficients determined by  
 101 the trained weights.

102 **Rational per-instance normalization.** RMS normalization contains an input-dependent square  
 103 root. Let  $m(x) = d^{-1} \sum_{j=1}^d x_j^2 + \epsilon$ , and let  $s_0 > 0$  be the frozen running mean-square scale. We  
 104 deploy a fixed rational approximation  $r(v) = P(v)/Q(v)$  to  $v^{-1/2}$  over  $v \in [0.1, 10]$ :

$$v = \frac{m(x)}{s_0}, \quad \text{RNorm}(x) = \frac{x}{\sqrt{s_0}} r(v). \quad (3)$$

105 Here *rational* means a ratio of two polynomials. Projective coordinates carry each activation as a  
 106 numerator-denominator pair  $(N, D)$  representing  $N/D$ . Linear, bilinear, residual, and normalization  
 107 operations update this pair. The model therefore keeps input-specific rescaling while remaining  
 108 exactly representable as a rational tensor network. Training and deployment use  $r$  itself. The  
 109 exactness claims below therefore concern faithful evaluation of the deployed rational computation,  
 110 not uniform agreement with RMSNorm. Appendix A.5 certifies pole-free denominators and deployed  
 111 arithmetic fidelity while retaining that boundary.




Figure 1: **Deployed  $\chi$ -VLA map.** The strongest checkpoints use one vision stream and a linear continuous-action head. Every learned component inside the dashed boundary is linear, bilinear, or rational.

## 2.2 Architecture and training

Both encoder variants follow the high-level VLA sequence of visual encoding, language and state embedding, joint processing, and action decoding [1]. They use a  $64 \times 64$  agent-view image, an eight-dimensional robot state, a 26-token task vocabulary, and eight action-query tokens. Both share a width-384, eight-block, twelve-head joint backbone and a linear head that maps each action-query state to seven continuous action dimensions. The 20.1M ViT uses four bilinear-attention vision blocks and has 20,137,352 learned parameters. It supplies the closed-loop capability, runtime operator census, exact-attention, and reduced-Gram results. The 19.2M convolutional variant instead uses three stride-2 bilinear convolution stages that produce an  $8 \times 8$  token grid. Three separate convolutional files receive the joint-attention and rational-normalizer certificates reported in Appendix A.5.

Training uses 66,984 demonstration frames, 40,000 AdamW updates, batch size 256, cosine decay, bfloat16 arithmetic, and an exponential moving average. Closed-loop evaluation uses a  $180^\circ$  camera rotation to match the demonstration convention, canonical LIBERO initial states, and fail-loud checks that verify every requested state was loaded.

## 2.3 Mechanical certificate on the seed-0 20.1M ViT checkpoint

Architecture definitions can differ from deployed code. We therefore audit the deployed seed-0 20.1M ViT checkpoint at runtime and reconstruct representative operations from saved weights.




Figure 2: **Mechanical audit of the seed-0 20.1M ViT rational checkpoint.** The runtime path invokes 112 linear maps, 74 rational normalizers, two embeddings, and one linear patch projection. No softmax, standard activation, or ordinary normalization is present. Reconstruction errors are measured over 54,784 real activation rows.

The normalization reconstruction has maximum absolute error  $1.64 \times 10^{-6}$ . A bilinear FFN reconstruction has median relative error  $1.40 \times 10^{-7}$ . These residuals are bounded by the deployed module’s float32 statistic. The audit certifies operator membership and local reconstruction. It does not materialize one compact polynomial for the full eight-layer transformer.

## 3 Protocol-aligned closed-loop capability

### 3.1 Evaluation protocol

We evaluate all ten LIBERO-Object tasks with the following fixed protocol:

- 50 canonical trials per task
- a 280-step cap
- three independently trained ViT checkpoints
- 500 rollouts per checkpoint and 1,500 rollouts in total




Figure 3: **Verified ViT capability.** Left: three independently trained checkpoints and their fixed elementwise prediction-mean ensemble. Right: per-task three-checkpoint mean and ensemble success. The ensemble uses three forward passes.

### 3.2 Results

| Method                                | Parameters       | Inference | Success                            |
|---------------------------------------|------------------|-----------|------------------------------------|
| $\chi$ -VLA, rational ViT             | 20.1M            | 1 pass    | <b>85.3% <math>\pm</math> 4.6%</b> |
| $\chi$ -VLA, prediction-mean ensemble | $3 \times 20.1M$ | 3 passes  | <b>93.4%</b>                       |

Table 1: **Closed-loop LIBERO-Object success.** Each ViT checkpoint receives 500 canonical trials. The single-model row is the mean  $\pm$  sample standard deviation across three checkpoints. The ensemble is one fixed composite evaluated on the same 500 states.

The three checkpoints reach  $426/500 = 85.2\%$ ,  $404/500 = 80.8\%$ , and  $450/500 = 90.0\%$ . Elementwise averaging of their predicted action chunks reaches  $467/500 = 93.4\%$ . This is 8.1 points above the member mean and 3.4 points above the strongest member. The ensemble requires three complete forward passes and is not one 20.1M-parameter policy. Every episode-level record, task index, and protocol field is retained. The artifact manifest binds each immutable result file and recorded checkpoint basename to a checkpoint hash verified separately. Member and ensemble shards span different recorded GPU families and predate explicit matrix-precision fields, so the arithmetic gaps are descriptive rather than a matched-hardware ensemble ablation.

### 3.3 What this result does and does not establish

The result shows that the rational architecture supports high closed-loop specialist capability. It is an existence claim for a compact tensor-decomposable policy, not a claim of broad robotic generalization or state of the art.

### 3.4 Prospectively frozen closed-loop instruction necessity

High task success alone does not show that a policy uses its instruction. Before these outcomes were observed, we fixed a paired test on canonical episodes 40–49 for every task and the same three checkpoints. Each checkpoint-state pair was rolled out for the full 280-step horizon under the correct familiar instruction, one deterministic outcome-independent co-present control instruction, and an empty instruction. All conditions shared the exact initial physical input and their order was chosen by

a frozen hash. Co-presence was verified from simulator observation fields and does not guarantee unoccluded camera visibility. The empty string encodes tokenizer BOS plus 31 unmasked PAD tokens, together with the model’s learned common BOS.




Figure 4: **Closed-loop instruction necessity on familiar tasks.** Left: exact successes among 100 paired canonical states per checkpoint. Right: correct-instruction success minus each control, pooling the three checkpoint outcomes for each task. All ten task gaps are positive for both controls at the task-aggregate level.

Across 100 unique canonical states and three checkpoints, the correct instruction succeeds on 247/300 = 82.3% of checkpoint-state pairs. The co-present control succeeds on 72/300 = 24.0% and the empty instruction on 78/300 = 26.0%, giving paired gaps of 58.3 and 56.3 points. Correct-instruction success is 83/100, 77/100, and 87/100 by checkpoint, and both paired gaps are positive for every checkpoint aggregate. All ten pooled task aggregates are strictly positive for both controls. One-sided exact McNemar tests give  $p = 2.02 \times 10^{-45}$  and  $p = 1.52 \times 10^{-44}$ . These pooled tests use 300 checkpoint-state pairs and are not cluster-robust inference. A frozen 20,000-draw task-stratified state-cluster bootstrap resamples 100 unique task-episode states while retaining all three checkpoint outcomes. Its 95% intervals are [51.0, 65.7] and [50.0, 62.3] points. Every prospectively frozen gate passes.

This establishes instruction necessity for familiar LIBERO-Object prompts, objects, scenes, tasks, and checkpoints. Appendix A.1 gives a disjoint paired test that isolates the learned lexical-embedding path. Neither test establishes counterfactual goal following, paraphrase robustness, unseen-object or compositional grounding, cross-suite or physical transfer, or a benefit unique to tensor decomposability.

Within its stated boundary, this closes the possibility that the policy’s familiar-task success is independent of its instruction. It does not turn the specialist benchmark into a general language grounding result.

## 4 Exact weight-derived structure through attention

### 4.1 Exact layerwise construction

Equation 2 contains two query-key products and one value stream. Expanding the linear maps gives a high-order coefficient table. Each entry says how strongly one combination of query, key, and value coordinates contributes to an output. It is *weight-derived* because no examples or fitted probe are needed to construct it. On a tiny four-token, width-four, one-head layer, direct evaluation and an independent explicit polynomial agree to  $2.00 \times 10^{-15}$ . The internal core identity agrees to  $1.24 \times 10^{-14}$ .

The deployed attention also normalizes each query and key with Equation 3. Each normalization multiplies all channels of one token by the same scalar rational factor. Those factors can be pulled outside the bilinear dot products. The polynomial attention core is unchanged, while explicit numerator and denominator factors carry the normalization. The rational extension agrees to  $3.61 \times 10^{-16}$ . On one fixed corrected cached input per checkpoint, we numerically replay all four vision

193 and eight joint attention modules in each of three records that name the same ViT checkpoint artifacts  
 194 as the capability results. This covers 36 module-input cases and 384 architectural heads. The worst  
 195 relative  $\ell_2$  error is  $2.016119 \times 10^{-7}$  and the worst absolute difference is  $2.44 \times 10^{-5}$ . This residual  
 196 comes from the deployed implementation’s intentional float32 statistic, while the reconstruction itself  
 197 uses the learned coefficients. The algebraic identity is input-general. The one-input-per-checkpoint  
 198 records provide a layerwise numerical replay audit, not an exact end-to-end policy decomposition.

199 We separately certify three convolutional checkpoints from the same rational architecture family.  
 200 Their vision stacks contain no attention, so we independently rebuild all eight joint modules and all  
 201 twelve heads per module from the saved weights. Sixteen deterministically selected cache inputs  
 202 per checkpoint yield 4,608 head-input cases. Every output is finite and the worst relative  $\ell_2$  error is  
 203  $6.36 \times 10^{-16}$ , below the frozen  $10^{-6}$  gate. On the same inputs, independent CP-factor evaluation  
 204 covers all 384 joint-FFN module-input cases at worst relative error  $1.3871 \times 10^{-15}$  against a  $10^{-10}$   
 205 gate. Independent complete-block equations cover all 384 block-input cases at worst relative error  
 206  $1.5024 \times 10^{-7}$  against a  $10^{-4}$  gate. The complete-block comparison crosses a deployed PyTorch  
 207 float32 block call and a NumPy float64 reconstruction. It is input-conditioned and per-block, not a  
 208 sequential end-to-end reconstruction of the joint transformer or policy.

209 A separately frozen ViT audit partitions joint attention into vision, instruction, robot-state, and  
 210 action-query source tokens. Across 128 deterministic, official-task-balanced inputs for each of the  
 211 three capability checkpoints, all 36,864 head-input cases and 147,456 source-group evaluations are  
 212 finite, with worst reconstruction error  $5.0741 \times 10^{-7}$  against a  $10^{-6}$  gate. Across the 3,072 projected  
 213 module-input cases, mean coherent-energy shares are 67.7% vision, 19.2% robot state, 7.9% action  
 214 query, and 5.1% instruction. These input-conditioned magnitudes are descriptive rather than causal  
 215 importance. Familiar-prompt permutations are also descriptive and do not establish grounding or  
 216 closed-loop success.

217 Finally, a prospectively frozen sequential replay reaches the learned action output on the same three  
 218 capable ViT checkpoints. For 16 deterministic, official-task-balanced corrected inputs per checkpoint,  
 219 an independent NumPy float64 evaluation starts from the exact prepared model tensors and rebuilds  
 220 the learned patch projection, four vision blocks, multimodal embeddings and projections, eight  
 221 joint blocks, final RationalNorm, and linear action head. Across 48 inputs and 816 recorded stage  
 222 evaluations, the worst final-action relative  $\ell_2$  error is  $3.9822 \times 10^{-7}$  against a frozen  $10^{-3}$  gate.  
 223 Every manual PyTorch traversal is bitwise identical to an unmodified full model call. This is an  
 224 input-conditioned full learned-forward numerical certificate, not one compact symbolic contraction  
 225 or an input-general proof. It excludes preprocessing, fixed action unnormalization, chunk execution,  
 226 simulator dynamics, and success.

227 Naively materializing the real width-384 coefficient tensor would require approximately  $3.2 \times 10^{15}$   
 228 values. We instead contract the CP factors into an  $O(d^2)$  reduced Gram. This smaller matrix plays  
 229 the role of a covariance matrix: its leading eigenvectors rank the directions that carry the most  
 230 coefficient energy without constructing the full table. At toy multi-head scale, it matches brute-force  
 231 materialization to  $2.13 \times 10^{-14}$ .

## 232 4.2 Trained policy result

233 The reconstruction certificate above spans three ViT checkpoints. The following subspace-ranking  
 234 experiment is a descriptive analysis of the seed-0 ViT checkpoint. We apply the exact weight-derived  
 235 calculation to all twelve heads in early, middle-late, and final backbone blocks 0, 6, and 7. Random  
 236 projector banks are sampled once, then fixed and reused for all evaluation examples. The weights  
 237 choose the subspace. Real cached activations are used only afterward to measure how faithfully that  
 238 subspace preserves the layer’s computation.




Figure 5: **Exact weight-derived attention structure.** Ratios above one favor the exact weight-derived subspace over a random equal-rank subspace. Each point in the right panel is one trained attention head.

At ranks 8, 32, and 128, respectively, 80.6%, 91.7%, and 100% of the 36 block-head pairs favor the exact subspace. Their median ratios are 1.028, 1.171, and 2.441. At rank 128, the mean ratio is 3.033 and the strongest head reaches 10.135. The rank-128 block means are 3.08, 3.45, and 2.56 for blocks 0, 6, and 7. The 36 block-head pairs are architectural units within one checkpoint, not independent model-level replicates.

This extends exact weight-only construction through individual attention layers. It is not an exact compact symbolic end-to-end decomposition of the complete policy. Symbolic composition across all eight blocks still faces rapid degree and representation growth.

A separate intervention tests behavioral sufficiency at the visual-to-joint interface rather than deriving a projector from the exact attention factors. An action-Jacobian projector estimated on official Object tasks 0–3, with rank 96 fixed before corrected outcomes, retains 259/272 full-policy successes on tasks 8–9 across three ViT checkpoints. These tasks are absent from corrected projector construction. An activation-energy projector retains 267/272, while three equal-rank random projectors retain 37/272, 9/272, and 22/272. All ten tasks appeared during policy training. This establishes a structured low-dimensional visual bottleneck relative to random projectors, not unseen-task generalization, unique action-Jacobian specificity, or a causal link to the exact weight factors. Appendix A.4 reports the paired controls and stability intervals.

## 5 Related work

**Capability, mechanisms, and robustness.** OpenVLA and OpenVLA-OFT establish pretrained VLA and continuous-control baselines [1, 2], while  $\chi$ -VLA studies a compact algebraically restricted specialist. Polynomial and tensor networks provide the multilinear foundation [12–14]. Bilinear MLPs and  $\chi$ -nets expose low-rank features and compositional circuits from trained weights [3, 4]. We extend that setting to closed-loop control and exact reduced softmax-free attention. Bilinear policy analysis and recent activation, circuit, and rollout studies make our distinction empirical rather than universal [5–8]. Because LIBERO success can coexist with weak robustness and language use [9–11], our claims remain bounded to the specialist setting.

## 6 Limitations and conclusion

$\chi$ -VLA shows that strong specialist control can coexist with exact access to trained structure. The paired controls establish familiar-task instruction necessity, not counterfactual goal following, broad generalization, or physical transfer. The learned path is replayed on fixed inputs, not given as one compact symbolic contraction or input-general proof, and the strongest ensemble costs three passes. Common ViT checkpoints connect capability, exact attention, and reduced-Gram analysis, while separate convolutional files replicate the attention equation certificate. We do not claim a coefficient-edit interface. The anonymous artifact and Appendix A.2 provide full reproducibility details.



## References

- [1] M. J. Kim et al. *OpenVLA: An Open-Source Vision-Language-Action Model*. CoRL, 2024.
- [2] M. J. Kim, C. Finn, and P. Liang. *Fine-Tuning Vision-Language-Action Models: Optimizing Speed and Success*. 2025.
- [3] A. Rigg, T. Dooms, M. T. Pearce, J. Oramas, and L. Sharkey. *Bilinear MLPs Enable Weight-Based Mechanistic Interpretability*. ICLR, 2025.
- [4] T. Dooms, W. Gauderis, G. A. Wiggins, and J. Oramas. *Compositionality Unlocks Deep Interpretable Models*. 2025.
- [5] F. Oozeer, Y. Eriskin, A. Rigg, J. Oramas, and L. Sharkey. *Bilinear Convolution Decomposition for Causal RL Interpretability*. 2024.
- [6] B. Häon, K. Stocking, I. Chuang, and C. Tomlin. *Mechanistic Interpretability for Steering Vision-Language-Action Models*. 2025.
- [7] S. Marks et al. *Sparse Feature Circuits: Discovering and Editing Interpretable Causal Graphs in Language Models*. ICLR, 2025.
- [8] M. Kekić et al. *Learning Nonlinear Causal Reductions to Explain Reinforcement Learning Policies*. ICLR, 2026.
- [9] X. Zhou et al. *LIBERO-PRO: Towards Robust and Fair Evaluation of Vision-Language-Action Models Beyond Memorization*. 2025.
- [10] Y. Fang et al. *When Vision Overrides Language: Evaluating and Mitigating Counterfactual Failures in VLAs*. 2026.
- [11] C. Kim et al. *LIBERO-Para: A Diagnostic Benchmark and Metrics for Paraphrase Robustness in VLA Models*. 2026.
- [12] T. Kolda and B. Bader. *Tensor Decompositions and Applications*. SIAM Review, 2009.
- [13] E. Stoudenmire and D. Schwab. *Supervised Learning with Tensor Networks*. NeurIPS, 2016.
- [14] G. Chrysos et al. *Deep Polynomial Neural Networks*. CVPR and TPAMI, 2020 to 2021.

## A Supporting results and decisions

### A.1 Causal isolation of the learned instruction-embedding path

Token IDs are lookup addresses. A learned embedding table turns those addresses into vectors that the policy can combine with vision and robot state. Before observing outcomes, we fixed a paired test that isolates this conversion. On canonical episodes 30–39, disjoint from the instruction-necessity episodes, both arms received the identical correct 32-token ID tensor. The ablated arm replaced only the output of the token-embedding lookup with exact zeros. Sequence length, token positions, vision, robot state, embodiment, the separate learned common BOS, action queries, transformer weights, and action head were preserved and hash-checked.




Figure 6: **Learned lexical-embedding-path necessity on familiar tasks.** Left: closed-loop success for the intact path and exact post-lookup zeroing, using 100 paired states per checkpoint and 300 pooled pairs. Right: every task aggregate has a positive paired gap. The pooled marker shows the frozen task-stratified state-cluster 95% bootstrap interval.

308 The intact path succeeds on  $263/300 = 87.7\%$  of checkpoint-state pairs, compared with  $63/300 =$   
 309  $21.0\%$  after zeroing the learned token vectors. The paired gap is 66.7 points, with 203 intact-only  
 310 and 3 zeroed-only outcomes. Full-path success is 84%, 87%, and 92% by checkpoint, and the  
 311 corresponding gaps are 78, 57, and 65 points. All ten task aggregates are positive. A pooled one-sided  
 312 exact paired test gives  $p = 1.42 \times 10^{-56}$ , but this test is not cluster-robust because checkpoints  
 313 reuse the same 100 states. The frozen 20,000-draw task-stratified state-cluster bootstrap interval is  
 314  $[61.7, 71.3]$  points and is our cluster-aware inference. All eight prospectively fixed gates pass.

315 This establishes that the learned lexical-embedding path is necessary for these familiar LIBERO-  
 316 Object tasks and three fixed policies. It does not establish semantic grounding, paraphrase robustness,  
 317 unseen-object or compositional generalization, cross-suite or physical transfer, or a benefit unique to  
 318 tensor decomposability.

## 319 A.2 Reproducibility details

320 Training and evaluation use fixed task lists, explicit camera and action conventions, canonical-  
 321 state manifests, checkpointed moving-average weights, and per-run result files. Historical Modal  
 322 cache-backed analyses read dataset-native task metadata. The replacement Athena pipeline pins that  
 323 metadata and maps dataset identifiers to official identifiers before defining task splits. A source-level  
 324 audit compared all 271,996 cached frames across Object, Spatial, Goal, and LIBERO-10 with their  
 325 pinned LeRobot revisions. Task identifiers, language joins, states, actions, and resized images matched  
 326 exactly for every frame. The anonymous artifact releases raw structural-certificate JSONs and all  
 327 2,000 episode rows behind the ViT member and ensemble results. Immutable SHA-256 manifests  
 328 and standard-library verifiers independently recompute the capability totals, visual-bottleneck counts  
 329 and intervals, and the attention, FFN, complete-block, modality-partition, full learned-forward, and  
 330 RationalNorm certificate decisions. Preflight checks fail if canonical initial states cannot be loaded.

## 331 A.3 Expanded attention screen

332 Within the same seed-0 ViT checkpoint, the corrected fixed-control attention analysis covers all  
 333 eight blocks and twelve heads. At rank 128, 83/96 block-head pairs favor the exact weight-derived  
 334 subspace. The median ratio is 1.94, the mean is 2.58, and blocks 2 and 3 provide useful depth  
 335 controls. At ranks 8 and 32, only 54/96 and 60/96 pairs favor the exact subspace. The effect is  
 336 therefore depth-dependent and strongest at the wider tested rank.

## 337 A.4 Prospectively fixed-rank visual bottleneck

338 The behavioral intervention acts at the width-384 visual-to-joint interface and keeps rank 96. Projector  
 339 construction uses official Object tasks 0–3. Rank 96 was fixed before any corrected-basis outcome  
 340 was observed, and the frozen evaluation uses tasks 8–9. Across three checkpoints and 300 canonical  
 341 episodes, the full policy succeeds on 272 trials. The action-Jacobian projector retains  $259/272 =$

342 95.2% of those successes, with a paired checkpoint-task-episode resampling interval of 91.2–98.6%.  
 343 The activation-energy projector retains  $267/272 = 98.2\%$ , with an interval of 95.8–100.0%. Three  
 344 fixed equal-rank random projectors retain  $37/272 = 13.6\%$ ,  $9/272 = 3.3\%$ , and  $22/272 = 8.1\%$ .




Figure 7: **Corrected closed-loop visual bottleneck.** Both learned rank-96 projectors preserve nearly all baseline successes across three checkpoints, while equal-rank random projectors do not. Error bars in panel (a) are deterministic paired cluster-resampling intervals. Tasks 8–9 are absent from corrected projector construction, but not from policy training. This is not a clean rank-selection-holdout result.

345 The action-Jacobian basis does not outperform the matched activation-energy basis. The stronger  
 346 action-specificity gate therefore does not pass. The supported claim is structured visual sufficiency  
 347 relative to random projectors. This intervention is data-derived and remains separate from the exact  
 348 weight-only attention decomposition.

### 349 A.5 Numerical scope of rational conversion

350 The deployed degree- $[2/2]$  least-squares rational fit is exact as a rational computation even though it  
 351 approximates RMS normalization. Its numerator has coefficients (5.32997, 7.53552, 0.327274) and  
 352 the denominator has coefficients (1, 9.64827, 2.60664). Every denominator coefficient is positive,  
 353 which certifies  $Q(v) > 0$  for all  $v \geq 0$ . We instrument all 53 RationalNorm module instances,  
 354 each invoked once per forward pass, on 4,096 deterministic inputs for each of three convolutional  
 355 checkpoints. Across 544,542,720 activation rows, the minimum observed denominator is 1.0016  
 356 and no row is nonfinite. Comparing deployed float32 normalization with an otherwise identical pass  
 357 that recomputes only RationalNorm scales in float64, the largest checkpoint-level primary metric,  
 358 defined as the maximum of normalized- and physical-action NRMSE, is  $2.414 \times 10^{-6}$  against a  
 359 frozen  $10^{-3}$  gate. This primary certificate concerns pole-free denominators and floating-point fidelity  
 360 to the deployed rational operator. It does not claim uniform agreement with exact RMS scaling or  
 361 measure matched alternative-normalizer runtime overhead.

362 The reconstruction ladder extends from individual attention heads through FFNs and per-block replay  
 363 to a separate sequential learned-forward certificate. That final record starts at exact prepared model  
 364 tensors and reaches the linear action head on 48 fixed inputs. It is numerical and input-conditioned,  
 365 not a compact symbolic contraction or an input-general full-policy proof. The ViT source-partition  
 366 audit remains layerwise and descriptive. Figure 8 reports the frozen numerical gates and the depth  
 367 profile without assigning causal importance.




Figure 8: **Frozen reconstruction ladder and descriptive ViT source partition.** Panel (a) reports observed-to-gate ratios for independent Conv attention, joint-FFN, and complete joint-block reconstruction, the ViT source-partition identity, the sequential ViT learned-forward action replay, and primary RationalNorm action fidelity. All plotted primary numerical gates pass. Panel (b) shows descriptive mean coherent-energy shares after the attention output projection across 384 checkpoint-input cases per joint block. The depth trend is not a causal importance or grounding result. The learned-forward replay excludes preprocessing, fixed action unnormalization, chunk execution, simulator dynamics, and success.