


Conversion cost is nonzero, while compilation narrows the systems gap



Each capability bar uses the same 500 canonical trials on one GPU class. The χ deficit is 4.4 points on A30 and 5.8 on A6000. Compilation narrows the latency ratio to 1.73–1.74x on both GPUs.