



Exact weight-only decomposition exposes learned attention structure



Real trained weights and real LIBERO activations. Each point in panel b is one attention head. Ratios above 1 favor the exact weight-derived subspace over a random subspace of the same rank.